Plant the seed. *Harvest the potential.*

News & Headlines

Computers can now paint like Van Gogh and Picasso

INDY/TECH

FACEBOOK'S ARTIFICIAL INTELLIGENCE ROBOTS SHUT DOWN AFTER THEY START TALKING TO EACH OTHER IN THEIR OWN LANGUAGE

China seeks dominance of global AI industry

Beijing challenges US with plan to create $150bn artificial intelligence sector

Artificial Intelligence

Google's AI has written some amazingly mournful poetry

Google's poetry was written by an AI system after it was fed thousands of unpublished romantic novels

Where we are...

Last time we examined Multi-Layered Perceptrons...

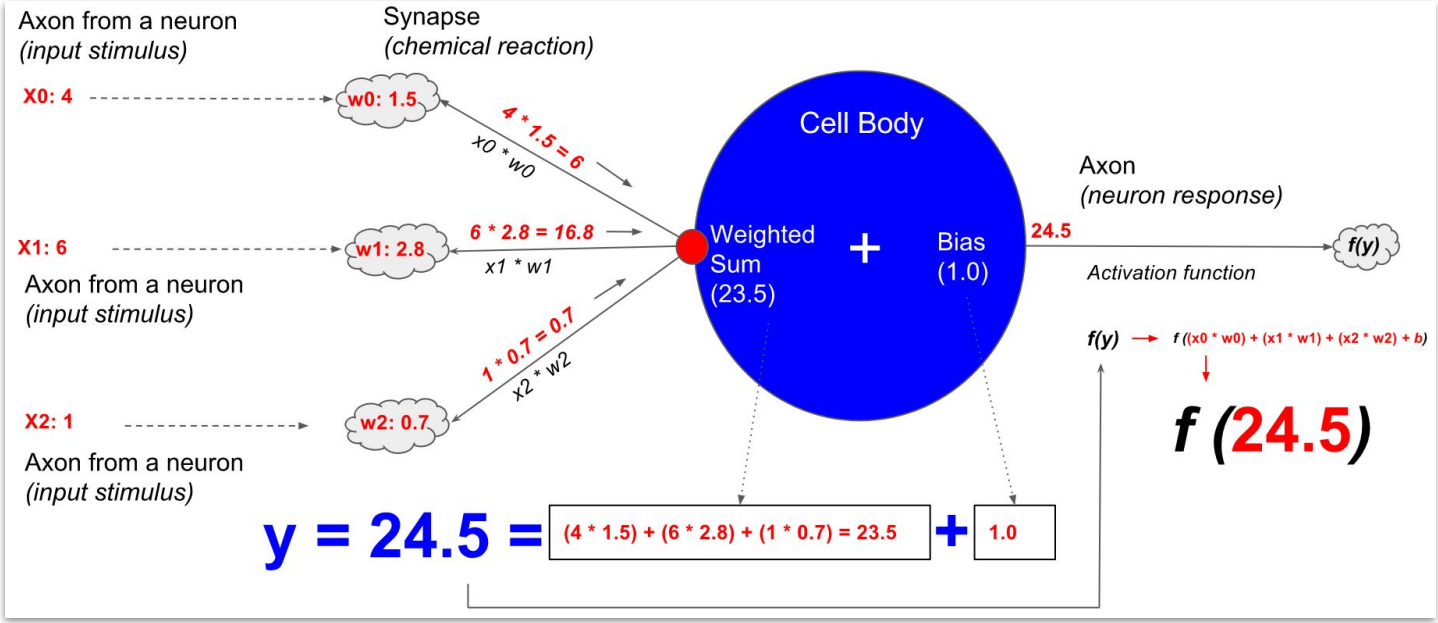

What is this?

"Tell me what you see"

Input Signal → *Neural Network* → *Prediction*

Dimensions: 640 x 480
Depth: 3 (Red, Green, Blue)

Neural Network (black box)

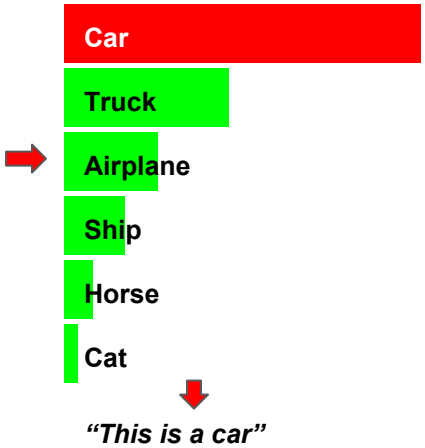

"This is a car"

What is this?

"Tell me what you see"

Input Signal → *Neural Network* → *Prediction*

Dimensions: 640 x 480
Depth: 3 (Red, Green, Blue)

Neural Network (black box)

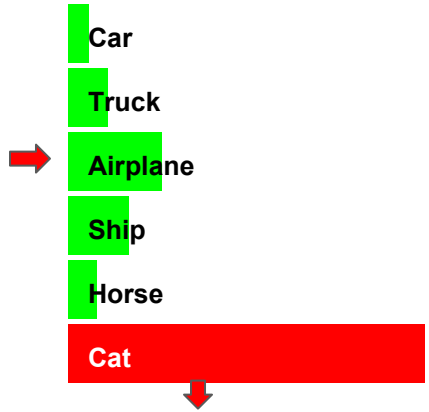

"This is a cat"

What is this?

"Tell me what you see"

Input Signal → *Neural Network* → *Prediction*

Dimensions: 640 x 480
Depth: 3 (Red, Green, Blue)

Neural Network (black box)

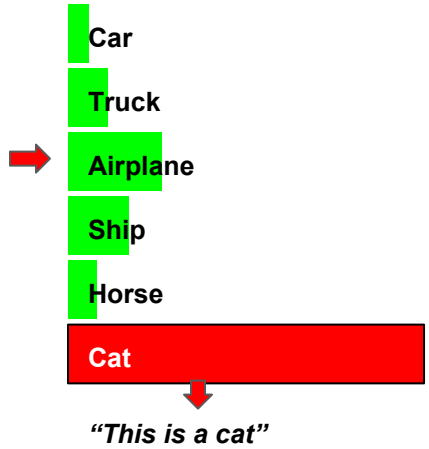

"This is a cat"

What is this?

"Tell me what you see"

Input Signal → *Neural Network* → *Prediction*

Dimensions: 640 x 480
Depth: 3 (Red, Green, Blue)

Neural Network (black box)

Classification Error

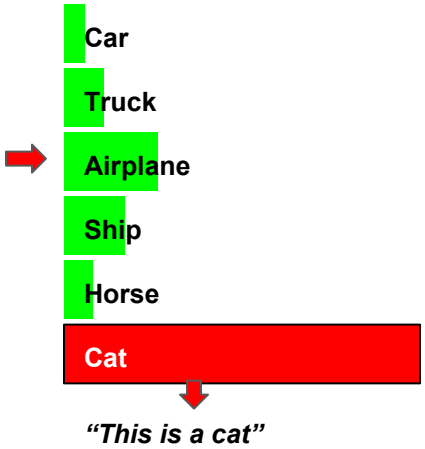

"This is a cat"

What is this?

My first shot: "I will try with a multi-layered perceptron…"

Assume the hidden layer is **2.8 times** bigger than the input layer.

Input Signal → ***Neural Network*** → ***Prediction***

Dimensions: 640 x 480
Depth: 3 (Red, Green, Blue)

640 x 480 x 3

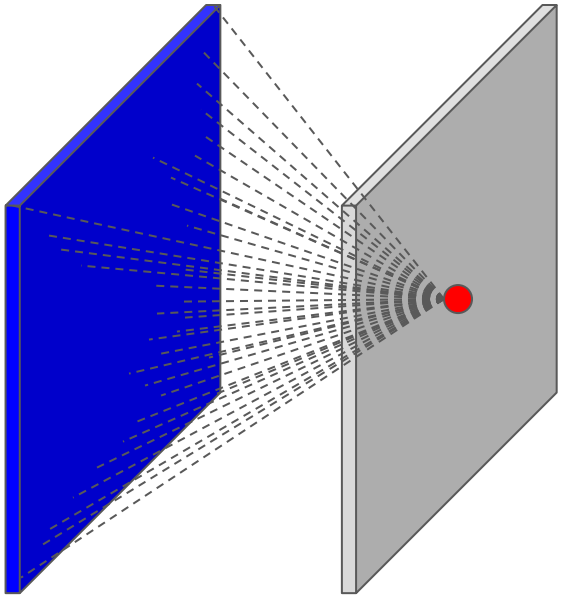

Input Layer
(921k neurons)

Hidden Layer
(2.5 million neurons)

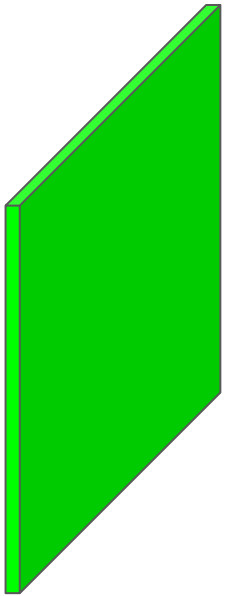

Output Layer

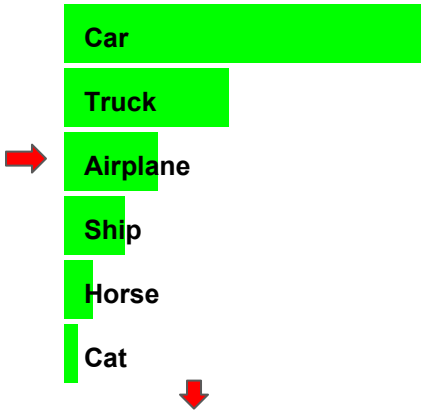

"This is a car"

What is this?

My first shot: "I will try with a multi-layered perceptron…"

Assume the hidden layer is **2.8 times** bigger than the input layer.

Input Signal → ***Neural Network*** → ***Prediction***

Dimensions: 640 x 480
Depth: 3 (Red, Green, Blue)

A minimum of 2.3 Trillion parameters!

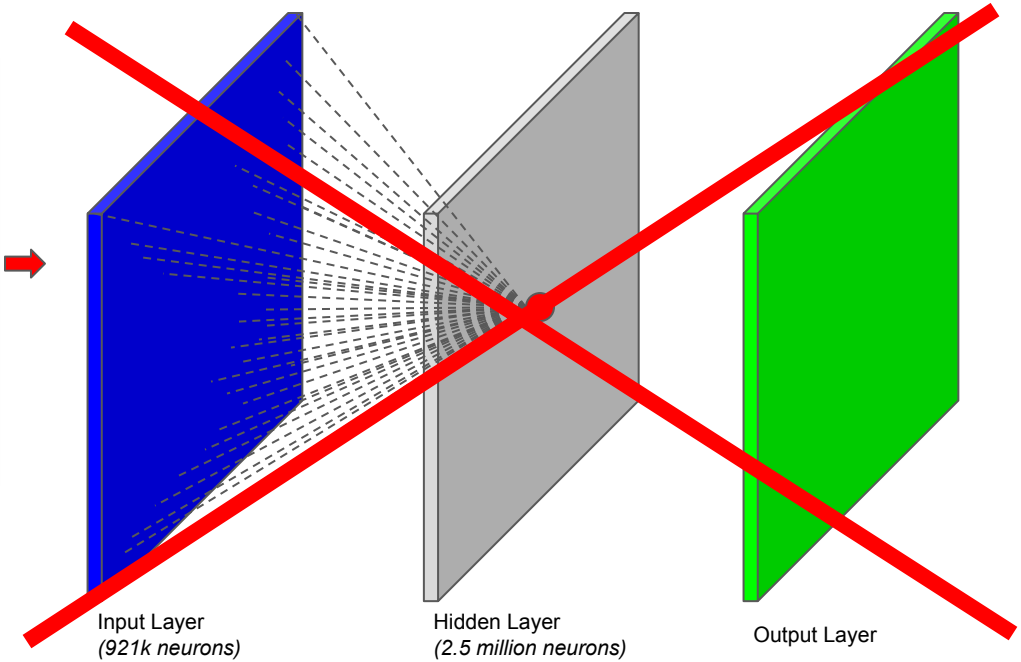

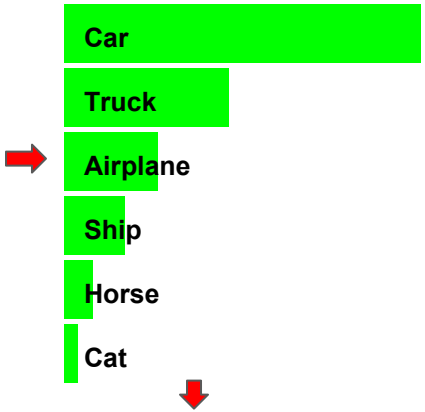

This won't work. Too many parameters!

What is this?

Let's try again → Second shot → "Reduce the input size and add more layers"

Assume the hidden layer is **2.8 times** bigger than the input layer.

Input Signal → ***Neural Network*** → ***Prediction***

Dimensions: 160 x 120
Depth: 3 (Red, Green, Blue)

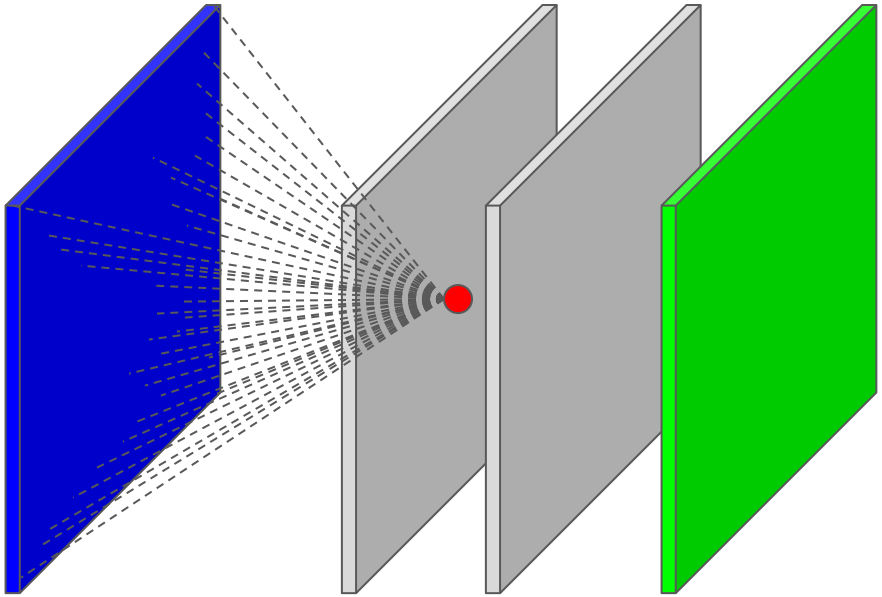

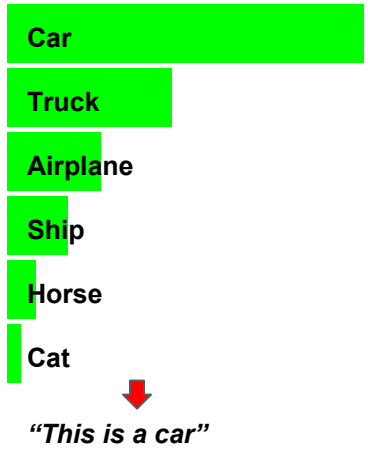

"This is a car"

What is this?

Let's try again → Second shot → "Reduce the input size and add more layers"

Assume the hidden layer is **2.8 times** bigger than the input layer.

Input Signal → ***Neural Network*** → ***Prediction***

Dimensions: 160 x 120
Depth: 3 (Red, Green, Blue)

A minimum of 9.2 Billion parameters!

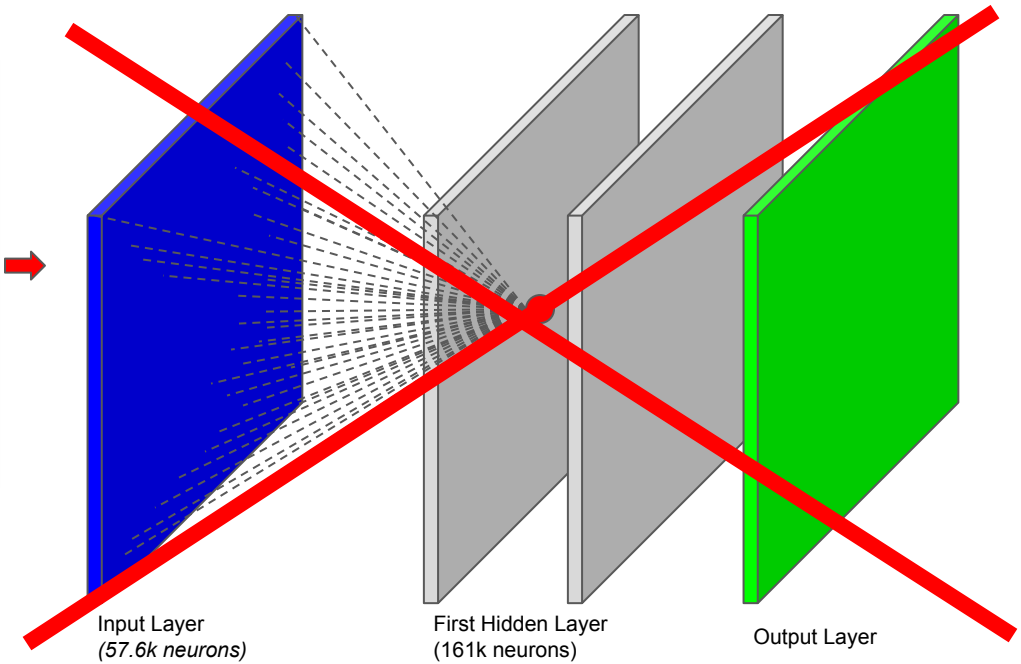

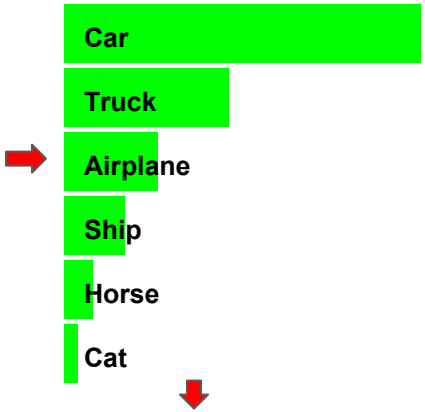

This won't work either. Too many parameters!

What is this?

One more time... → Third shot → "Reduce the input size again"

Assume the hidden layer is **2.8 times** bigger than the input layer.

Input Signal → ***Neural Network*** → ***Prediction***

Dimensions: 16 x 16
Depth: 3 (Red, Green, Blue)

The size of an icon

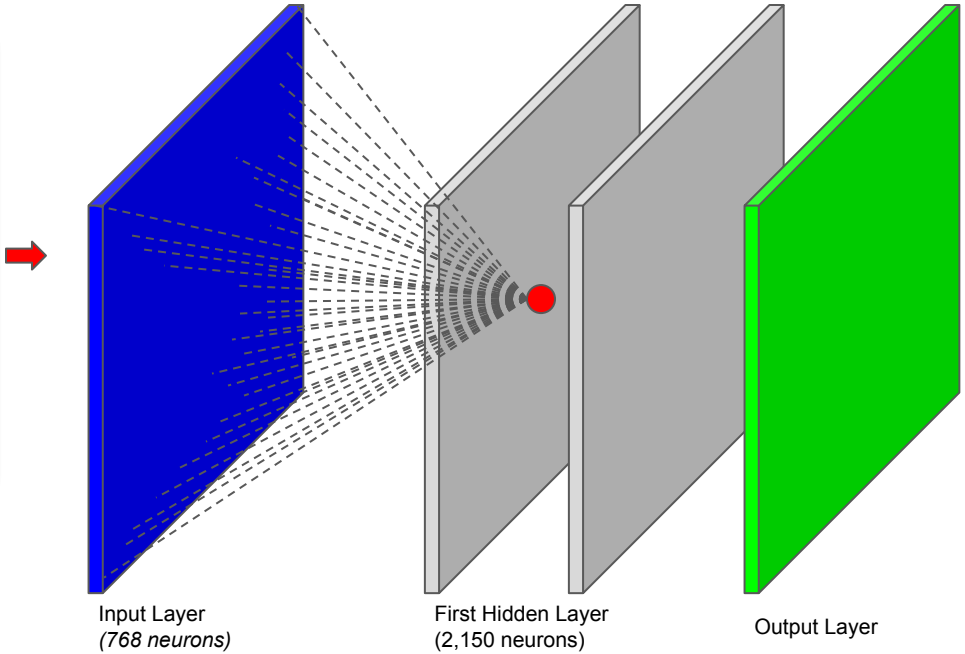

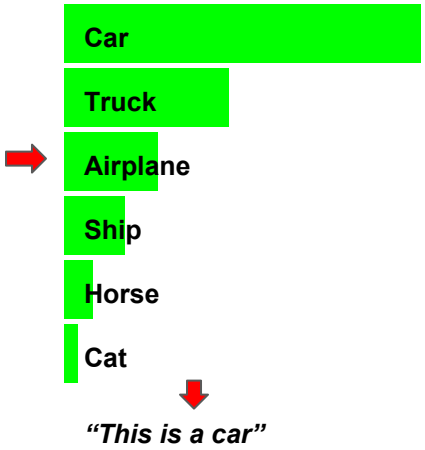

What is this?

One more time... → Third shot → "Reduce the input size again"

Assume the hidden layer is **2.8 times** bigger than the input layer.

Input Signal → ***Neural Network*** → ***Prediction***

Dimensions: 16 x 16
Depth: 3 (Red, Green, Blue)

A minimum of 1.6 Million parameters!

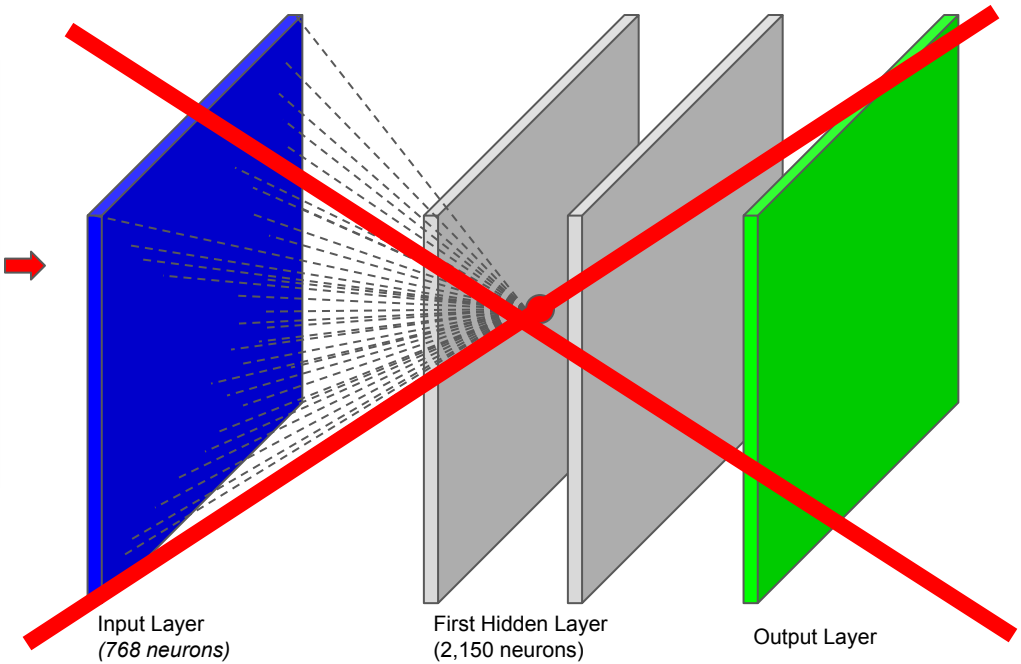

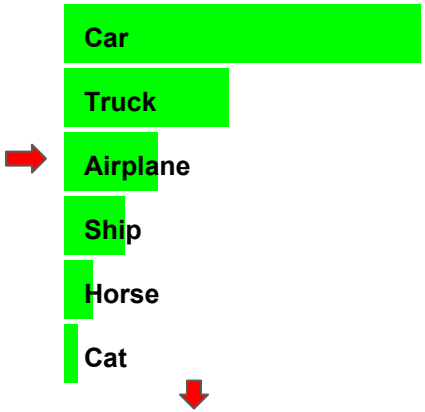

Still doesn't work. Too many parameters!

Too many parameters

So I can't solve this problem with an ordinary Multi-Layered Perceptron...

Input Signal → *Neural Network* → *Prediction*

Multi-Layered Perceptron:

1) *Too many parameters ("weights").*
2) *Lack of representational power, in this context.*
3) *Best case scenario their accuracy is extremely poor.*

...and this is just for starters.

It's just not the right tool for the job.

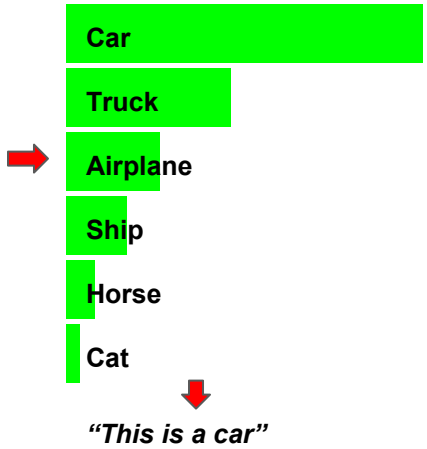

"This is a car"

Too many parameters

Perceptrons are not the right tool for the job. The number of parameters is a killer.

Too many parameters.

Width	Height	Depth	Total Inputs	Hidden Neurons	Input to Hidden Weights	Minimum Parameters
1024	768	3	2,359,296	6,606,029	15,585,577,323,725	15.5 trillions
800	600	3	1,440,000	4,032,000	5,806,080,000,000	5.8 trillions
640	480	3	921,600	2,580,480	2,378,170,368,000	2.3 trillions
320	240	3	230,400	645,120	148,635,648,000	148.6 billions
160	120	3	57,600	161,280	9,289,728,000	9.2 billions
80	60	3	14,400	40,320	580,608,000	580.6 million
32	32	3	3,072	8,602	26,424,115	26.4 million
16	16	3	768	2,150	1,651,507	1.6 million
8	8	3	192	538	103,219	100k

Assume the hidden layer is 2.8 times bigger than the input layer.

Convolutional Neural Networks

A *Convolutional Neural Network* is a sequence of convolutional layers, interspersed with activation functions.

They use **weight-sharing** in order to overcome the issue with parameters. It doesn't make sense to learn the same thing across all spatial locations.

Input Signal → ***Neural Network*** → ***Prediction***

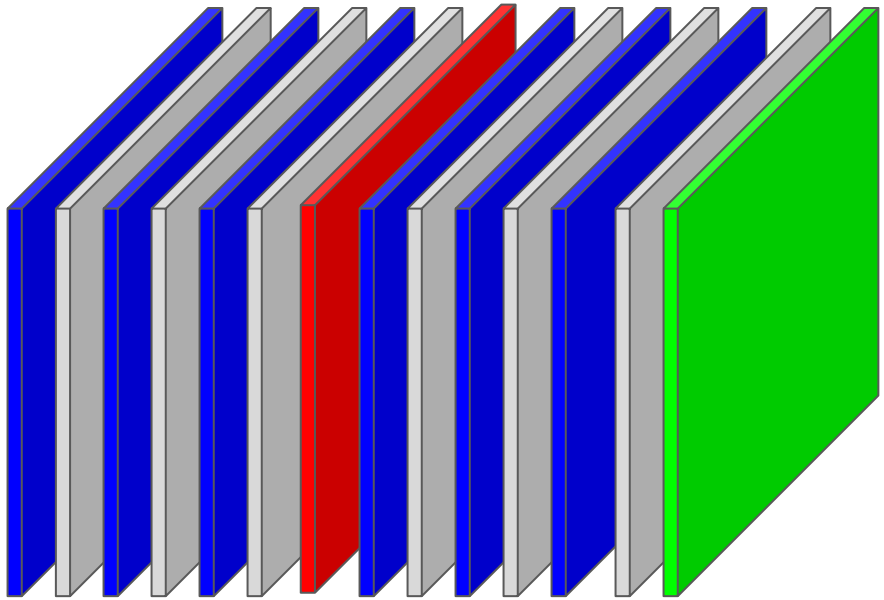

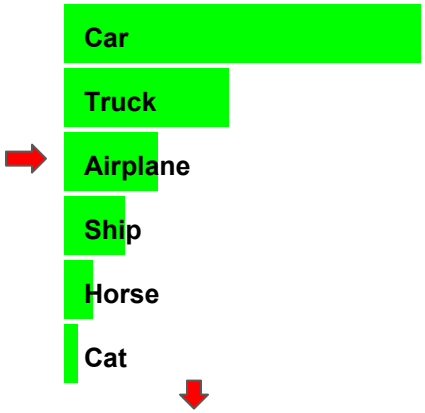

Conv
ReLU
Pooling
Fully Connected

Depth of the Network

"This network is 14 levels deep"

Convolutional Neural Networks

From low-level details to high-level objects...

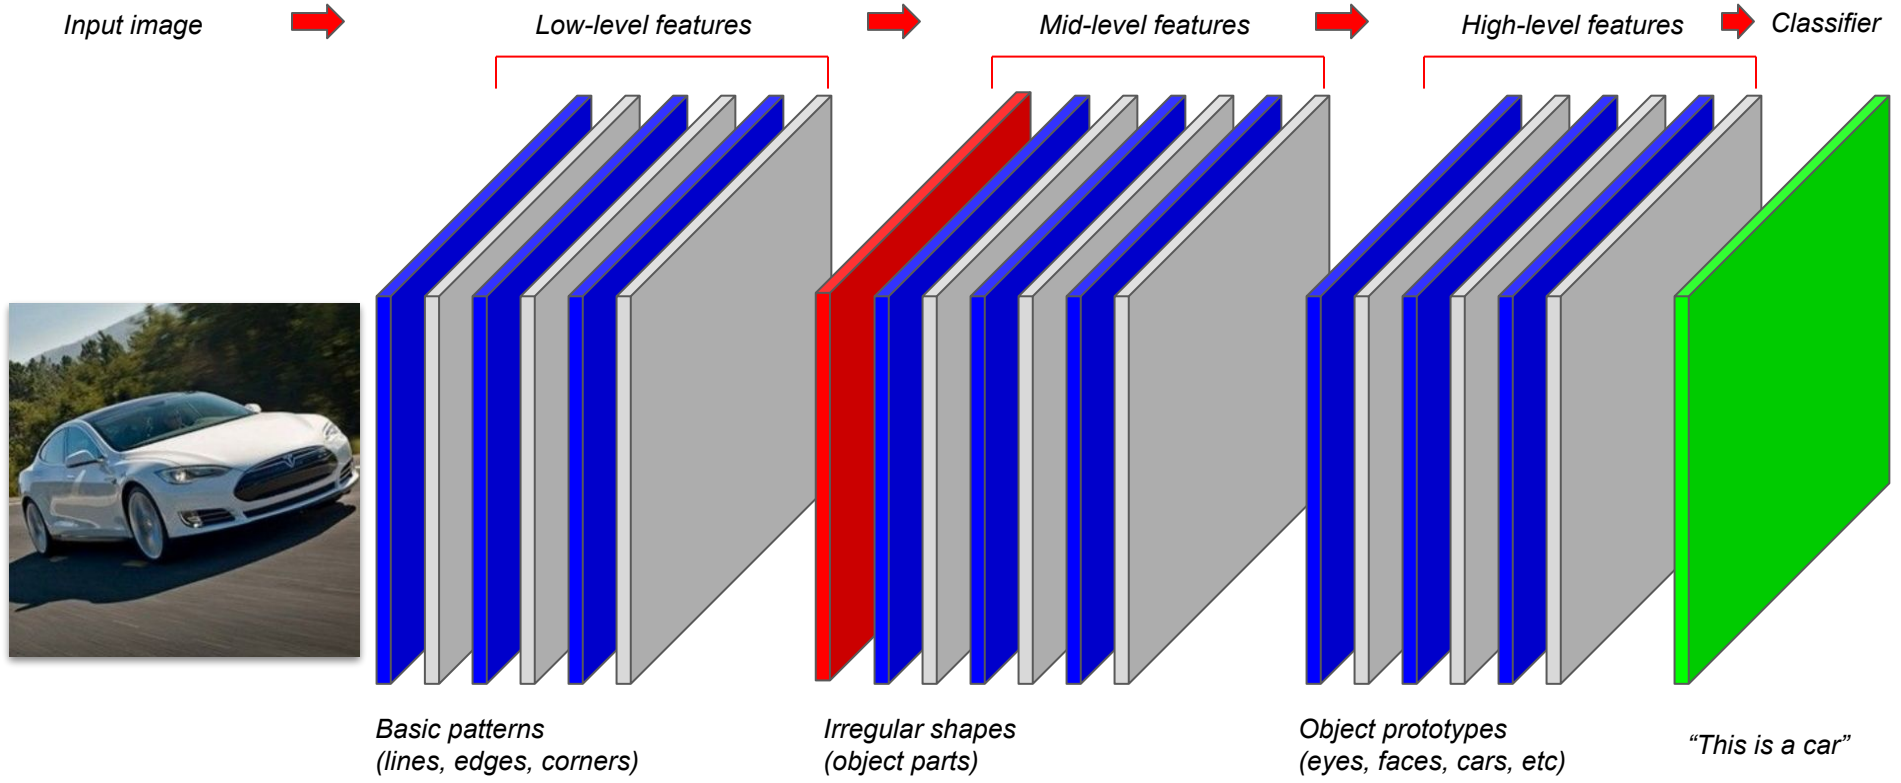

Convolutional Neural Networks

From low-level details to high-level objects...

Input image → *Low-level features* → *Mid-level features* → *High-level features* → *Classifier*

Basic patterns
(lines, edges, corners)

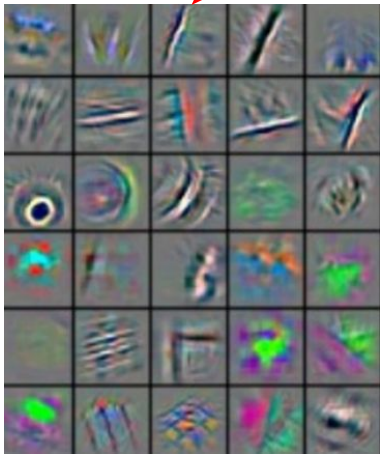

Irregular shapes
(object parts)

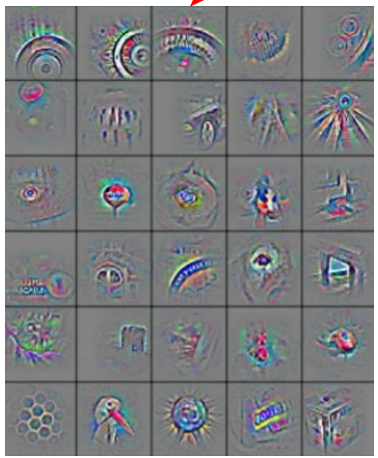

Object prototypes
(eyes, faces, cars, etc)

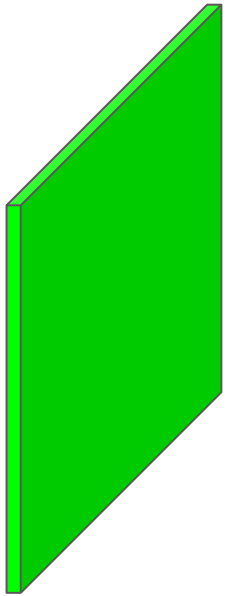

"This is a car"

Convolutional Neural Networks

Each neuron "*looks*" at a small portion of the input → *(Region Of Interest, or ROI for short)*
The ROI has a size *(could be 1x1, 3x3, 5x5, 7x7, etc.)* → This is the ***receptive field of the neuron***.

Input Signal → *Neural Network* → *Prediction*

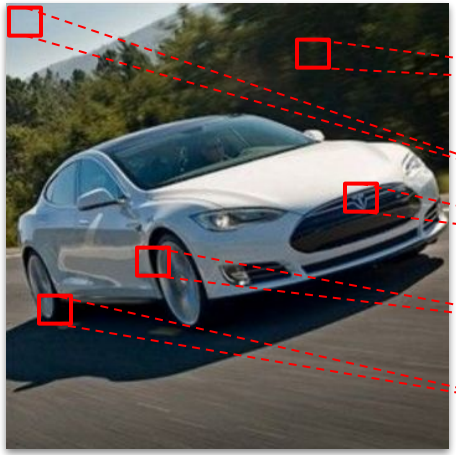

Convolve (slide) over all spatial locations

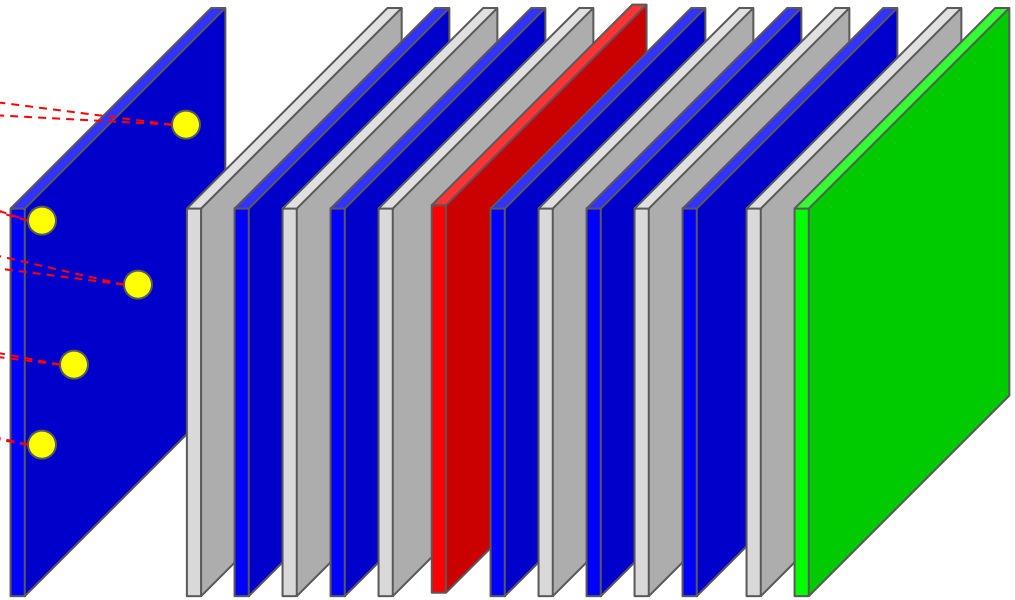

Depth of the Network

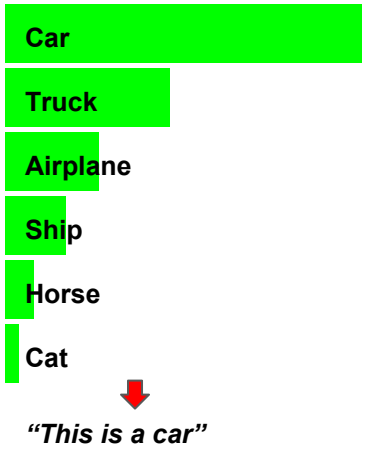

"This is a car"

Convolutional Neural Networks

The network will perform convolutions between an input patch and a set of filters

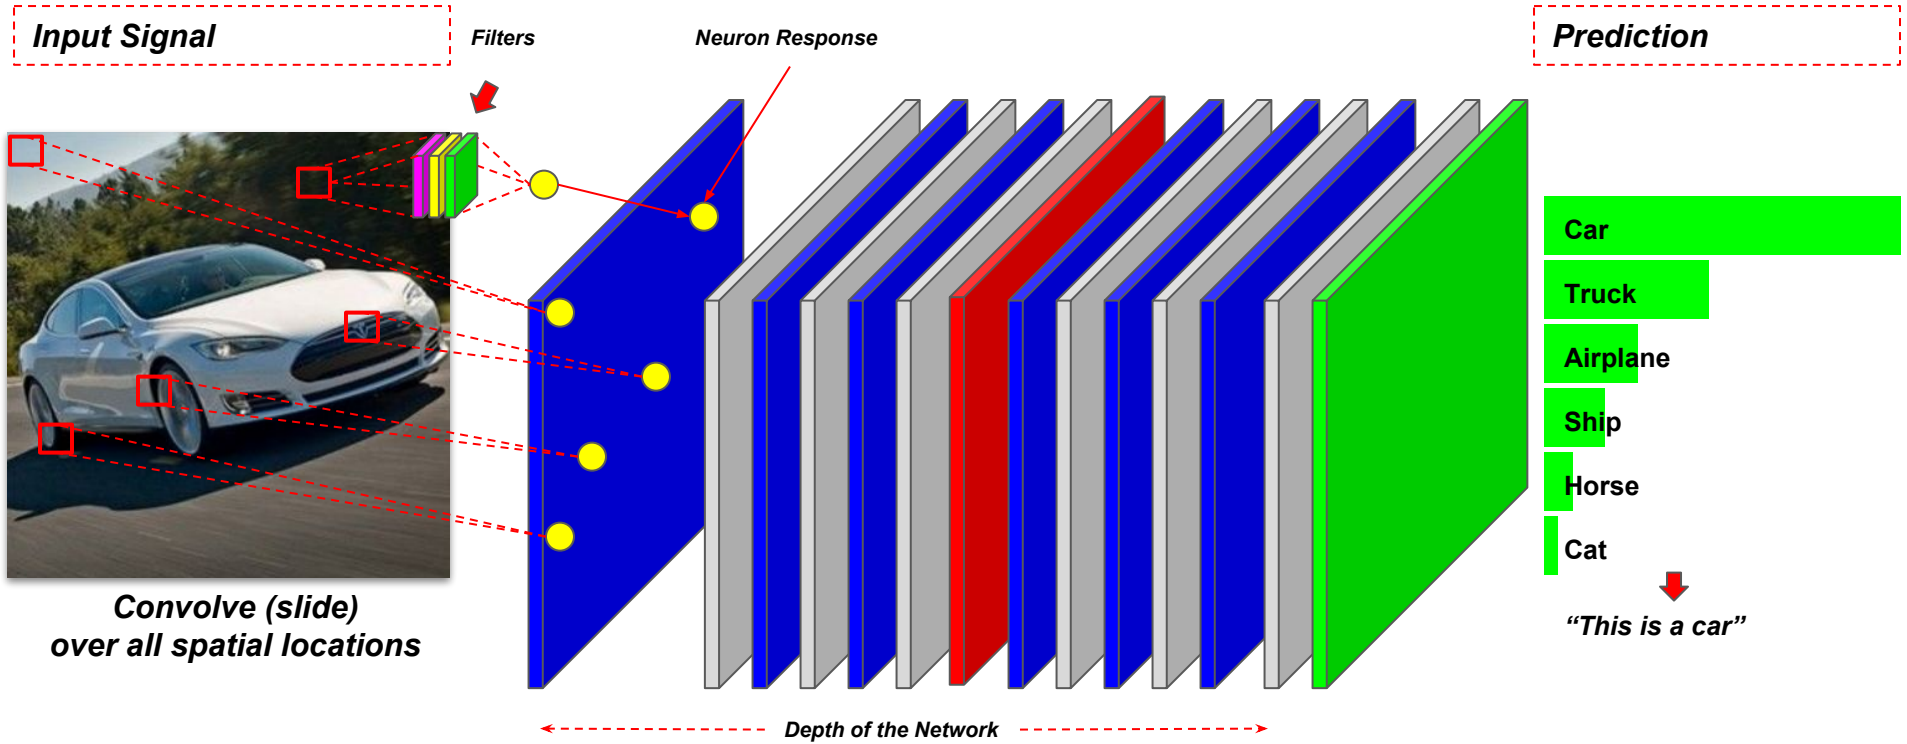

Convolutional Neural Networks

So to identify a car, the network will need to learn the *"weights"* required to produce each of the feature maps below.

Input Signal

These guys are called "Feature Maps"

Prediction

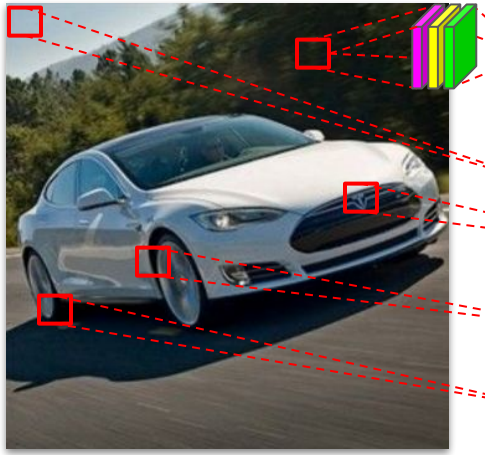

Convolve (slide) over all spatial locations

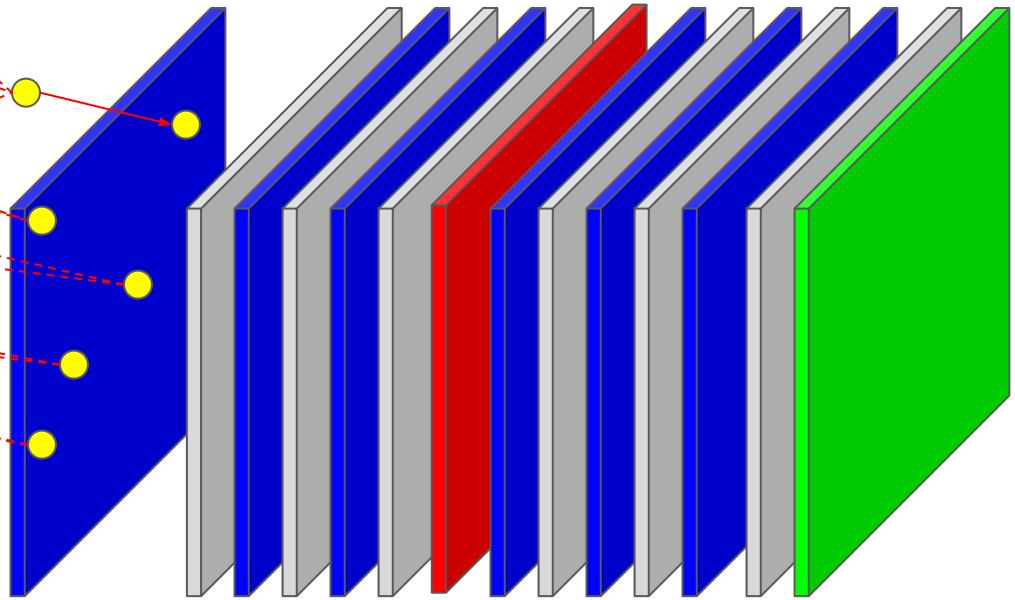

Depth of the Network

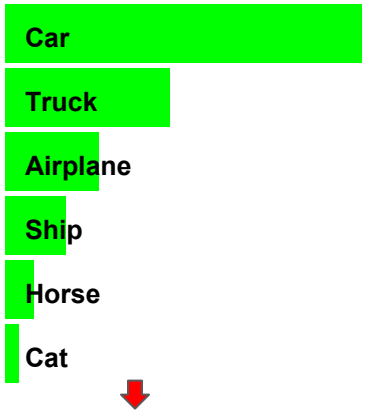

"This is a car"

Convolutional Neural Networks

So to identify a car, the network will need to learn the *"weights"* required to produce each of the feature maps below.

Input Signal

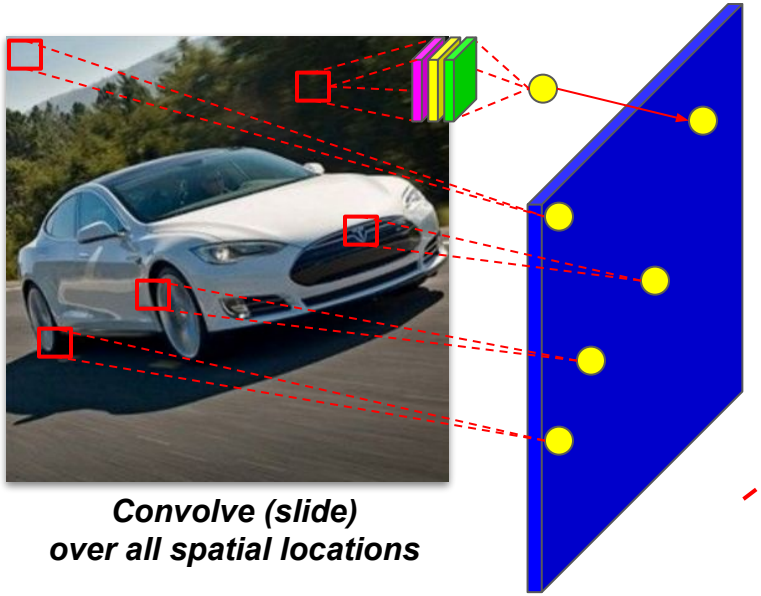

Convolve (slide) over all spatial locations

All of the neurons in the same layer share their weights.

This is called "*Weight sharing*"

Prediction

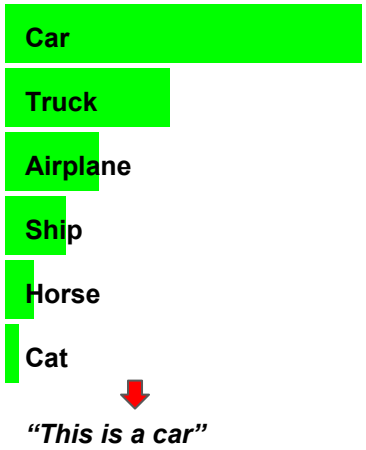

"This is a car"

Convolutional Neural Networks

Convolve (slide) over all spatial locations.

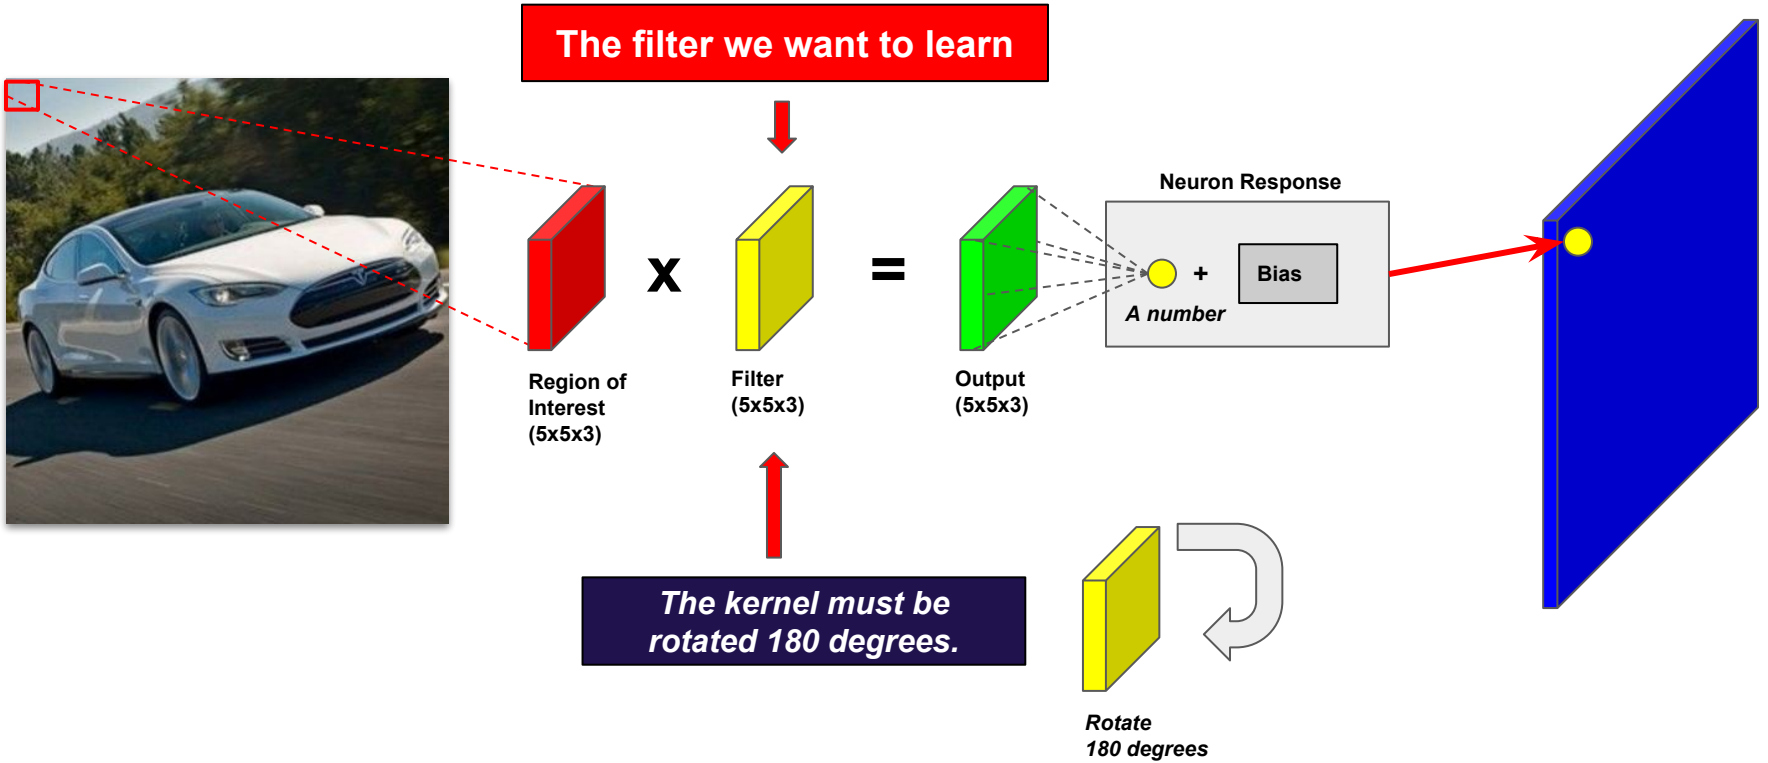

Convolutional Neural Networks

Convolve (slide) over all spatial locations.

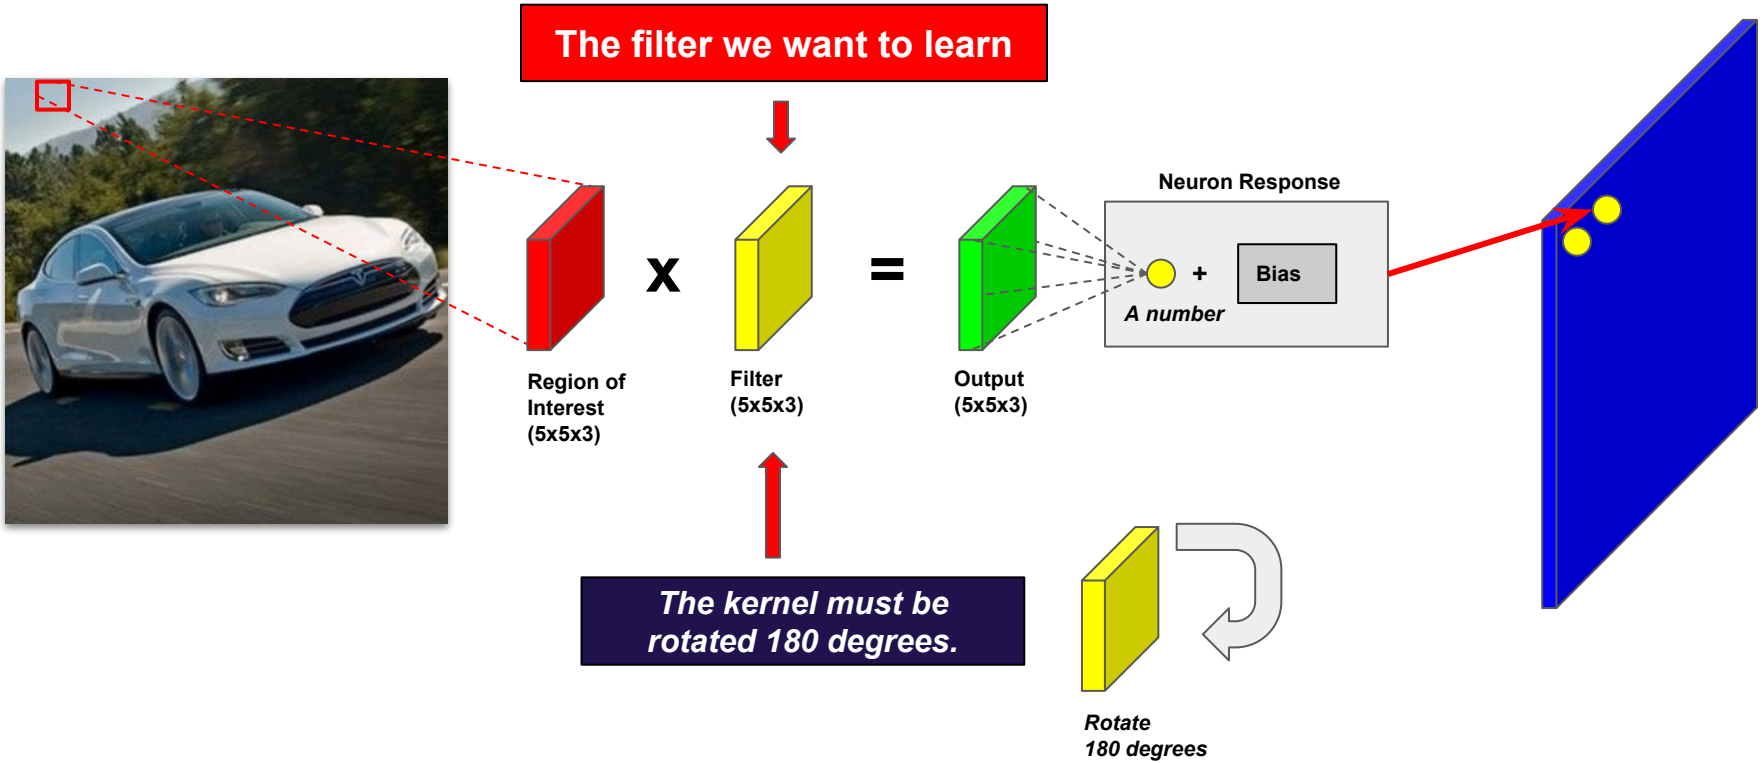

Convolutional Neural Networks

Convolve (slide) over all spatial locations.

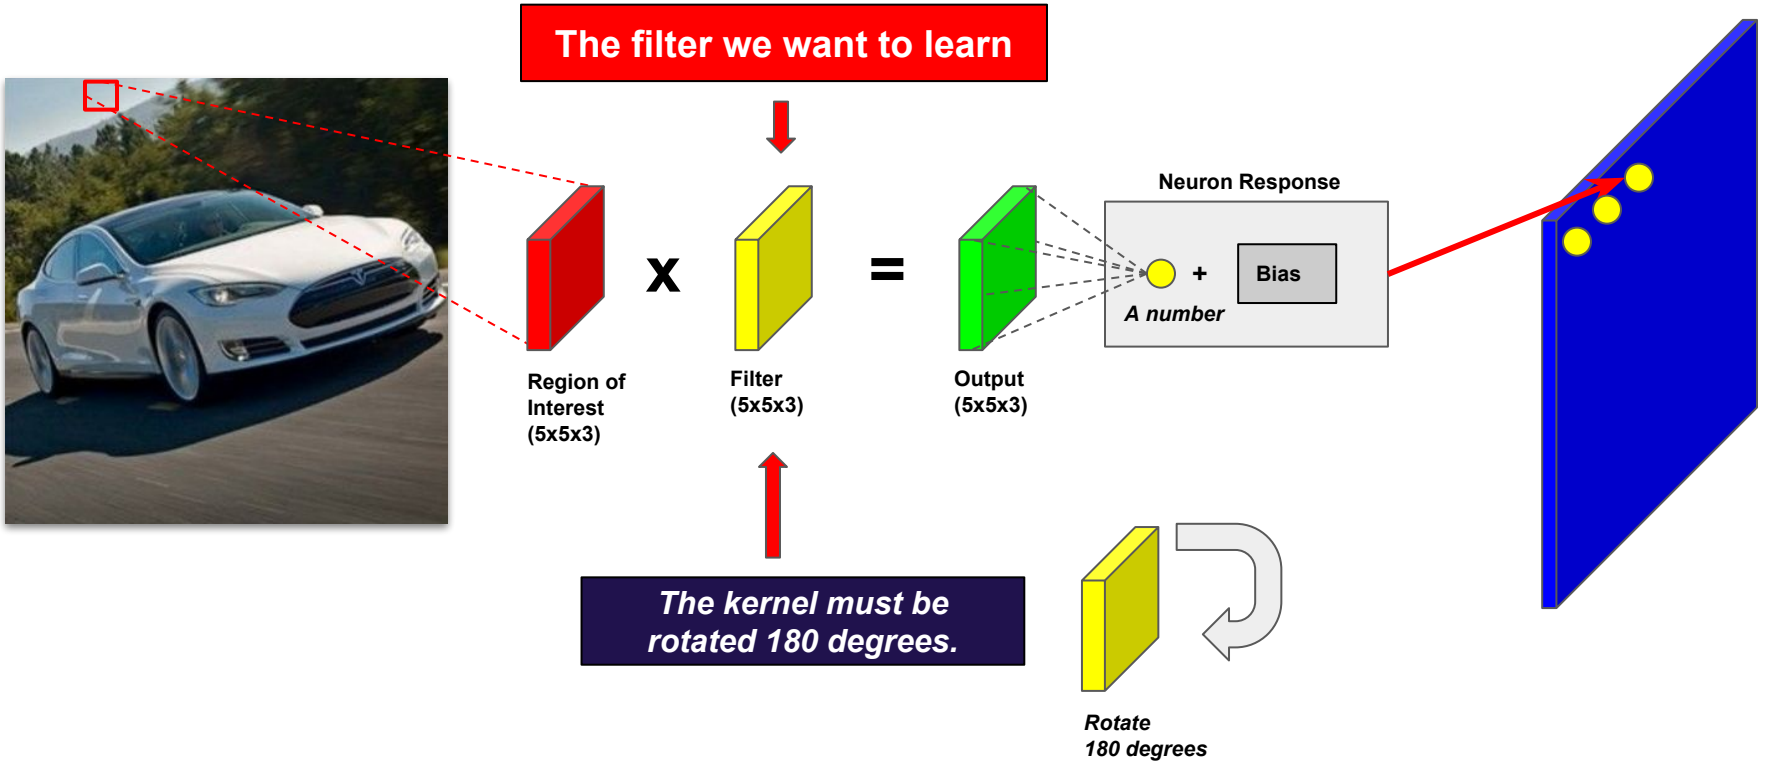

Convolutional Neural Networks

Convolve (slide) over all spatial locations.

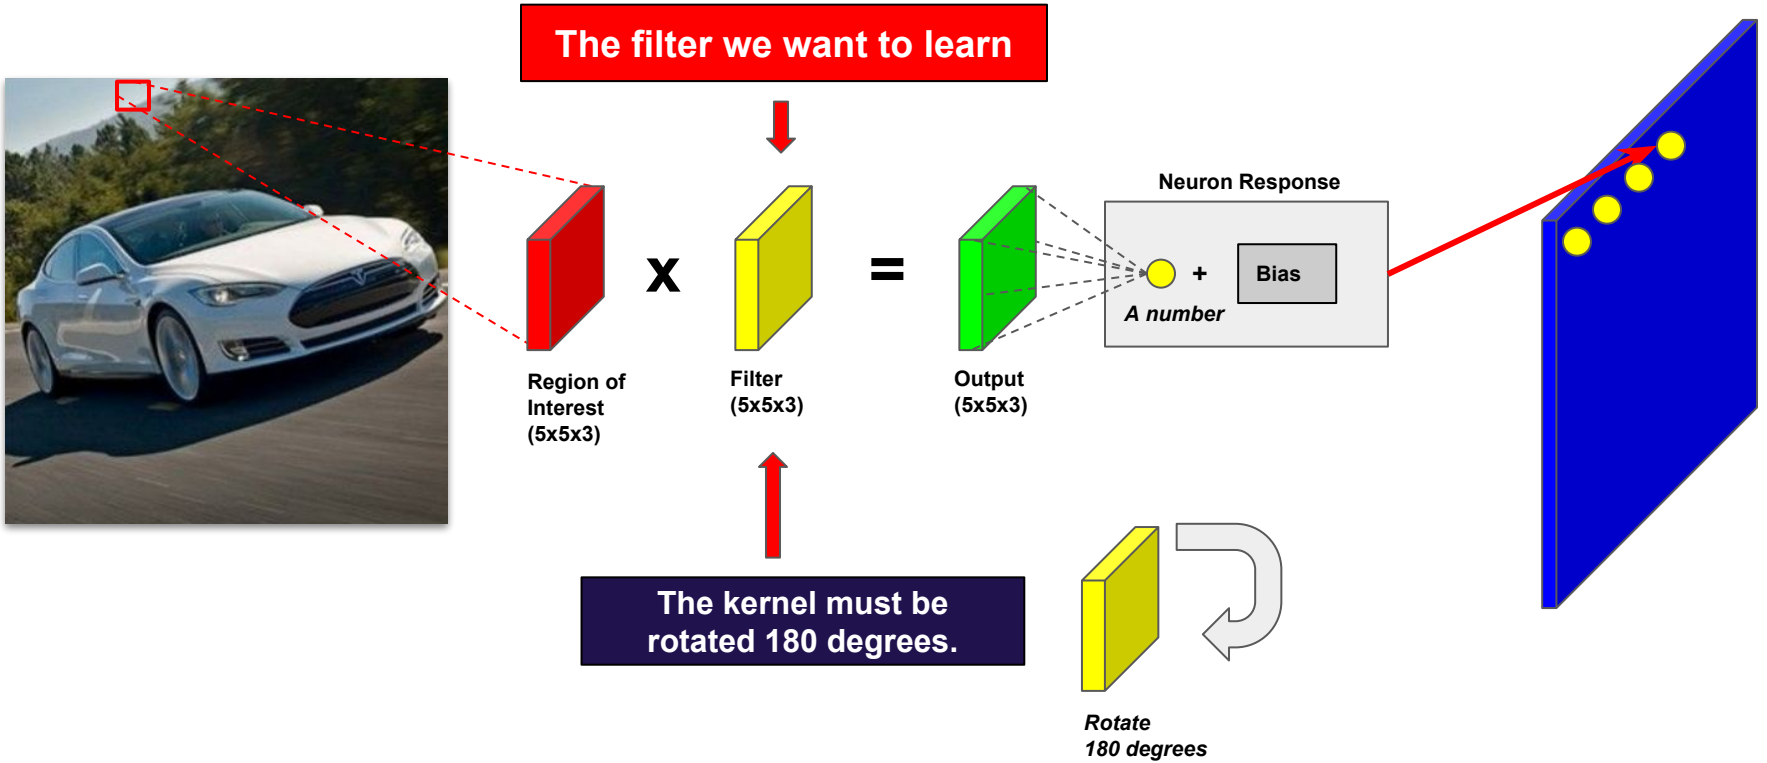

Convolutional Neural Networks

Before

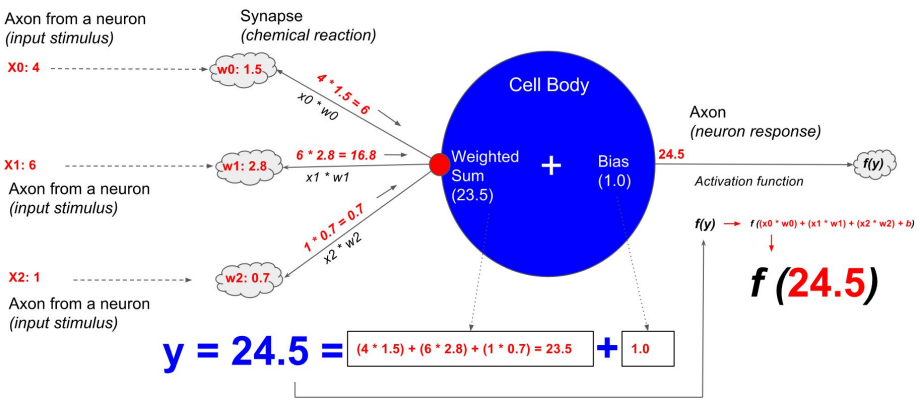

Now

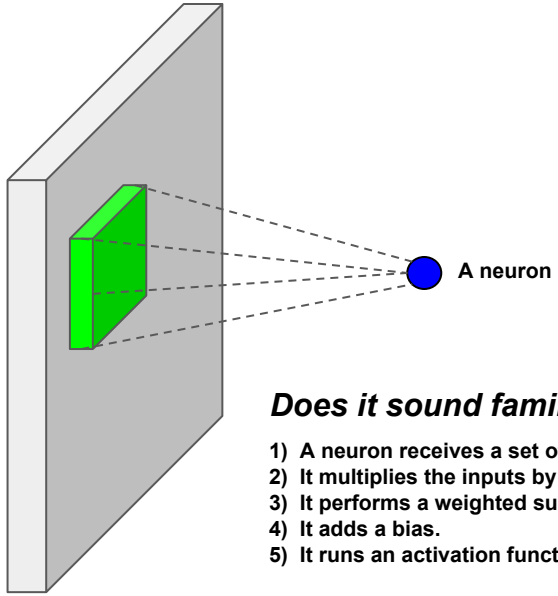

Does it sound familiar to you?

1) A neuron receives a set of inputs.
2) It multiplies the inputs by their respective weights.
3) It performs a weighted sum.
4) It adds a bias.
5) It runs an activation function.

Nothing changes!

Convolutional Layers

Say you've got an image. Let's pretend the image is in RGB format *(Red, Green, Blue)*

The image is 32 pixels width by 32 pixels height.

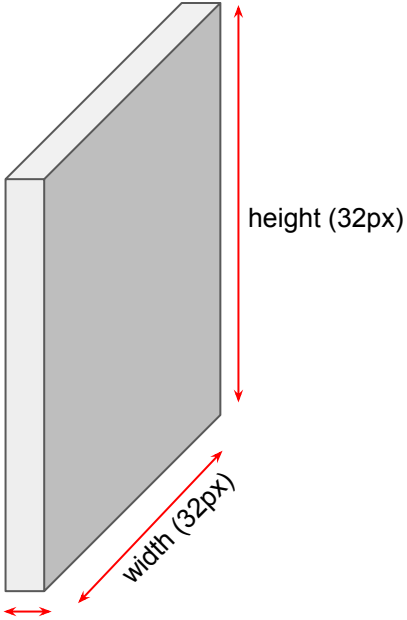

Depth = 3 channels
Red, Green and Blue

Original Image

Convolutional Layers

They produce *feature maps* by convolving an input volume with a bank of filters.

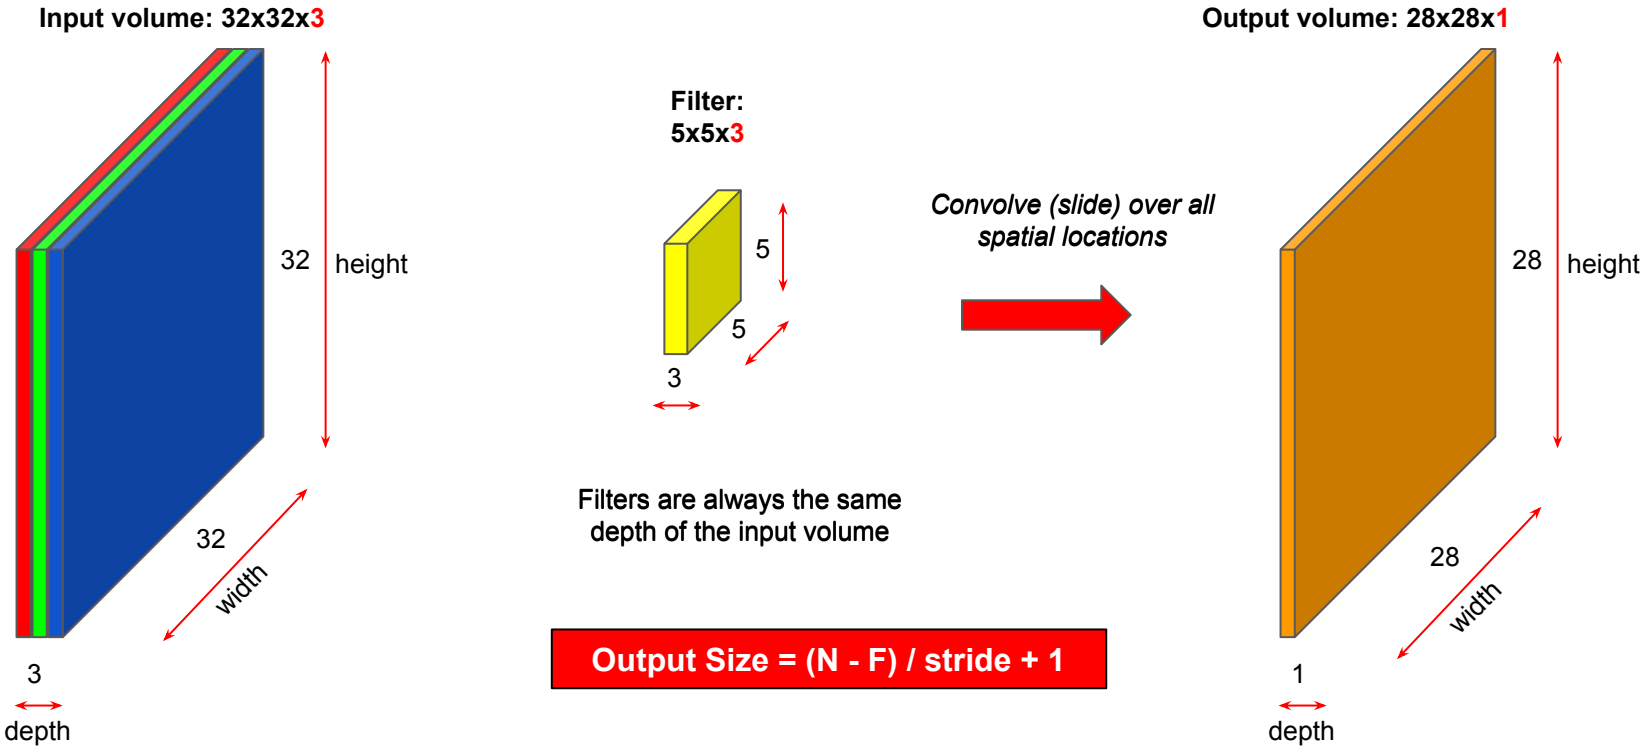

Convolutional Layers

Convolving with a second filter will produce another feature map

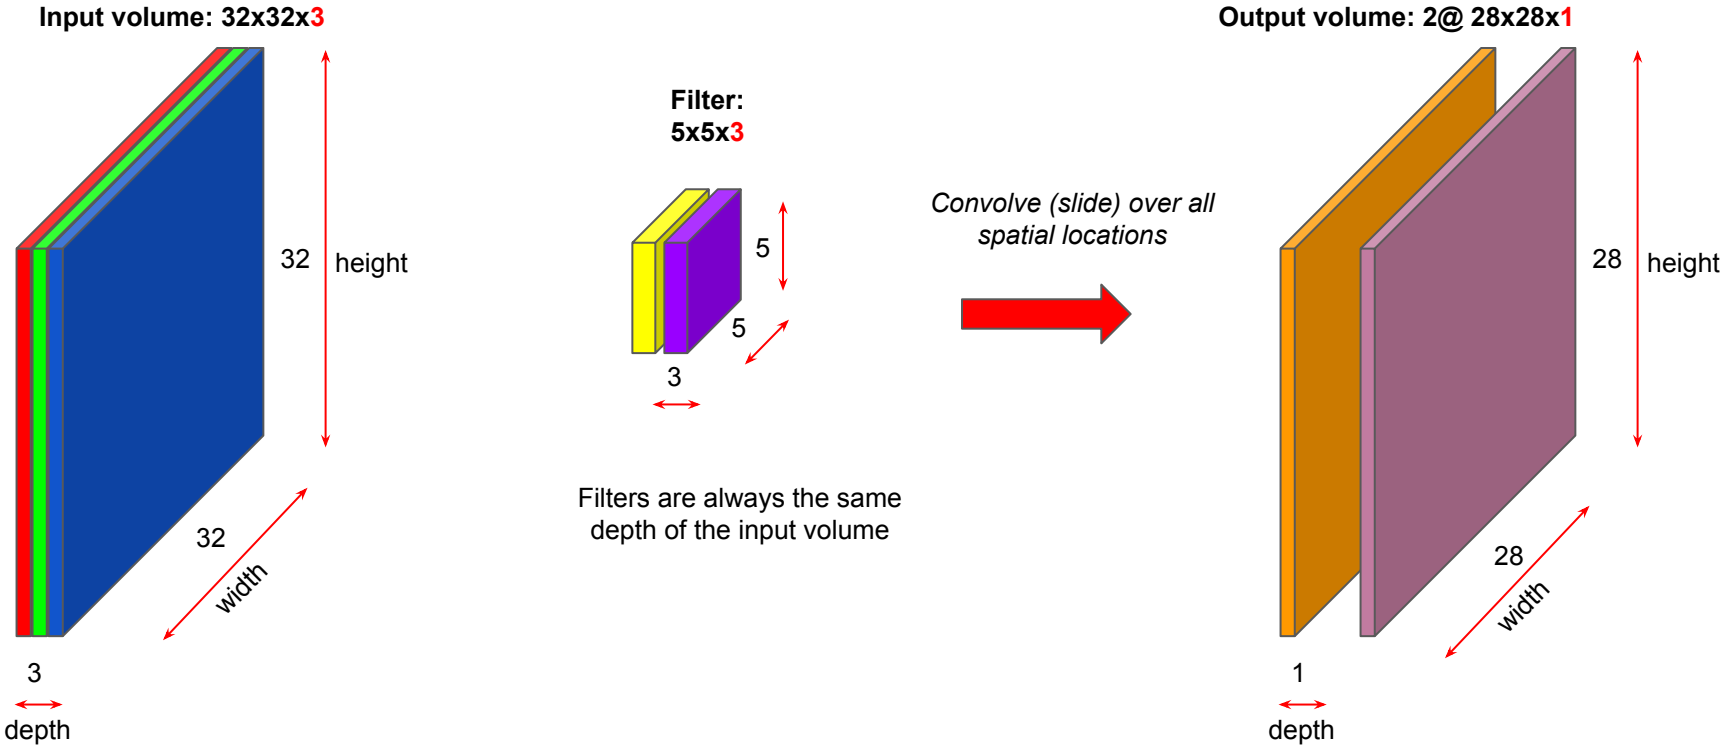

Convolutional Layers

If I add a third filter then I will get another feature map

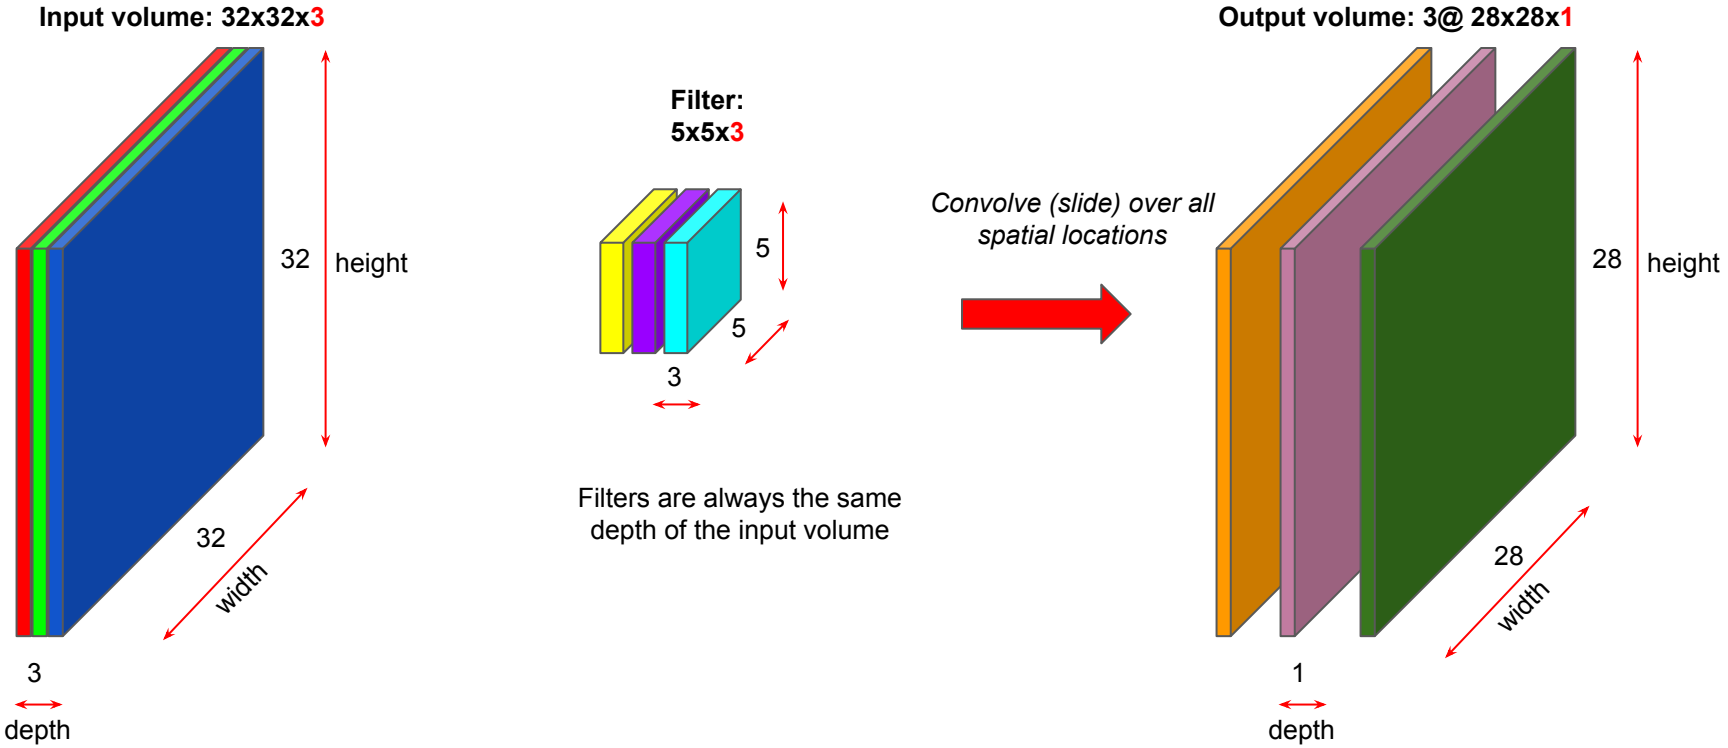

Convolutional Layers

They produce *feature maps* by convolving an input volume with a bank of filters.

Input volume: 32x32x3

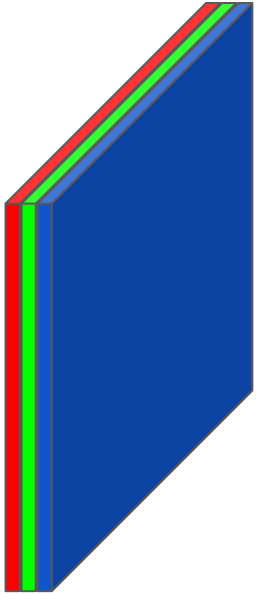

Filters bank:
5x5x3

Filters are always the same depth of the input volume

Convolve (slide) over all spatial locations

Output volume: 4@ 28x28x1

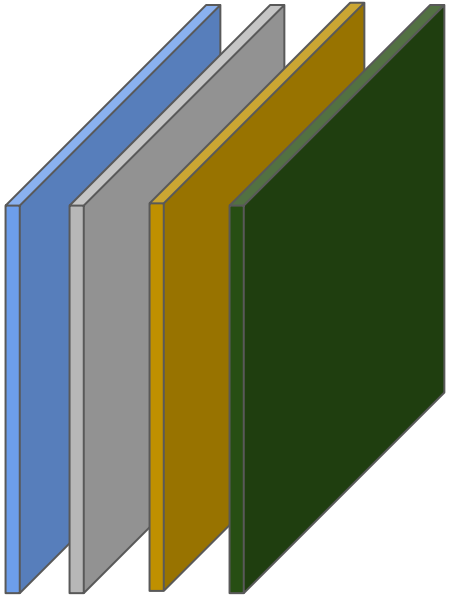

Feature maps

Convolutional Layers

How it works.

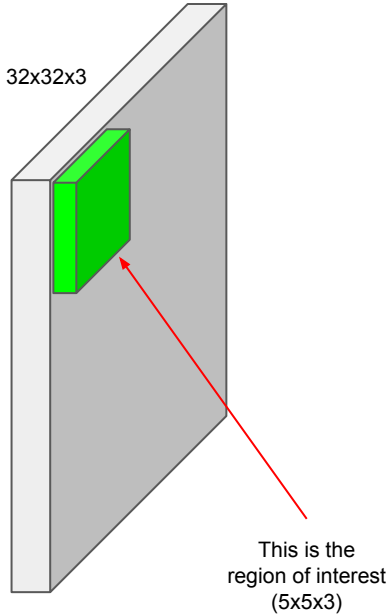

Convolutional Layers

How it works.

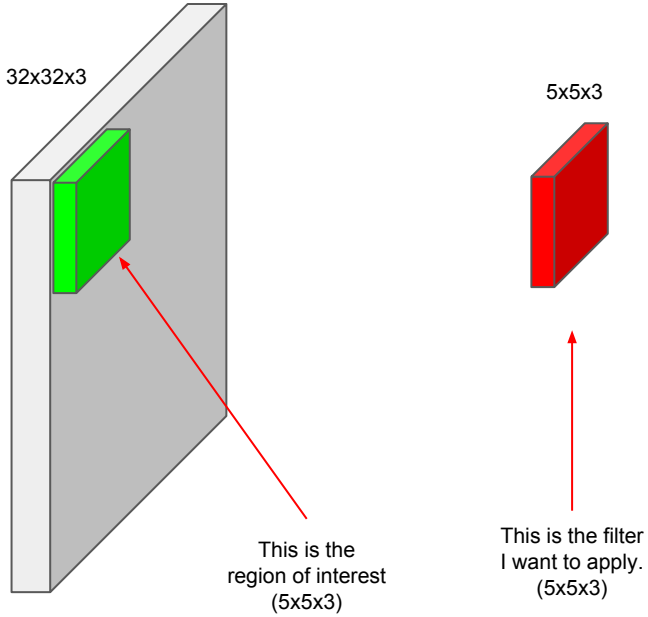

Convolutional Layers

How it works.

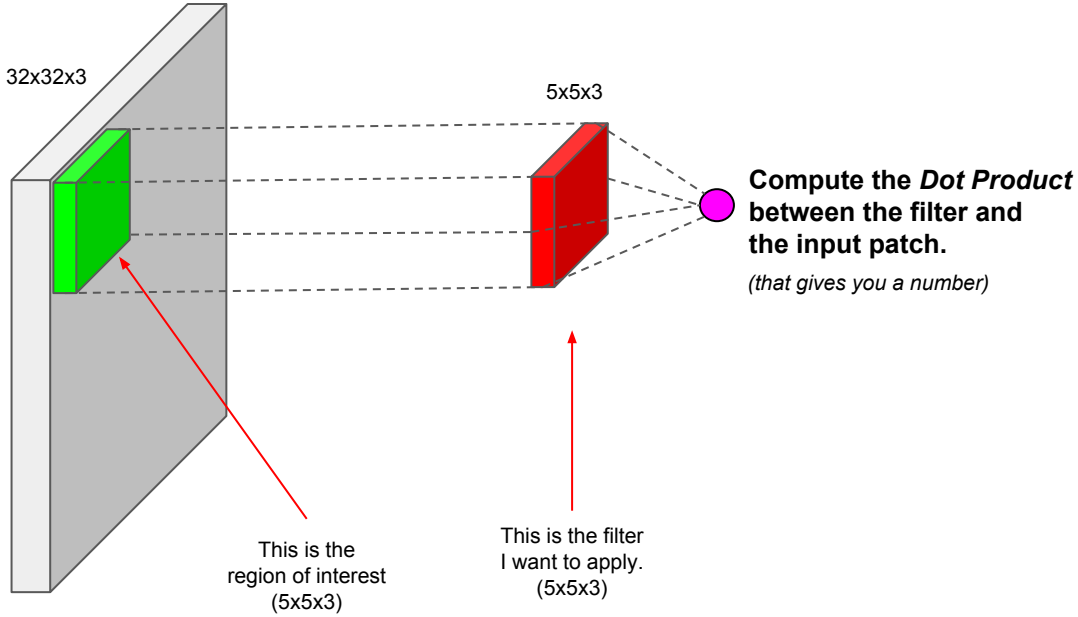

Convolutional Layers

How it works.

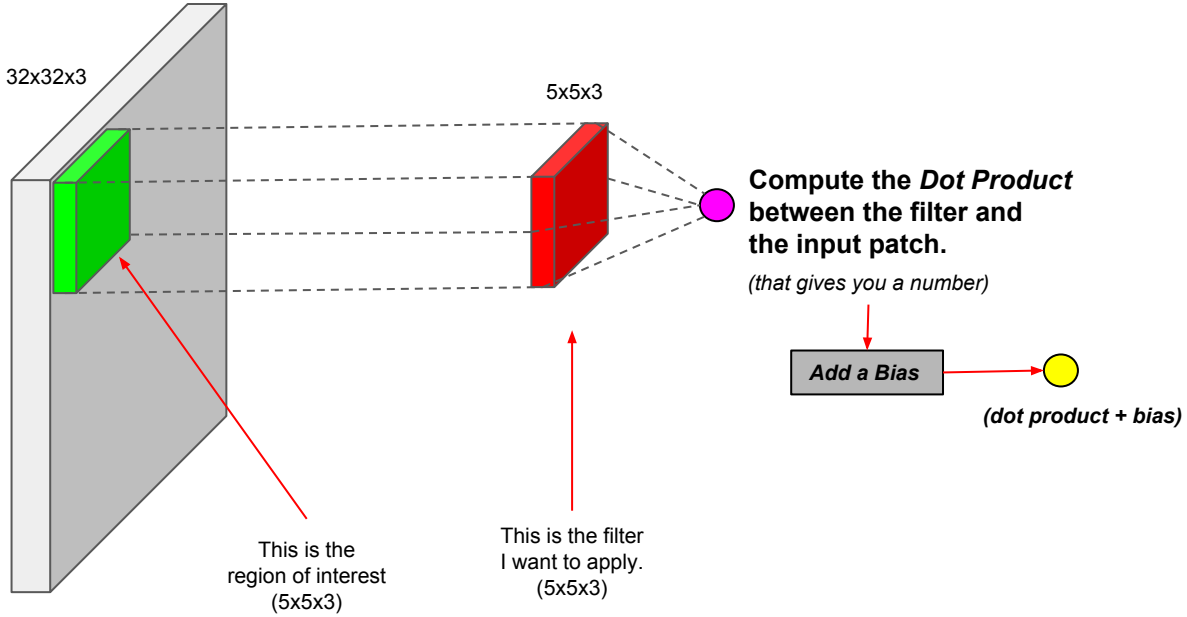

Convolutional Layers

How it works.

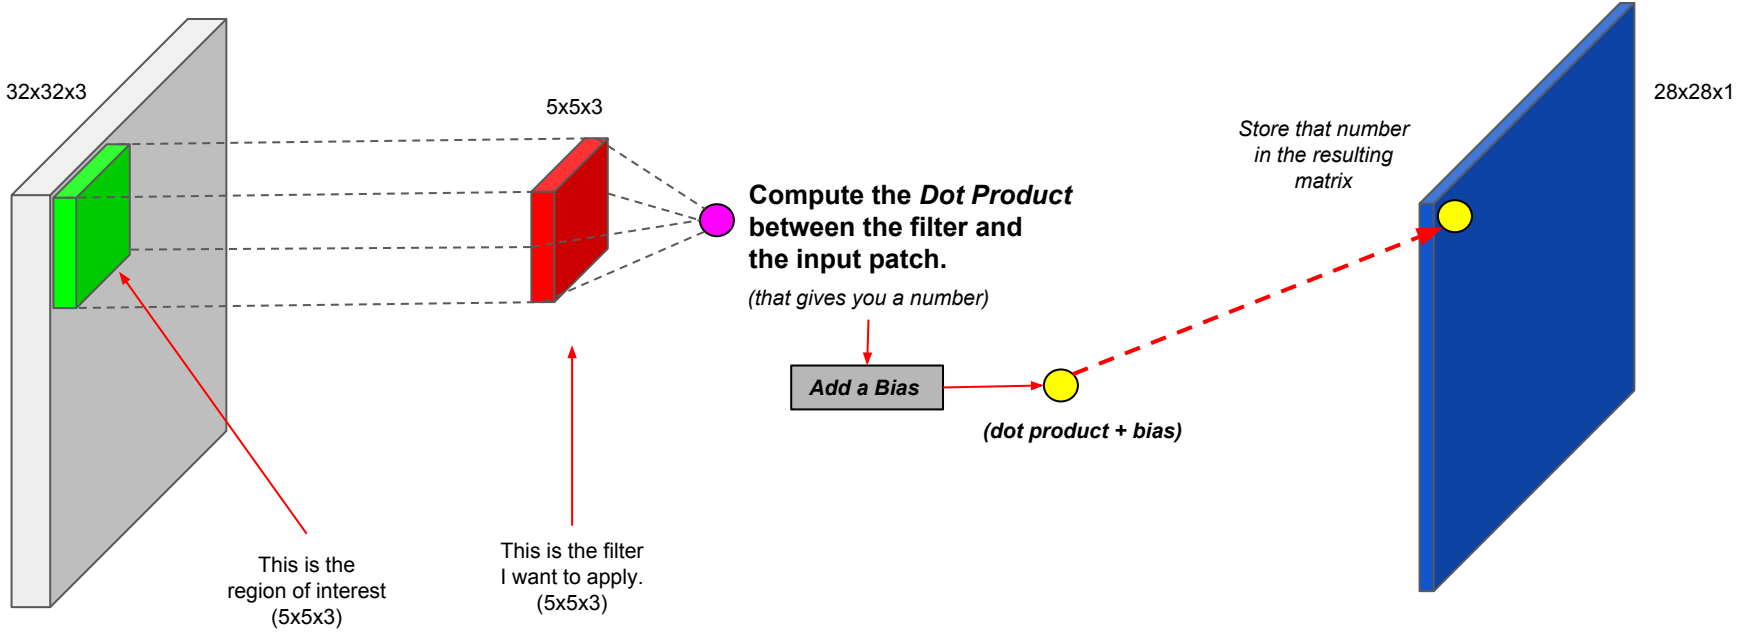

Convolutional Layers

How it works.

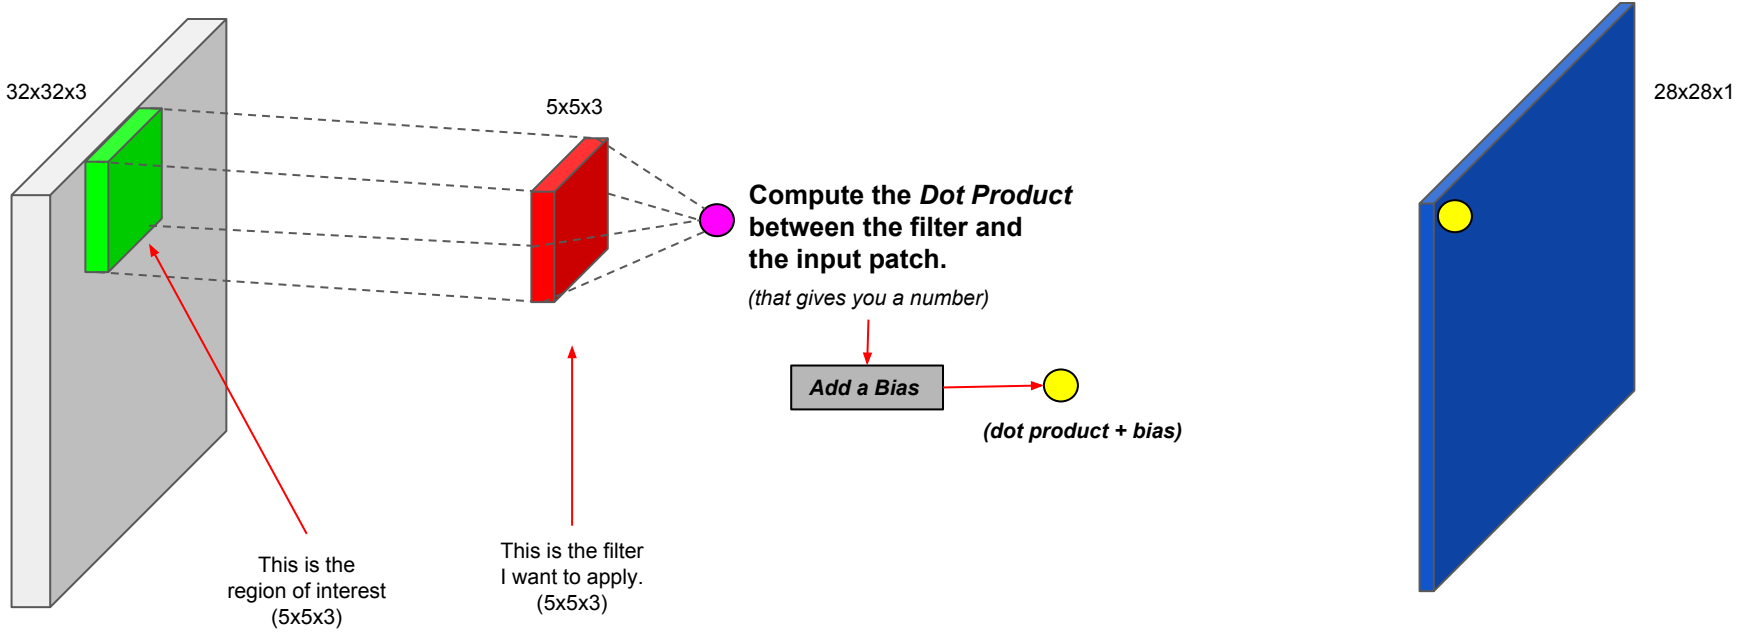

Convolutional Layers

How it works.

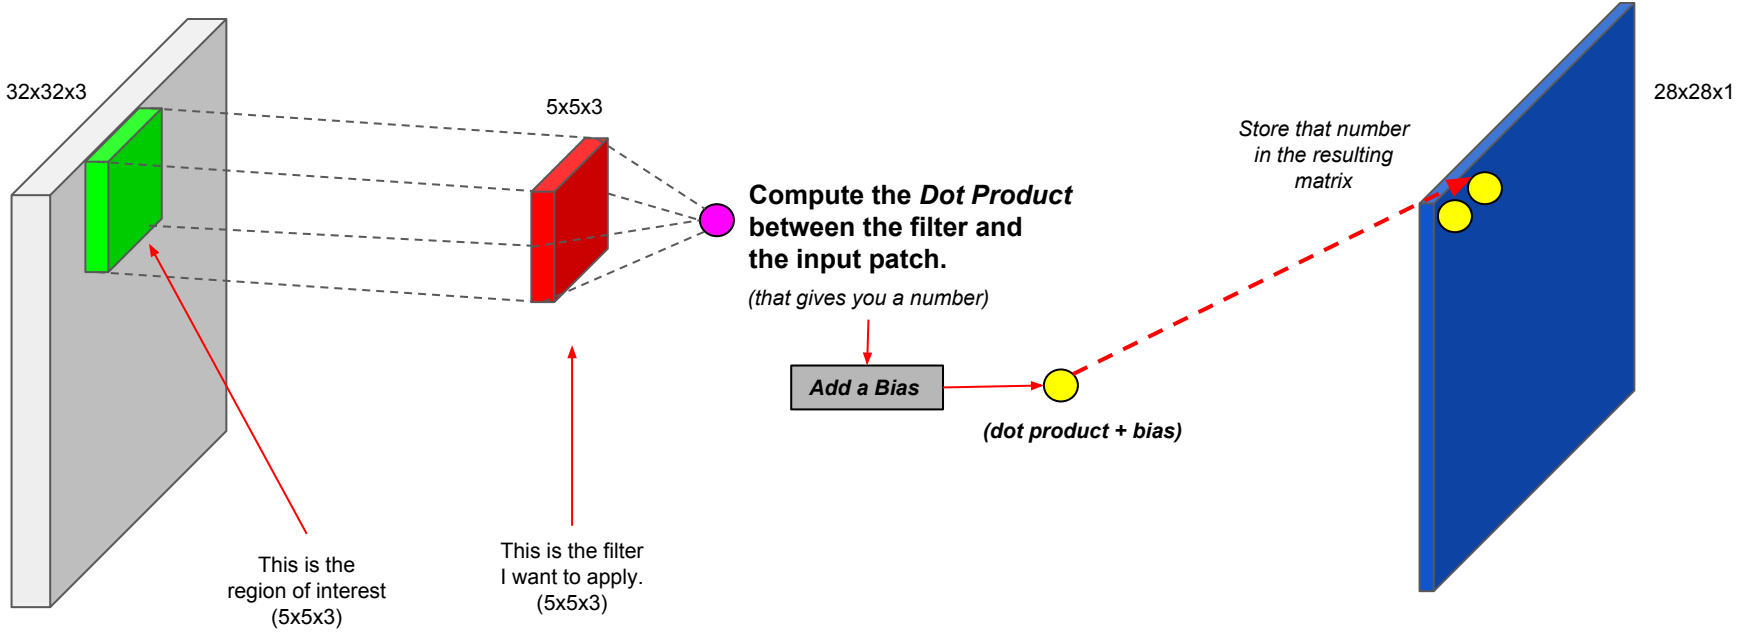

Convolutional Layers

How it works.

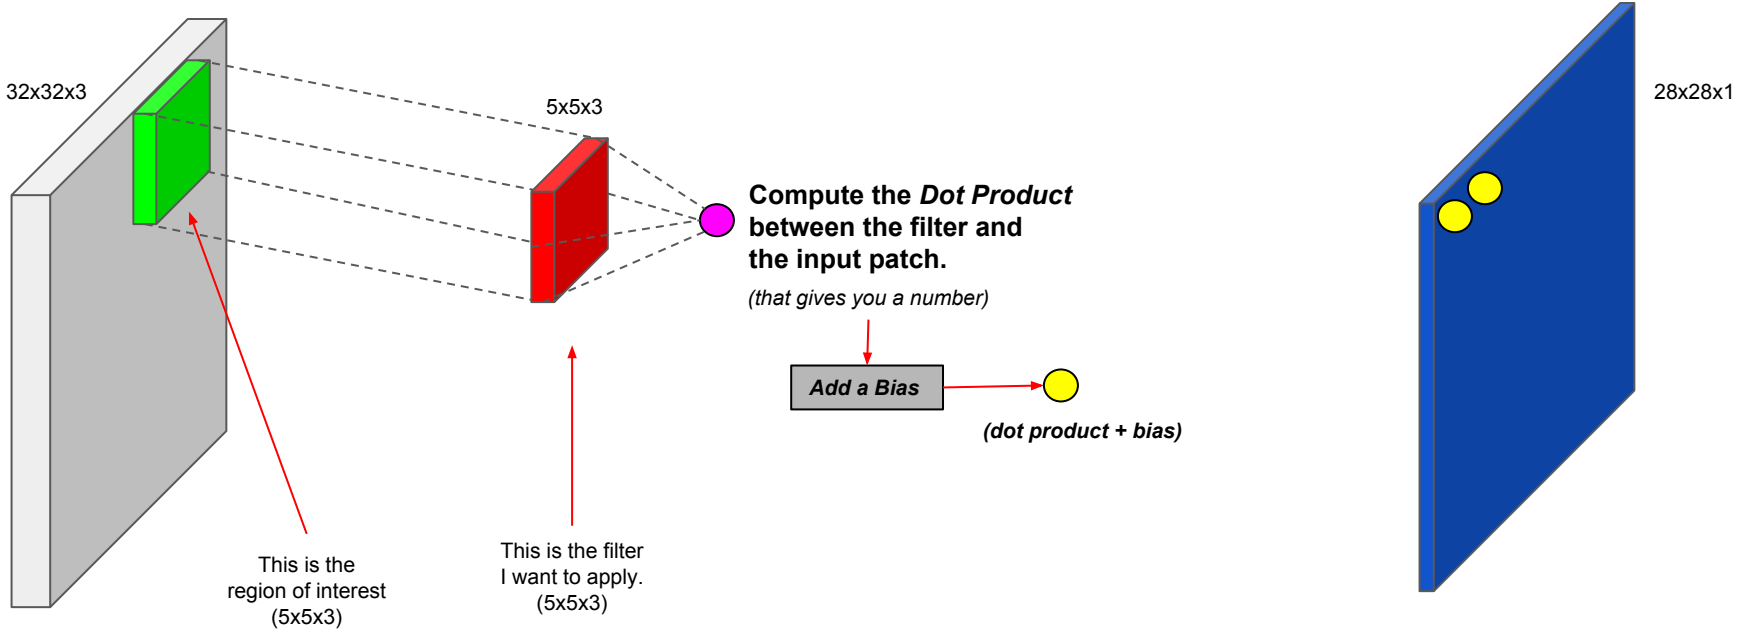

Convolutional Layers

How it works.

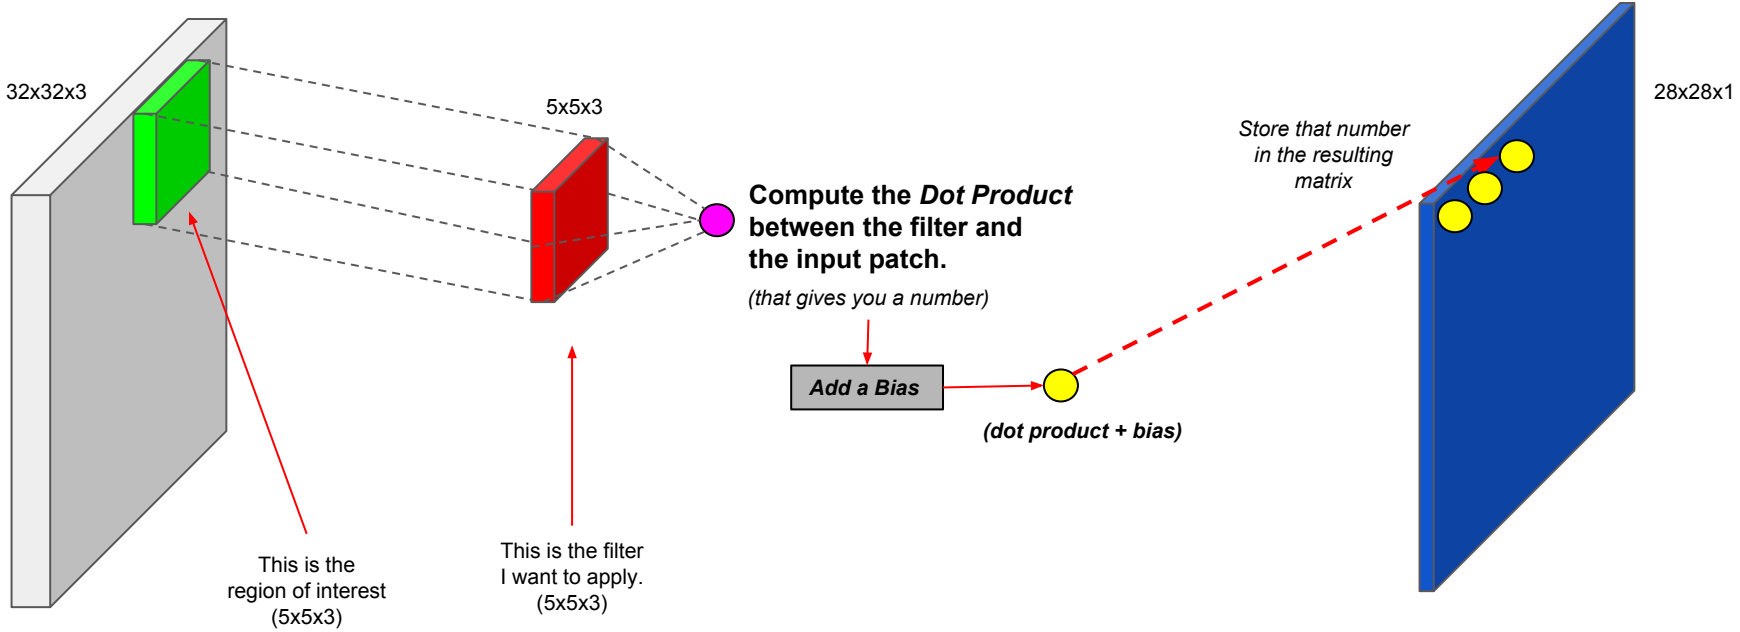

Convolutional Layers

How it works.

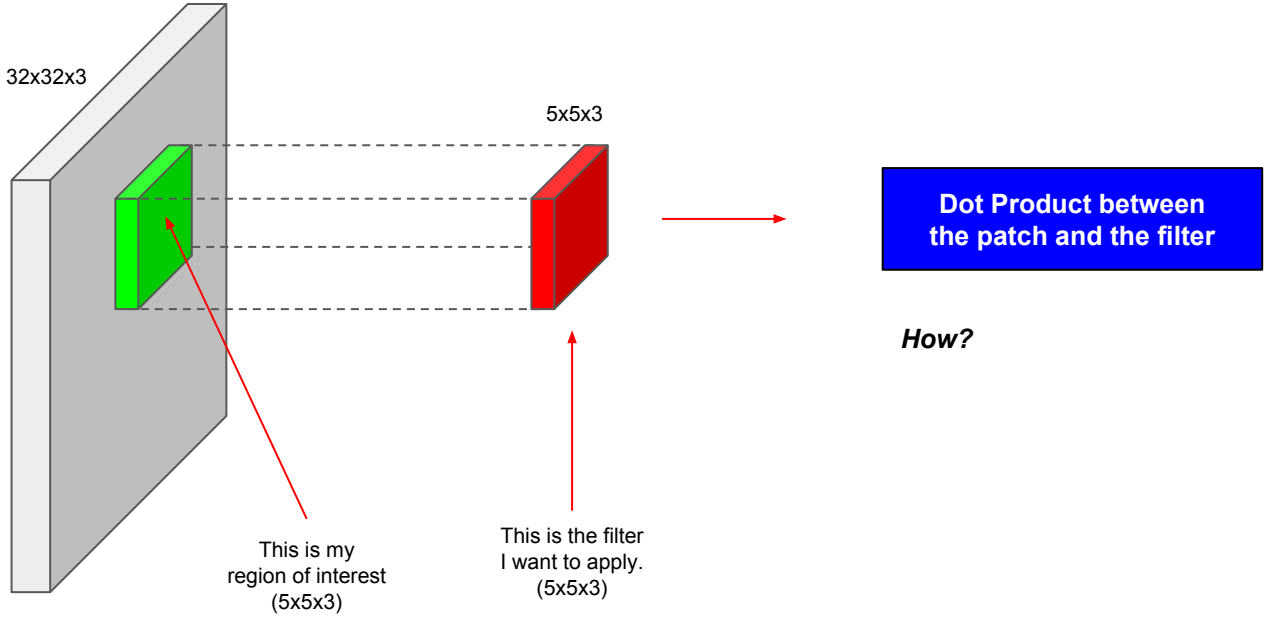

Convolutional Layers

How it works.

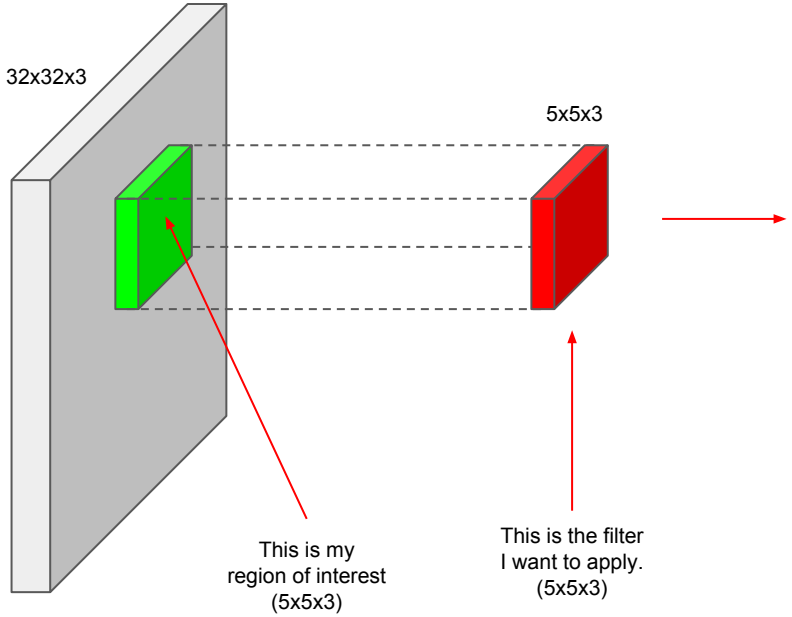

Step 1: First, I rotate the kernel 180 degrees.

Step 2: After that, I multiply both matrices.

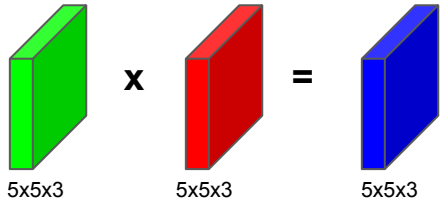

Rotate this matrix
180 degrees

Notice the resulting matrix
is also 5x5x3

Convolutional Layers

How it works.

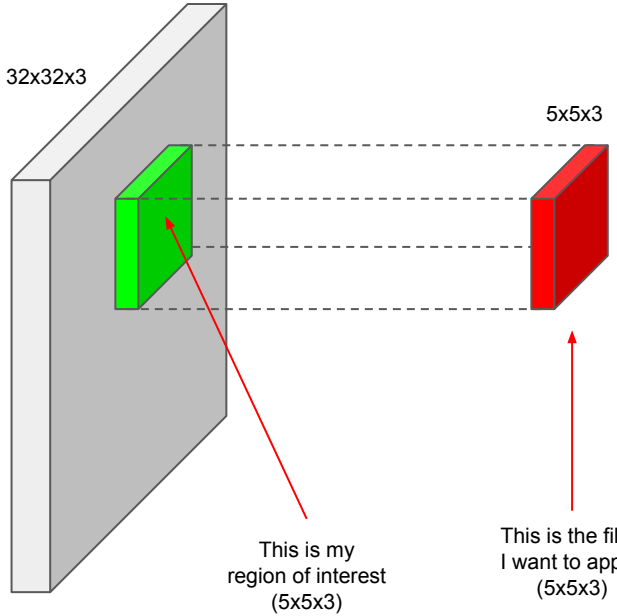

Then I sum up the values of the resulting matrix.

- *That gives me a number.*
- *That number is the dot product.*

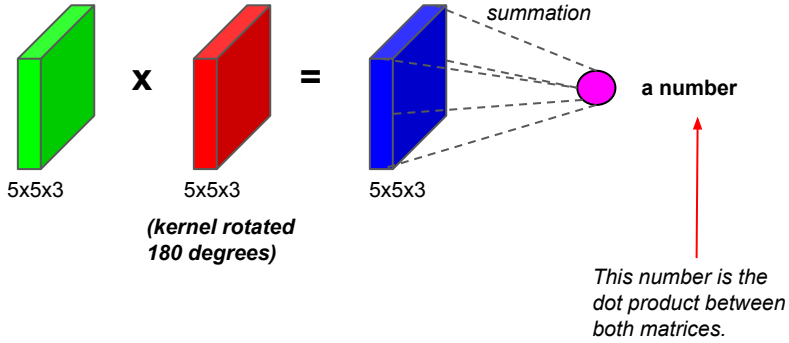

Convolutional Layers

Example → Compute the *cross-correlation* between two matrices of depth #1

Input patch

Matrix A		
1	2	3
4	5	6
7	8	9

Filter

Matrix B		
7	8	9
1	2	3
4	5	6

Matrix A multiplied by ***Matrix B***

Matrix C *(Matrix A multiplied by Matrix B)*		
7	16	27
4	10	18
28	40	54

Cross correlation			Total
7	16	27	
4	10	18	
28	40	54	
39	**66**	**99**	**204**

The result

7 + 16 + 27 + 4 + 10 + 18 + 28 + 40 + 54 = 204

Hands on →

Convolutional Layers

Another Example → Compute the cross-correlation between two matrices of size 3x3x3

A

Input Volume		
Red Channel		
255	50	67
80	44	21
33	77	95
Green Channel		
0	30	60
39	0	51
25	42	0
Blue Channel		
10	20	30
40	50	60
70	80	90

B

Filter		
Kernel for Red Channel		
1	2	3
4	5	6
7	8	9
Kernel for Green Channel		
9	8	7
6	5	4
3	2	1
Kernel for Blue Channel		
3	7	9
2	4	6
8	1	6

C = A x B

Multiplication		
Multiplication for Red Channel		
255	100	201
320	220	126
231	616	855
Multiplication for Green Channel		
0	240	420
234	0	204
75	84	0
Multiplication for Blue Channel		
30	140	270
80	200	360
560	80	540

D = SUM(C)

Cross-correlation	
Channel	Totals
Red	2924
Green	1257
Blue	2260
Result	**6441**

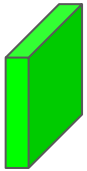

3x3x3

x

3x3x3

=

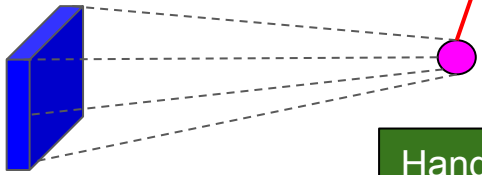

3x3x3

Hands on →

Convolutional Layers

In this case, the output response is the *cross-correlation* *(plus a bias)*

= **6441** =

Input Volume		
Red Channel		
255	50	67
80	44	21
33	77	95
Green Channel		
0	30	60
39	0	51
25	42	0
Blue Channel		
10	20	30
40	50	60
70	80	90

Filter		
Kernel for Red Channel		
1	2	3
4	5	6
7	8	9
Kernel for Green Channel		
9	8	7
6	5	4
3	2	1
Kernel for Blue Channel		
3	7	9
2	4	6
8	1	6

Multiplication		
Multiplication for Red Channel		
255	100	201
320	220	126
231	616	855
Multiplication for Green Channel		
0	240	420
234	0	204
75	84	0
Multiplication for Blue Channel		
30	140	270
80	200	360
560	80	540

Cross-correlation	
Channel	Totals
Red	2924
Green	1257
Blue	2260
Result	**6441**

The result is **6441**

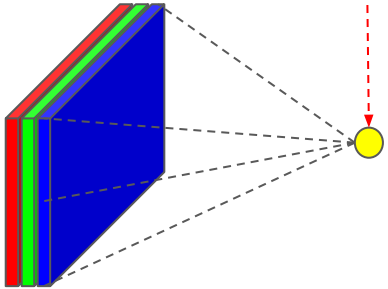

3x3x3

+ **bias**

(suppose the Bias is 1)

The value of this neuron is 6441 + 1
(cross correlation + bias)

Computing the *Dot Product*

A worked example, step by step

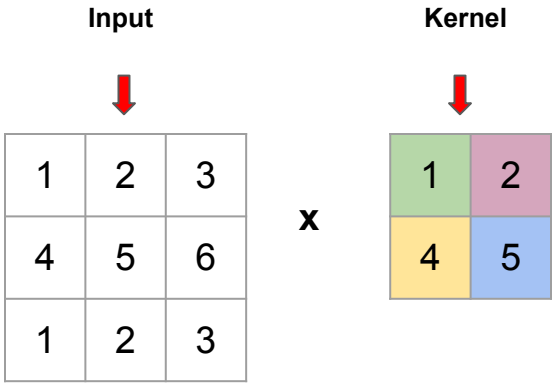

Computing the *Dot Product*

A worked example, step by step

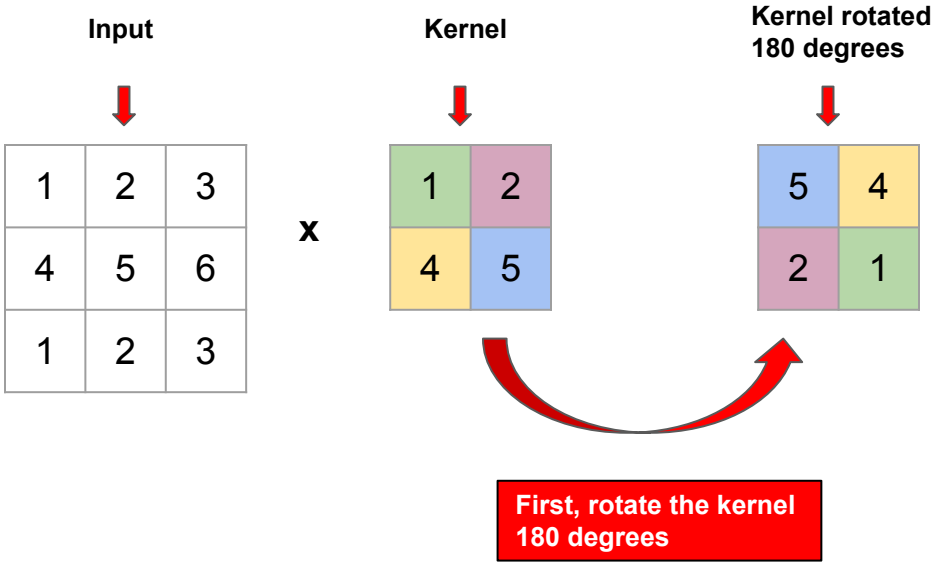

Computing the *Dot Product*

A worked example, step by step

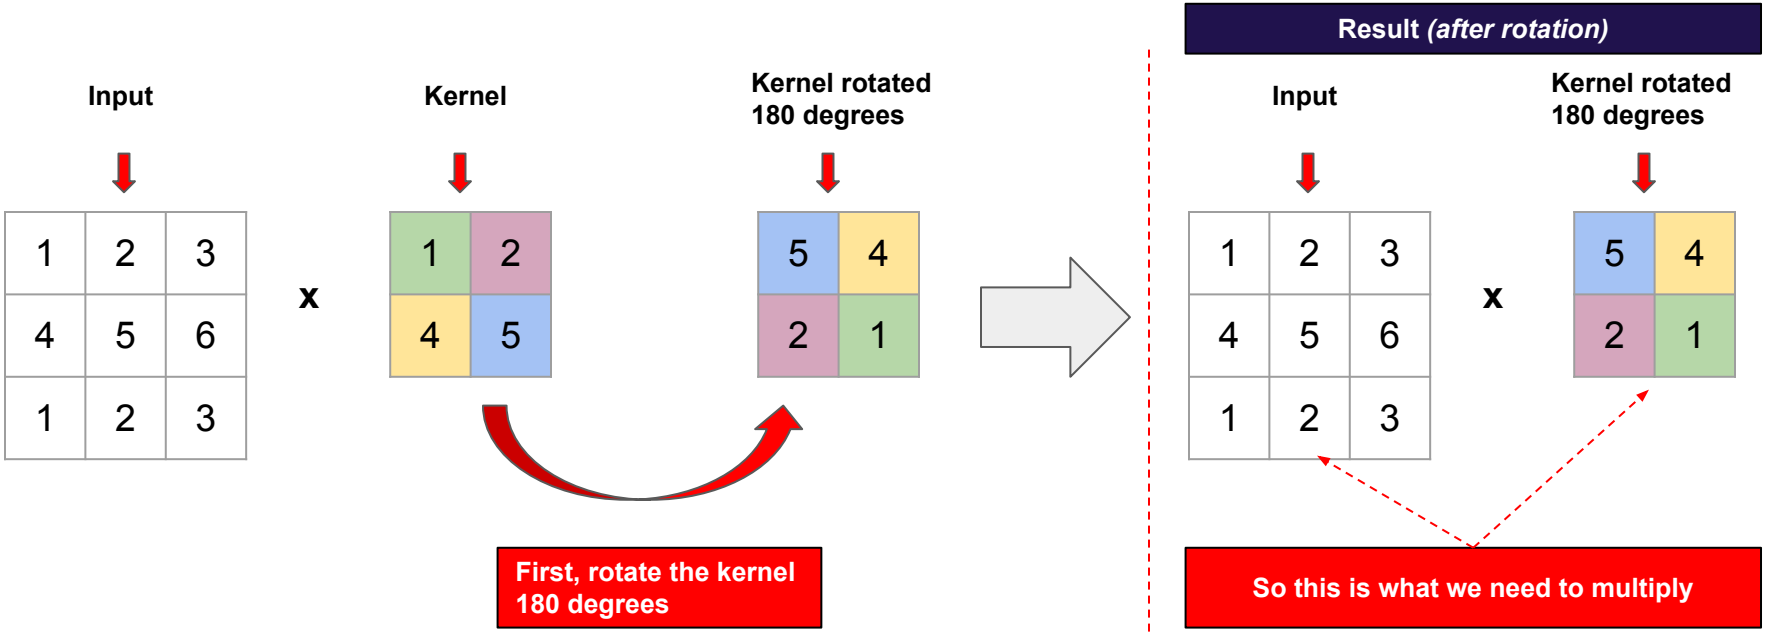

Computing the *Dot Product*

A worked example, step by step

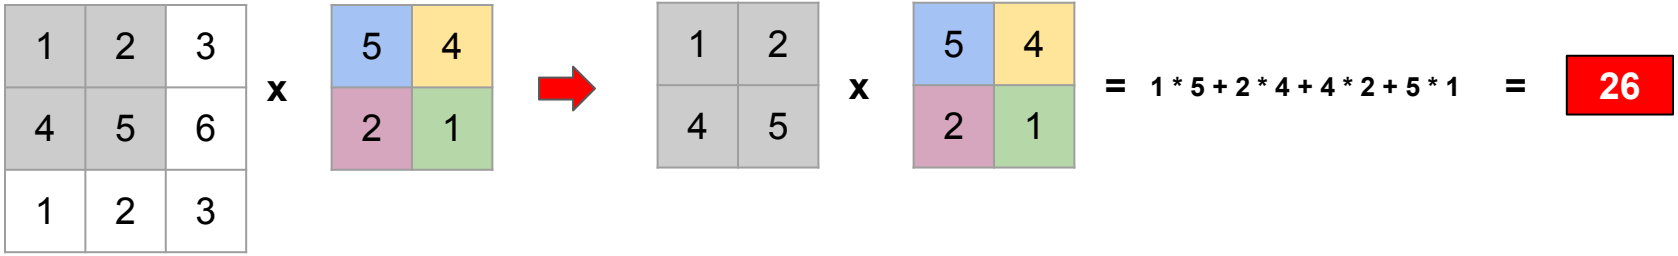

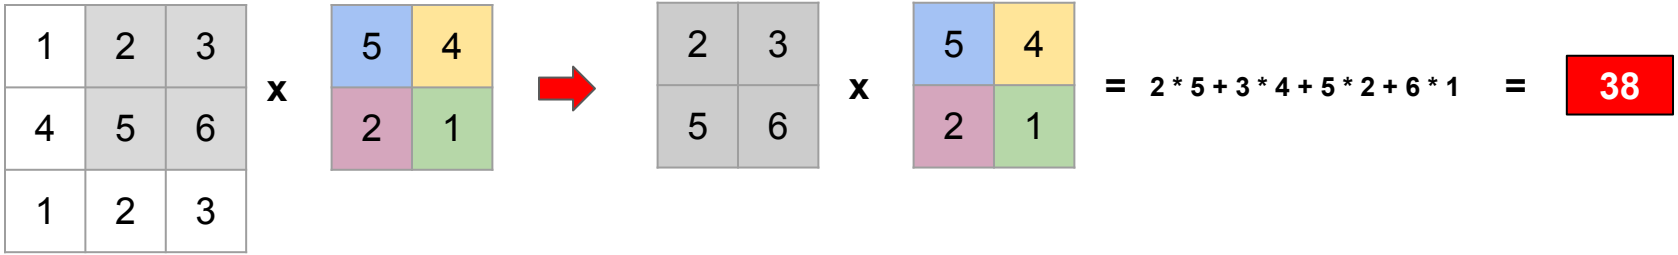

Computing the *Dot Product*

A worked example, step by step

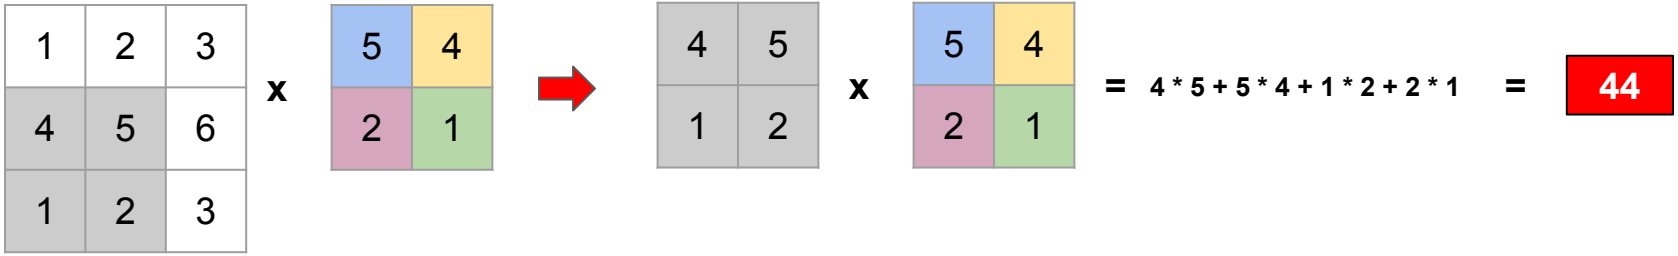

Computing the *Dot Product*

A worked example, step by step

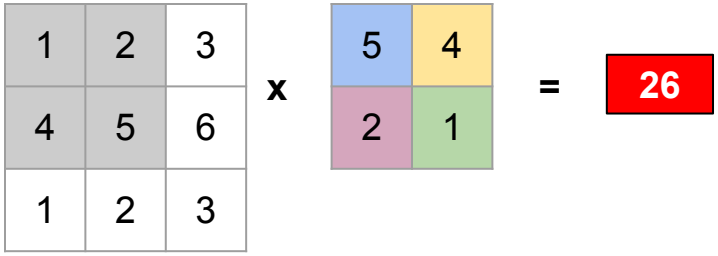

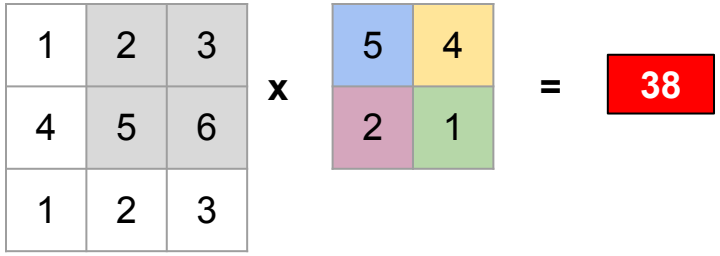

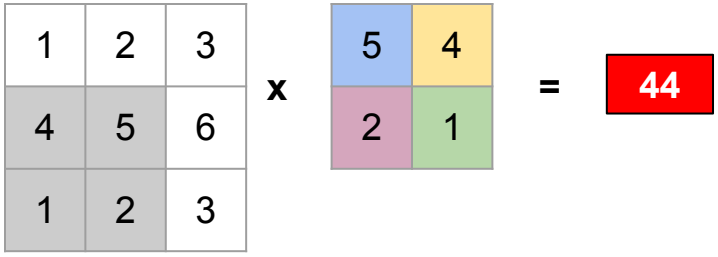

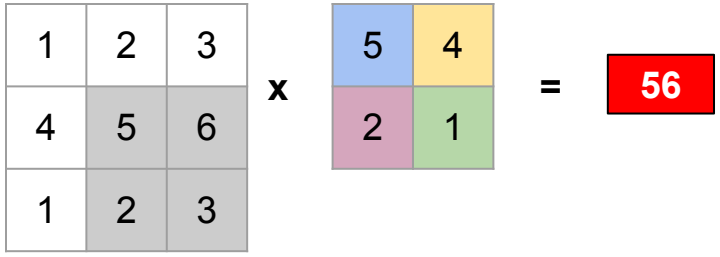

Computing the *Dot Product*

A worked example, step by step

1	2	3	x	5	4	=	26	1	2	3	x	5	4	=	38
4	5	6		2	1			4	5	6		2	1		
1	2	3						1	2	3					
1	2	3	x	5	4	=	44	1	2	3	x	5	4	=	56
4	5	6		2	1			4	5	6		2	1		
1	2	3						1	2	3					

1	2	3	x	5	4	=	26	38
4	5	6		2	1		44	56
1	2	3						

Convolutional Layers

Just to recap

I have an input patch
of depth 3
(it has 3 channels)

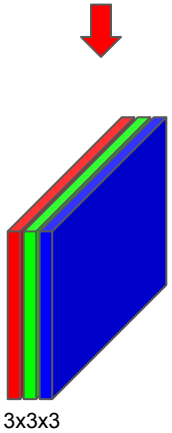

Convolutional Layers

Just to recap

I have an input patch of depth 3
(it has 3 channels)

I have a filter of the same size

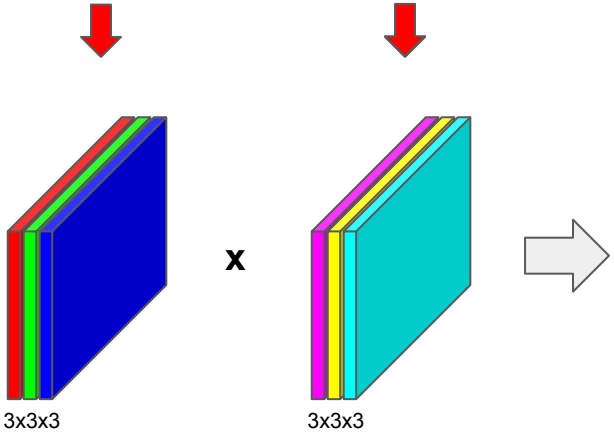

Convolutional Layers

Just to recap

I have an input patch of depth 3
(it has 3 channels)

I have a filter of the same size

Rotate the filter 180 degrees

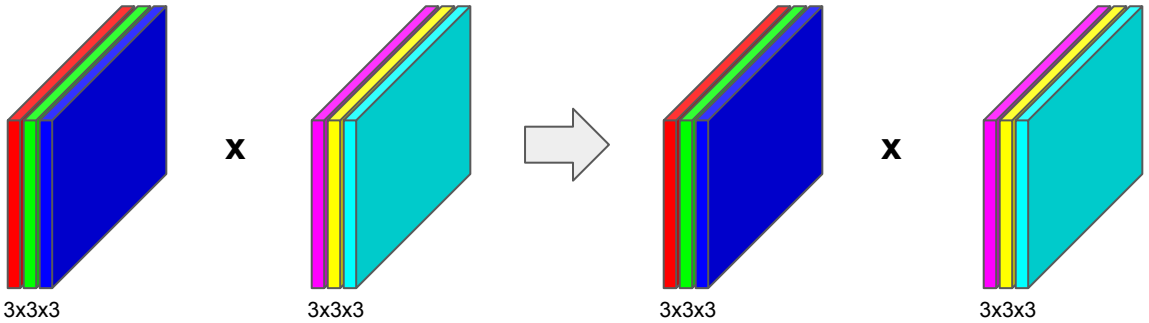

Convolutional Layers

Just to recap

I have an input patch of depth 3
(it has 3 channels)

I have a filter of the same size

Channel #1

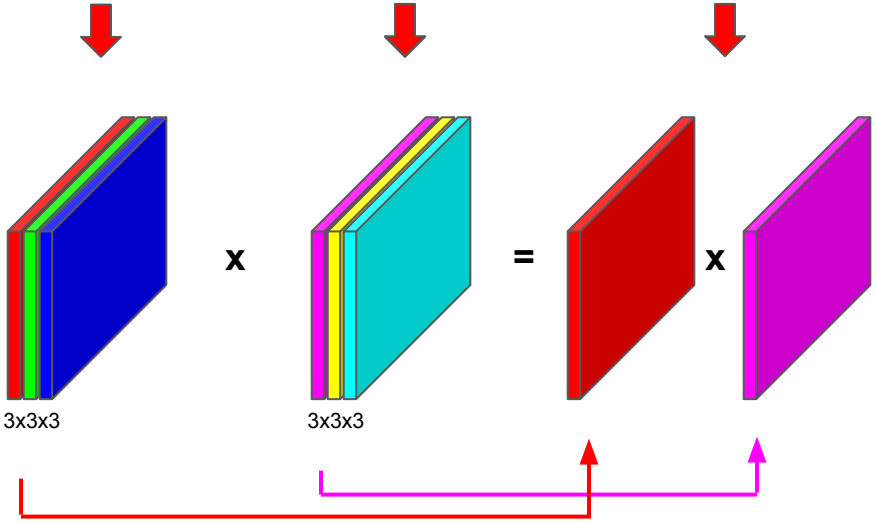

Convolutional Layers

Just to recap

I have an input patch of depth 3
(it has 3 channels)

I have a filter of the same size

Channel #1

Channel #2

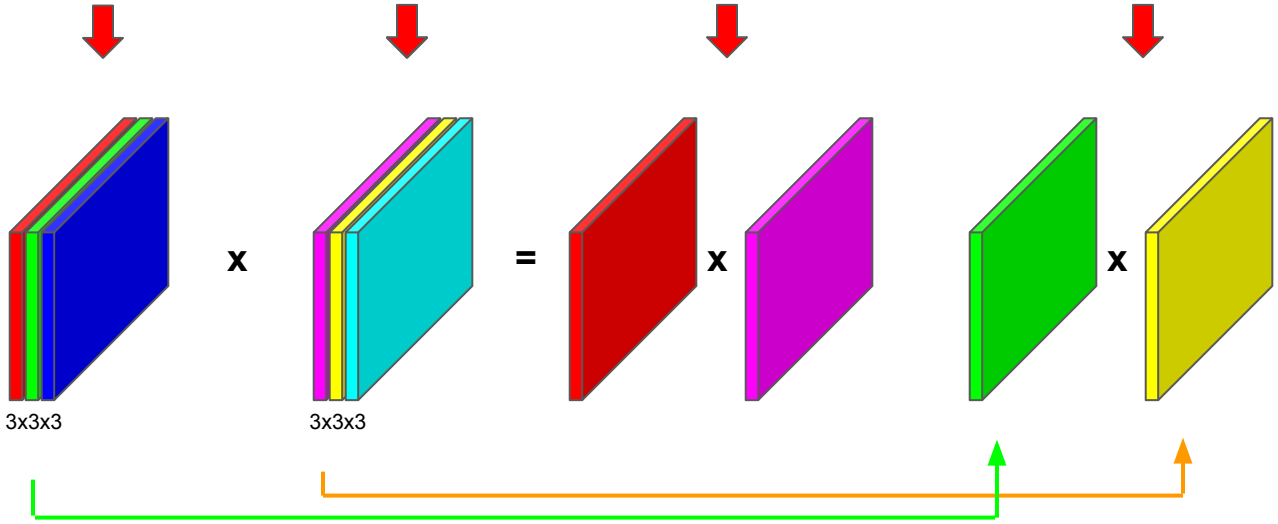

Convolutional Layers

Just to recap

I have an input patch of depth 3
(it has 3 channels)

I have a filter of the same size

Channel #1

Channel #2

Channel #3

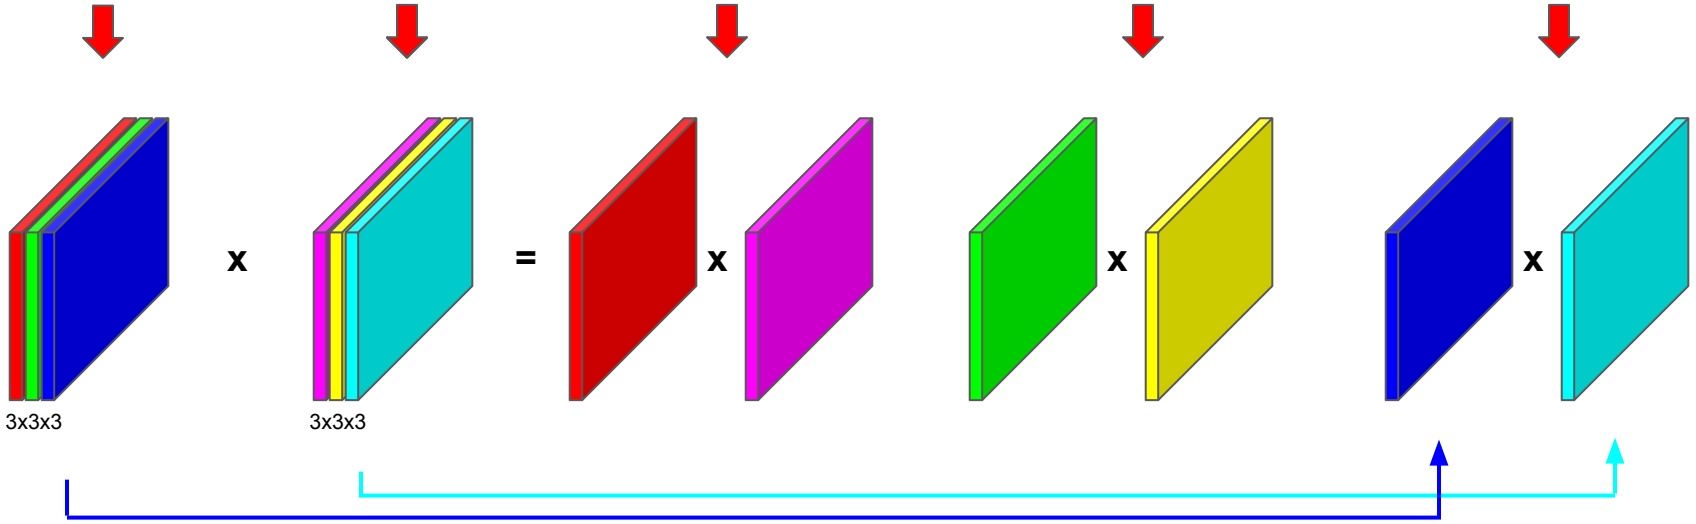

Convolutional Layers

Just to recap

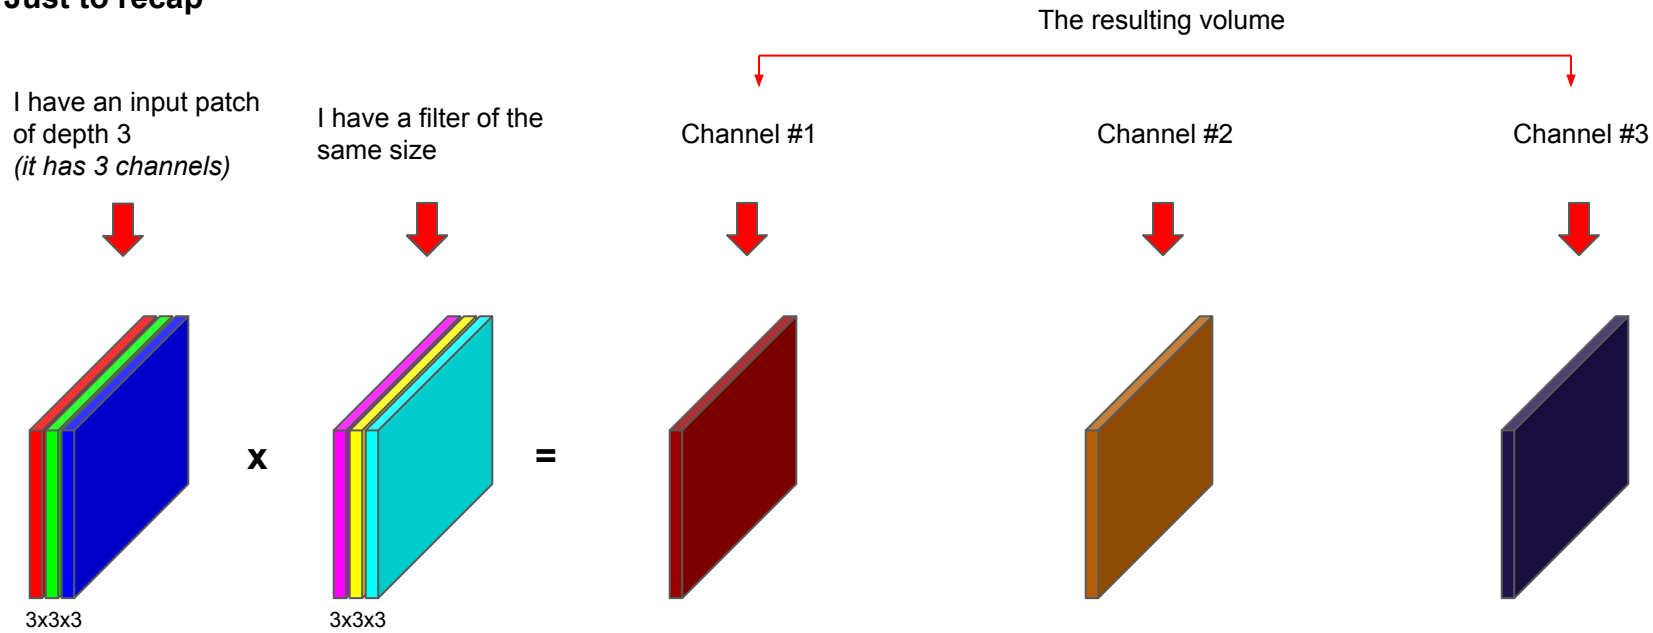

Convolutional Layers

Just to recap

I have an input patch of depth 3 *(it has 3 channels)*

I have a filter of the same size

The resulting volume C = A x B

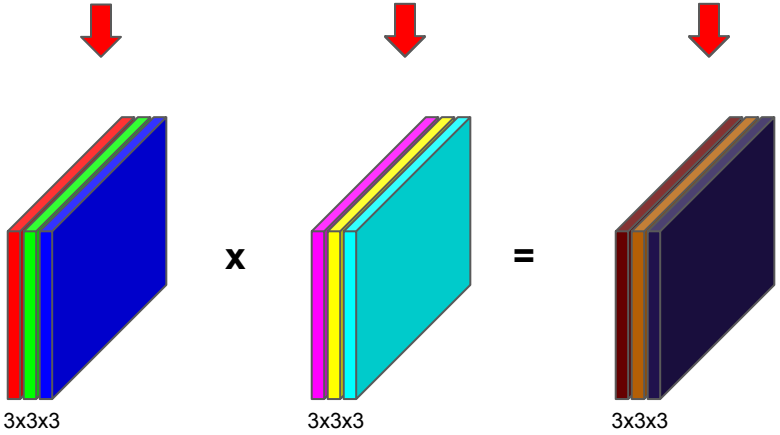

Convolutional Layers

Just to recap

The dot product between A and B

I have an input patch of depth 3 *(it has 3 channels)*

I have a filter of the same size

The resulting volume C = A x B

Summation of channels D = SUM(C)

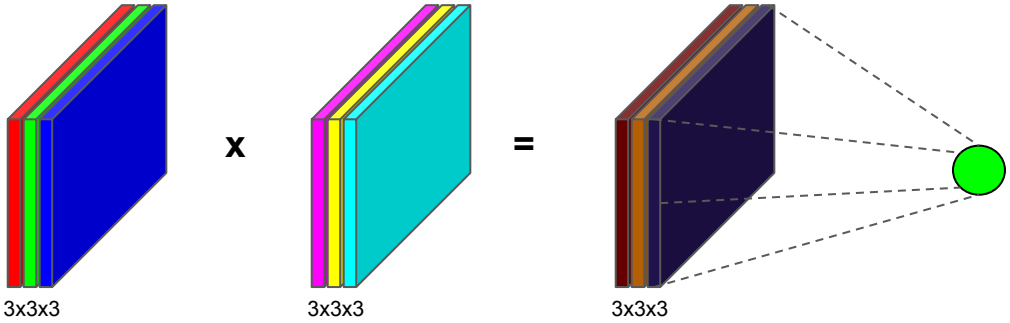

Convolutional Layers

Just to recap

The dot product between A and B

I have an input patch of depth 3 *(it has 3 channels)*

I have a filter of the same size

The resulting volume C = A x B

Summation of channels D = SUM(C)

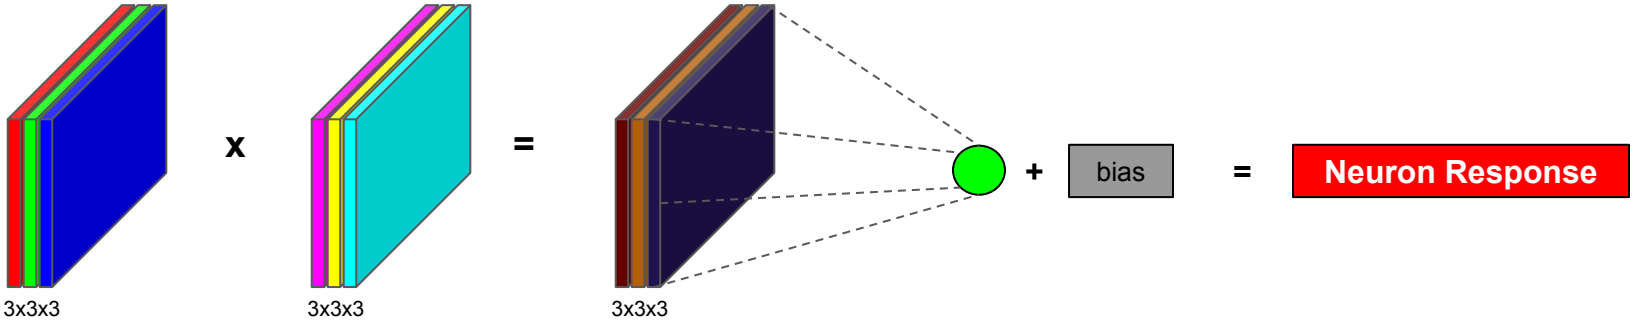

Convolutional Layers

They produce *feature maps* by convolving an input volume with a bank of filters.

Input volume: 32x32x3

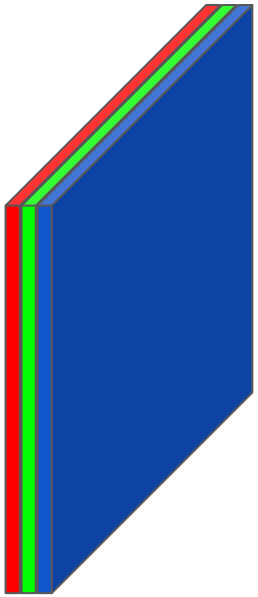

Filters bank:
4@ 5x5x3

Filters are always the same depth of the input volume

Convolve (slide) over all spatial locations

Output volume: 28x28x4

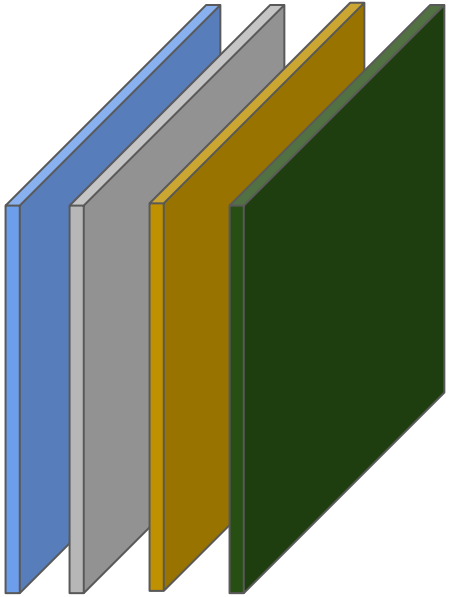

The output volume contains 1 feature map per filter

The goal of the network is to learn these filters

Convolutional Layers

Convolve (slide) over all spatial locations

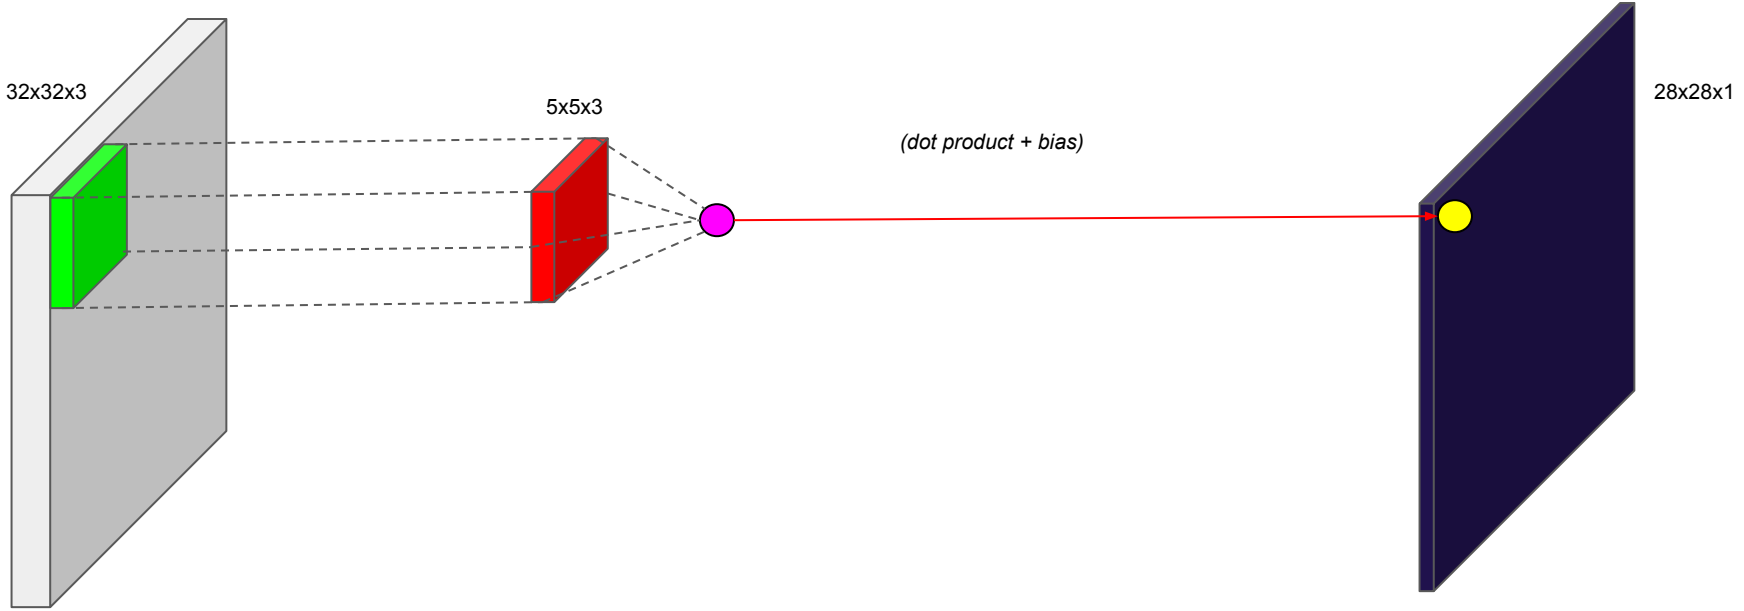

Convolutional Layers

Convolve (slide) over all spatial locations

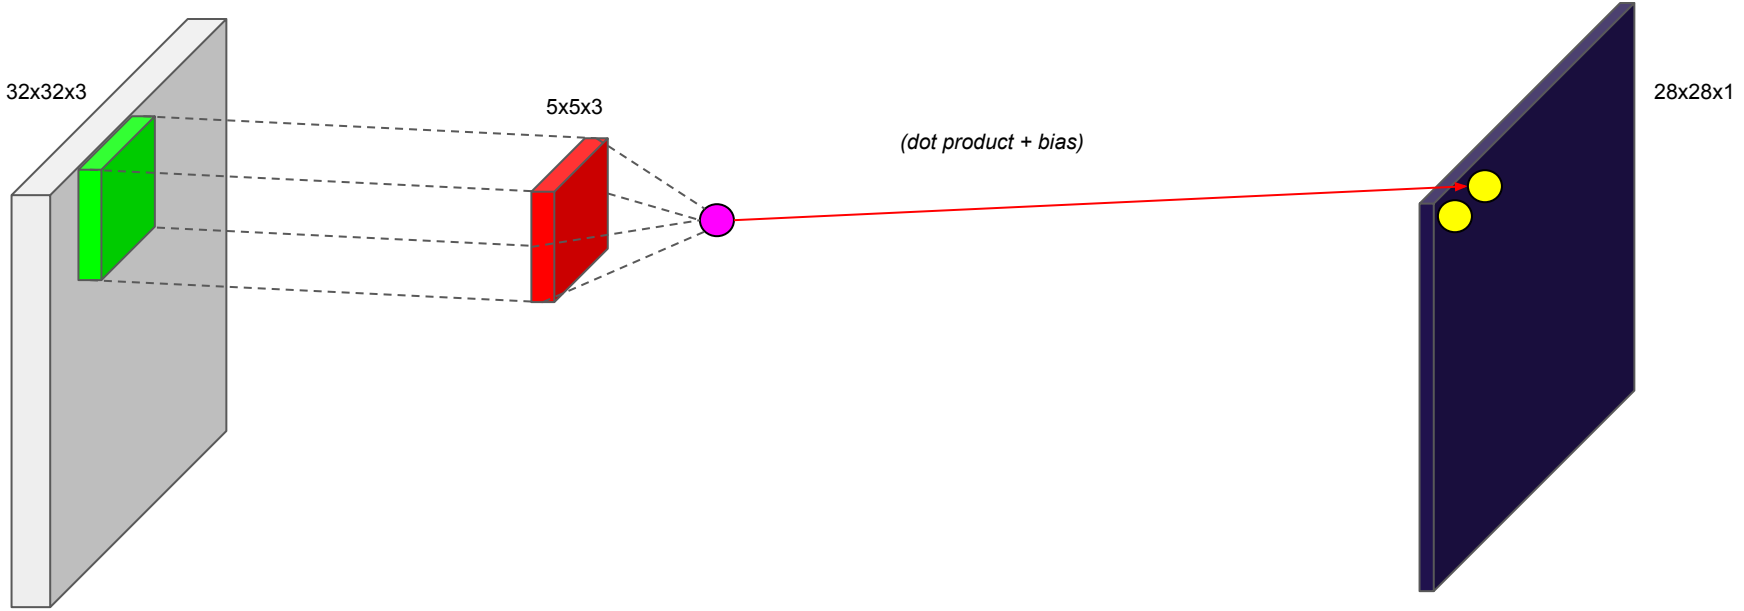

Convolutional Layers

Convolve (slide) over all spatial locations

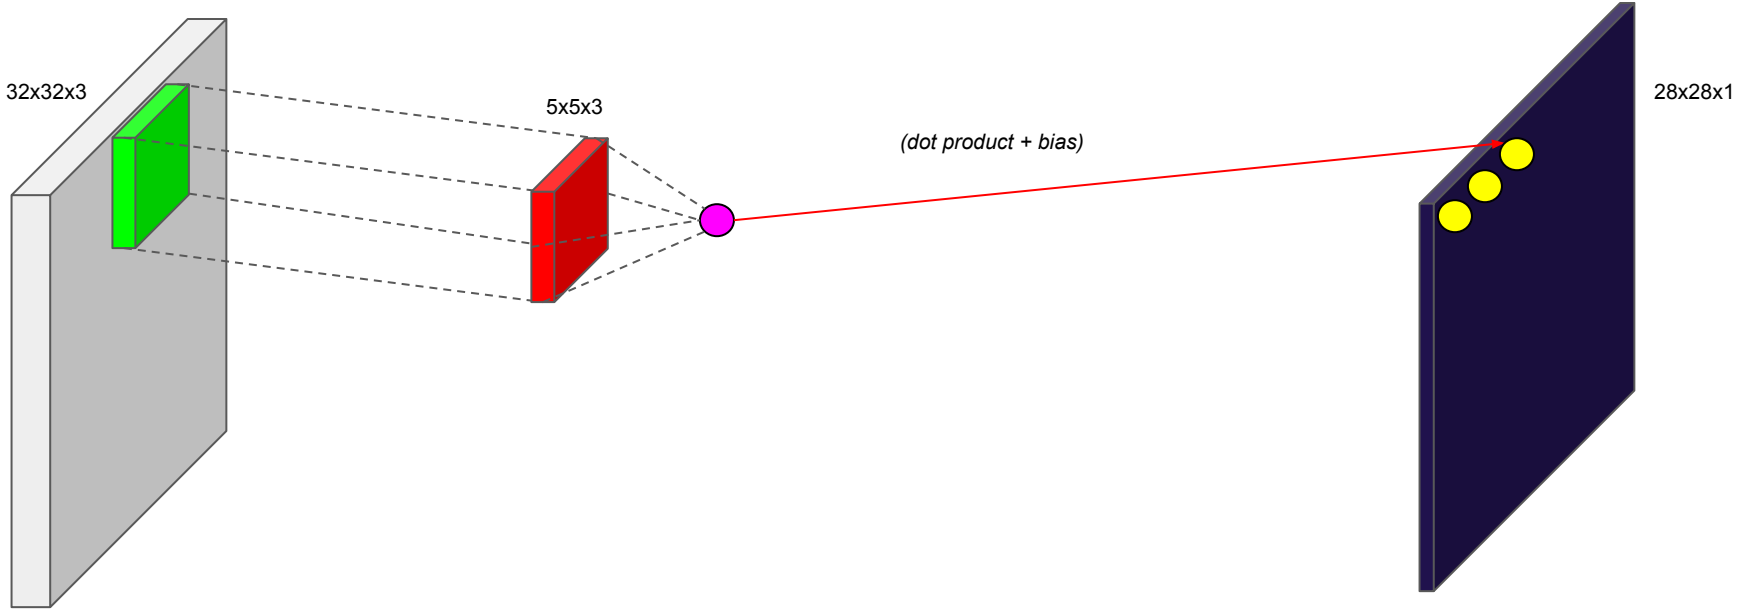

Convolutional Layers

Convolve (slide) over all spatial locations

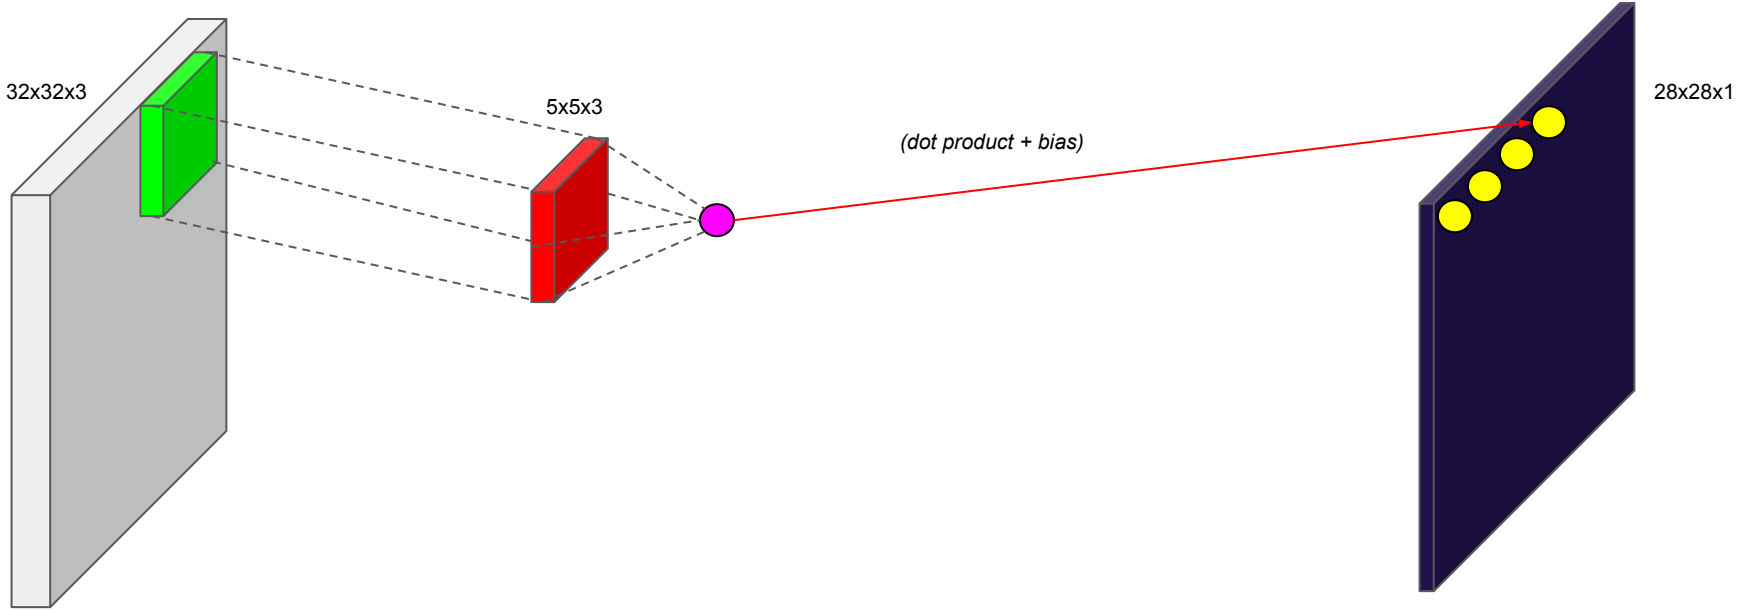

Convolutional Layers

Convolve (slide) over all spatial locations

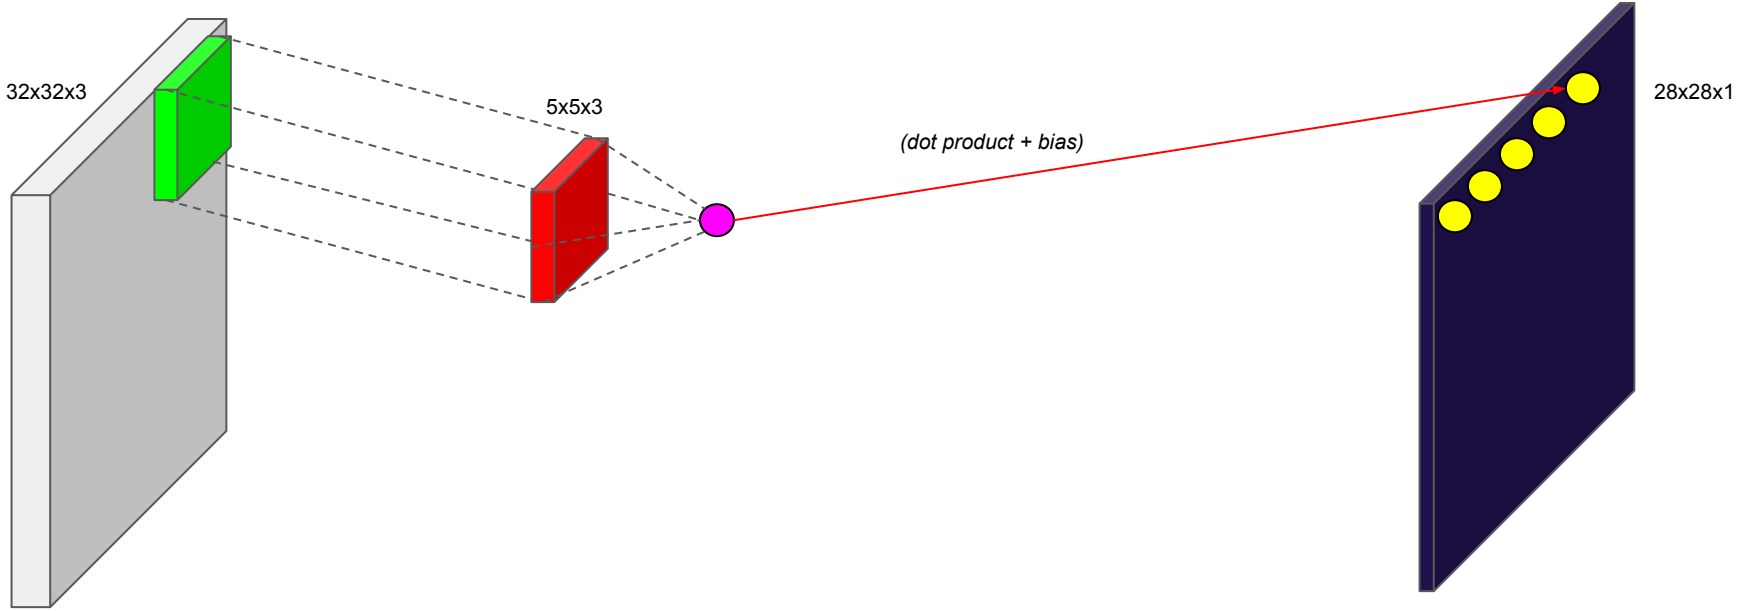

Convolutional Layers

Convolve (slide) over all spatial locations

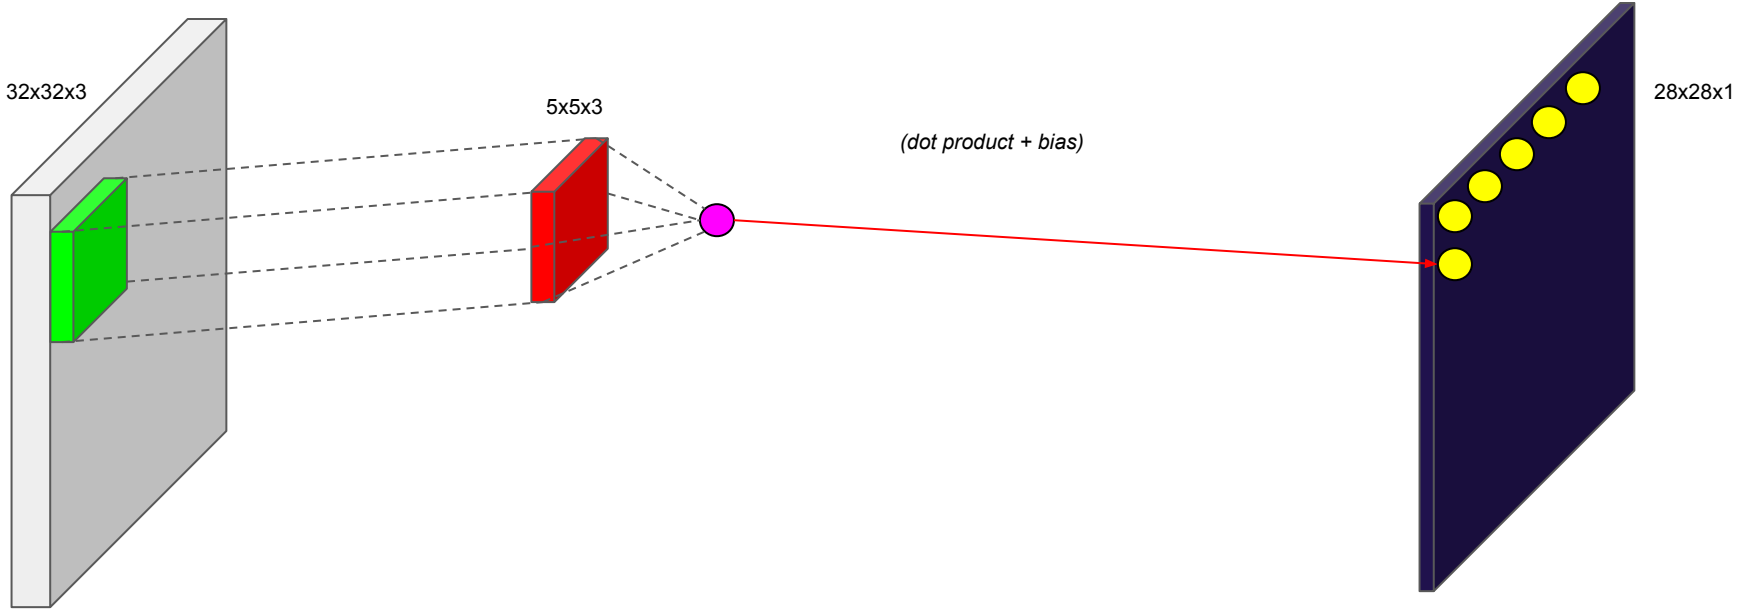

Convolutional Layers

Convolve (slide) over all spatial locations

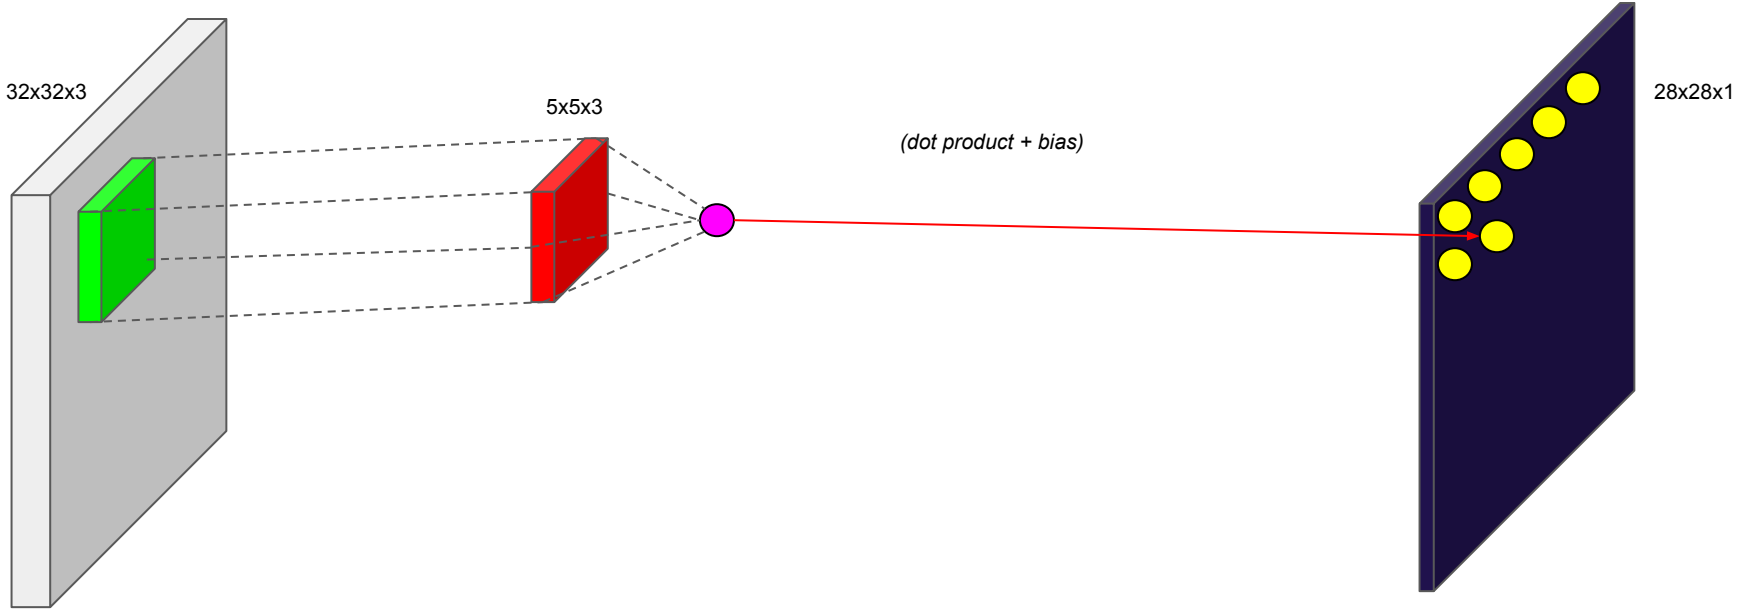

Convolutional Layers

Convolve (slide) over all spatial locations

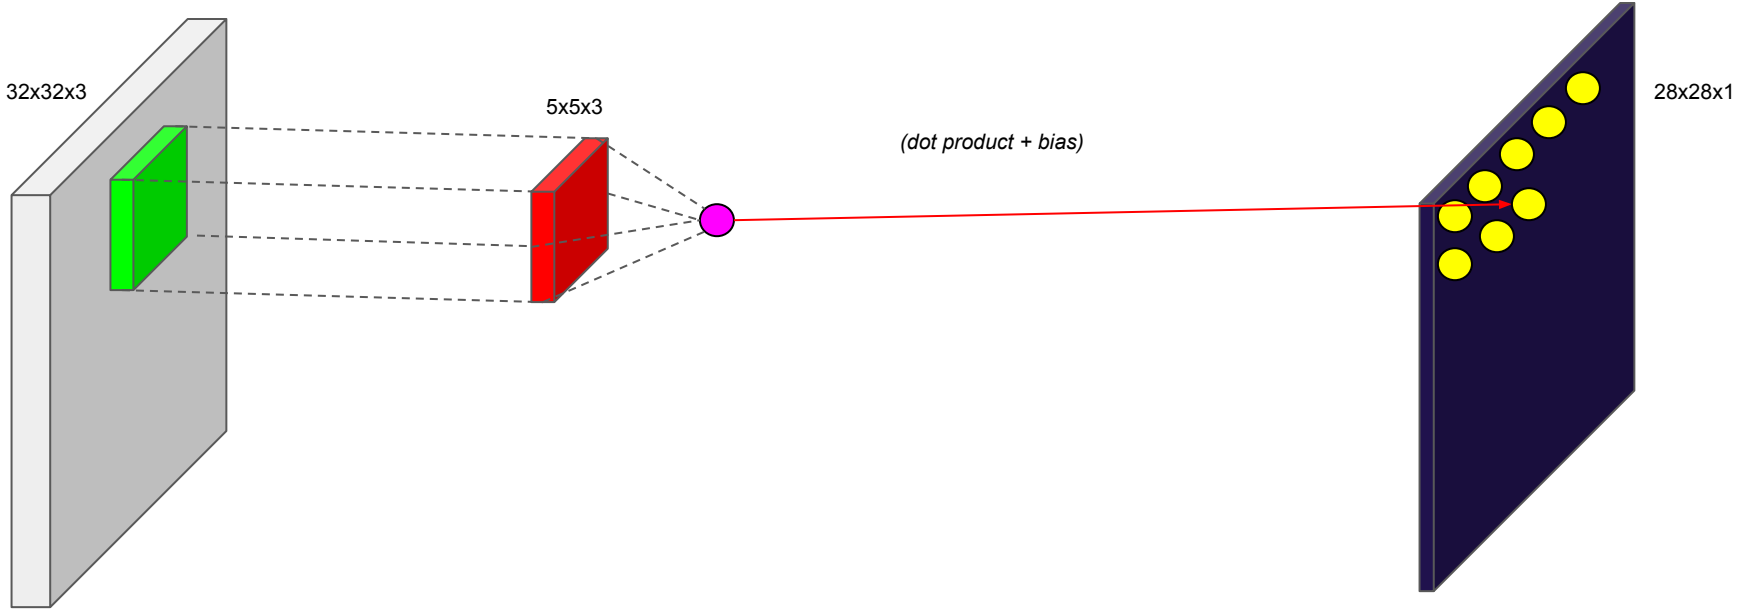

Convolutional Layers

Convolve (slide) over all spatial locations

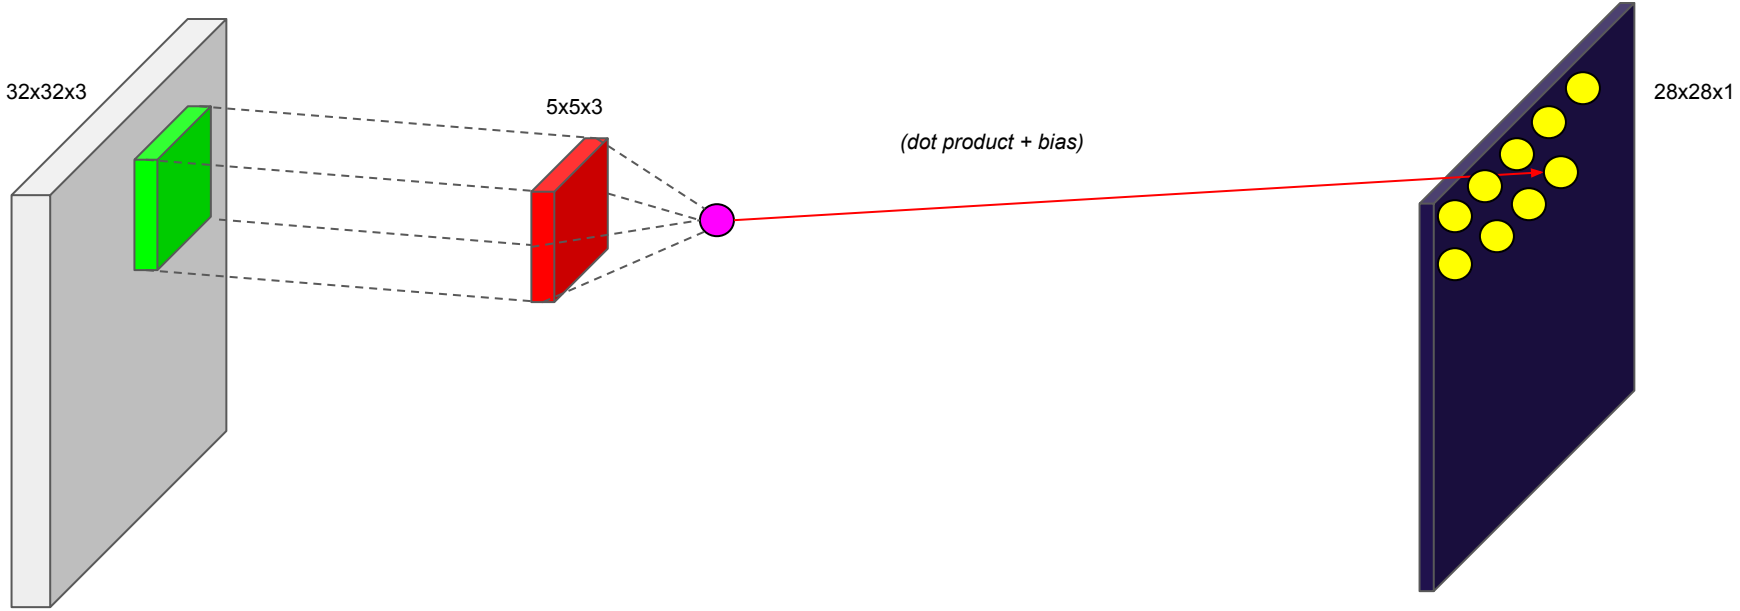

Convolutional Layers

Convolve (slide) over all spatial locations

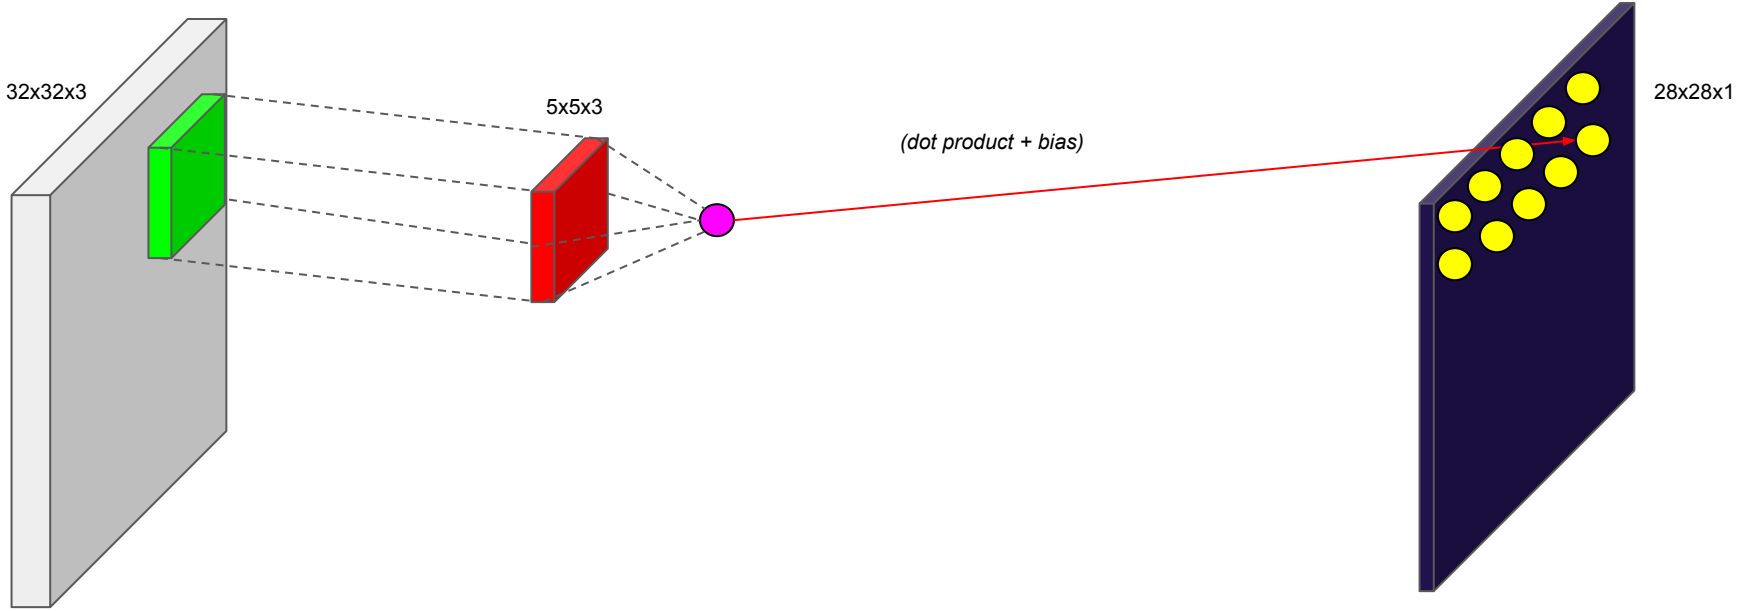

Convolutional Layers

Convolve (slide) over all spatial locations

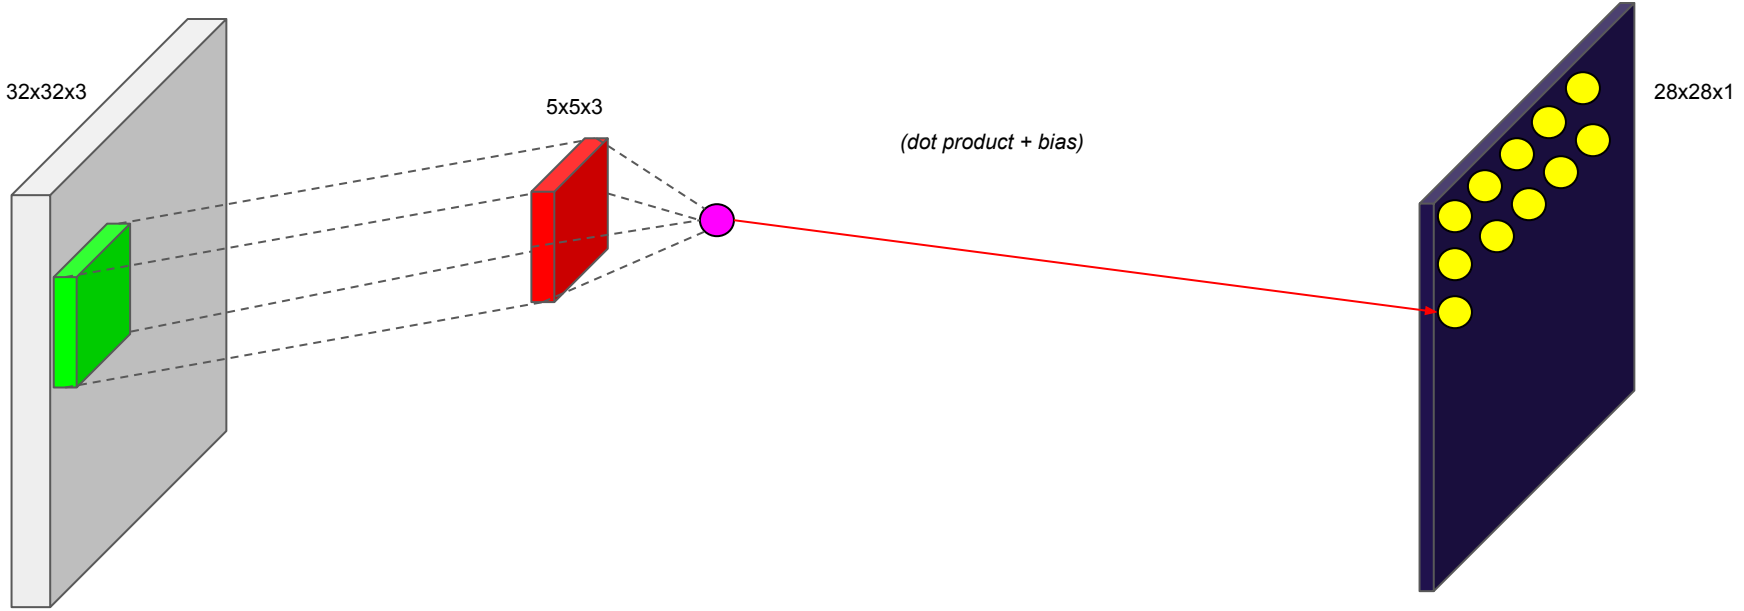

Convolutional Layers

Convolve (slide) over all spatial locations

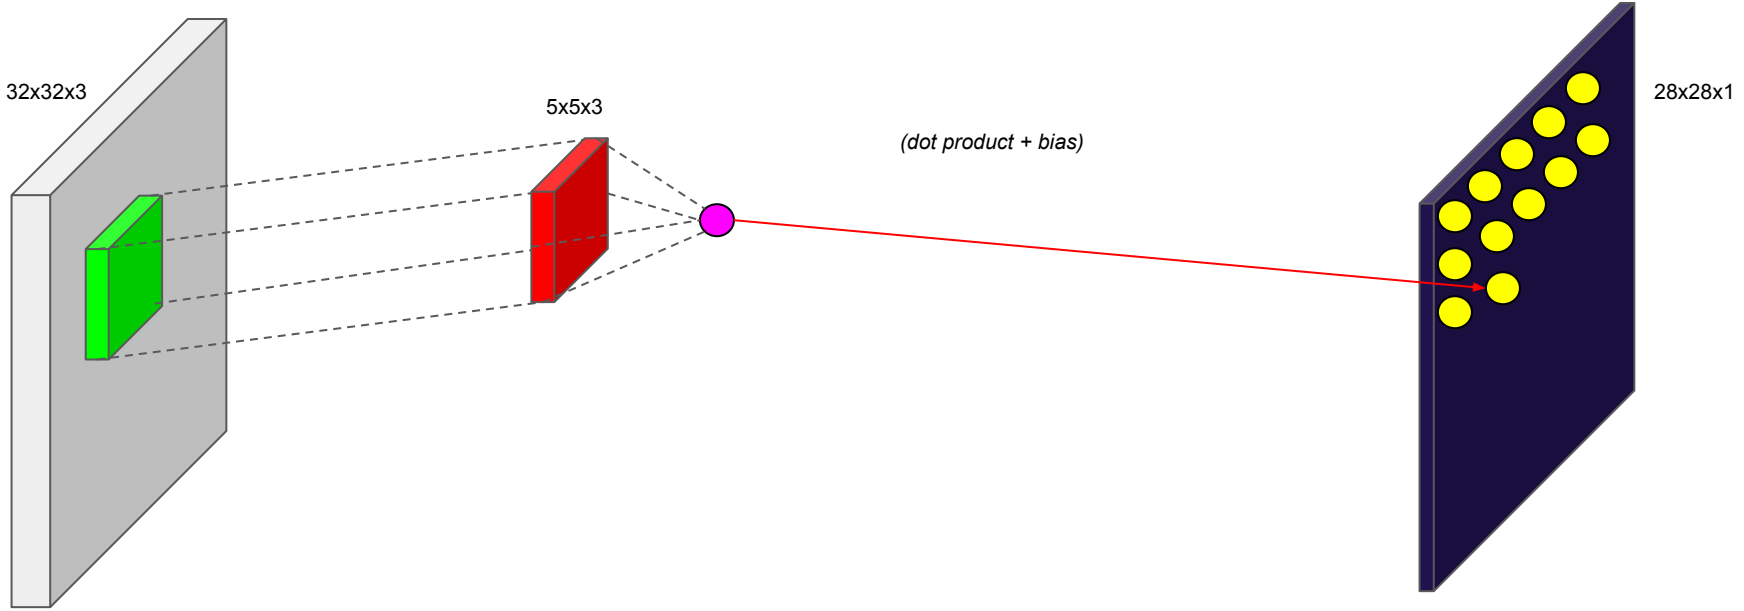

Convolutional Layers

Convolve (slide) over all spatial locations

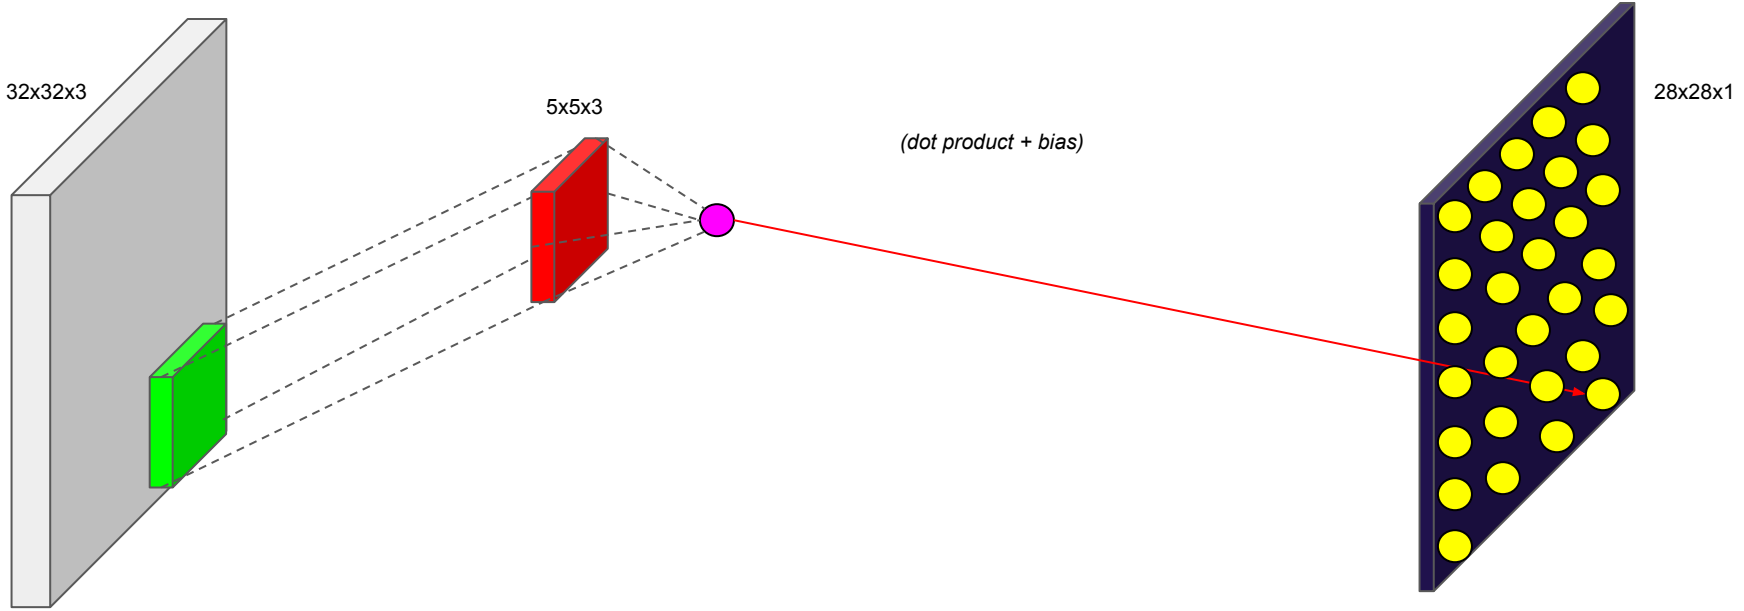

Convolutional Layers

The sample 32x32 input convolved repeatedly with 5x5 filters shrinks volumes spatially *(32 → 28 → 24)*

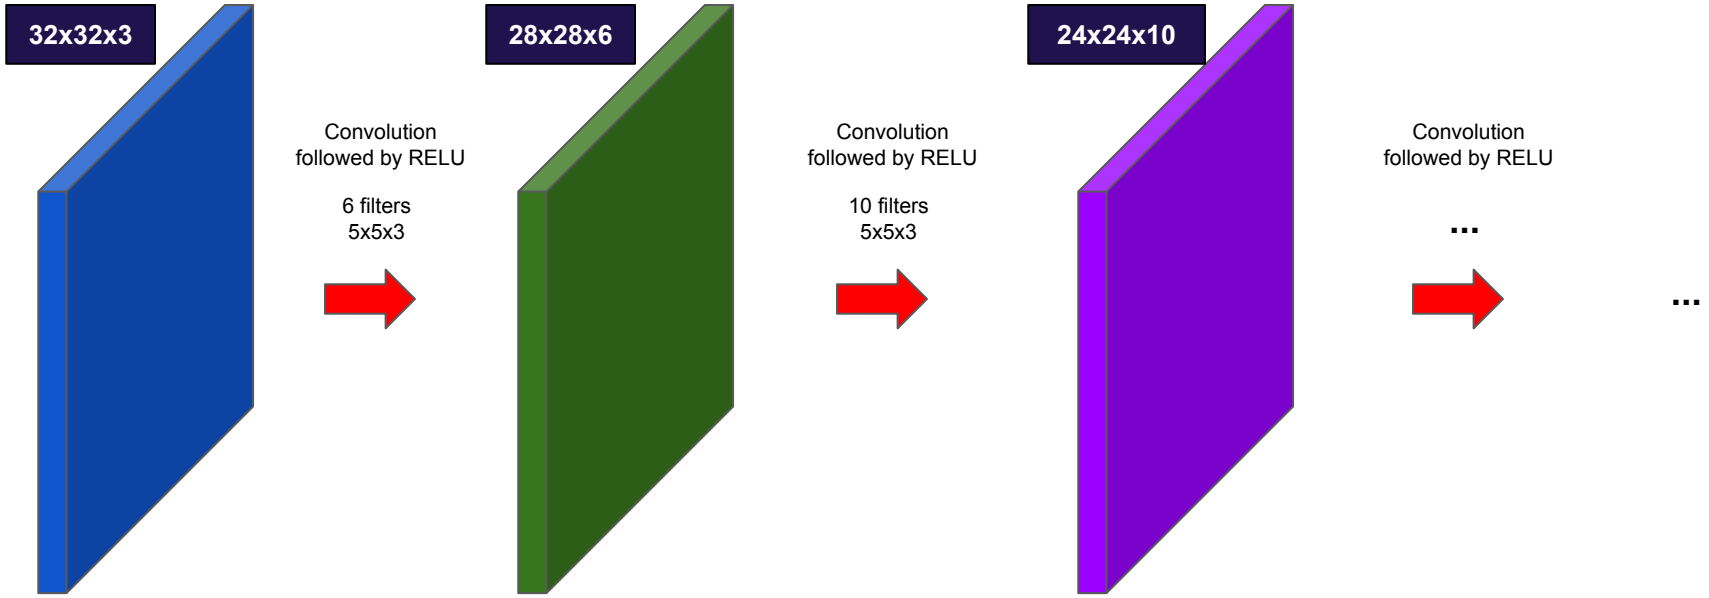

Shrinking too fast is not good. It doesn't work well.

Convolutional Layers

Hey! 1x1 convolutions make perfect sense

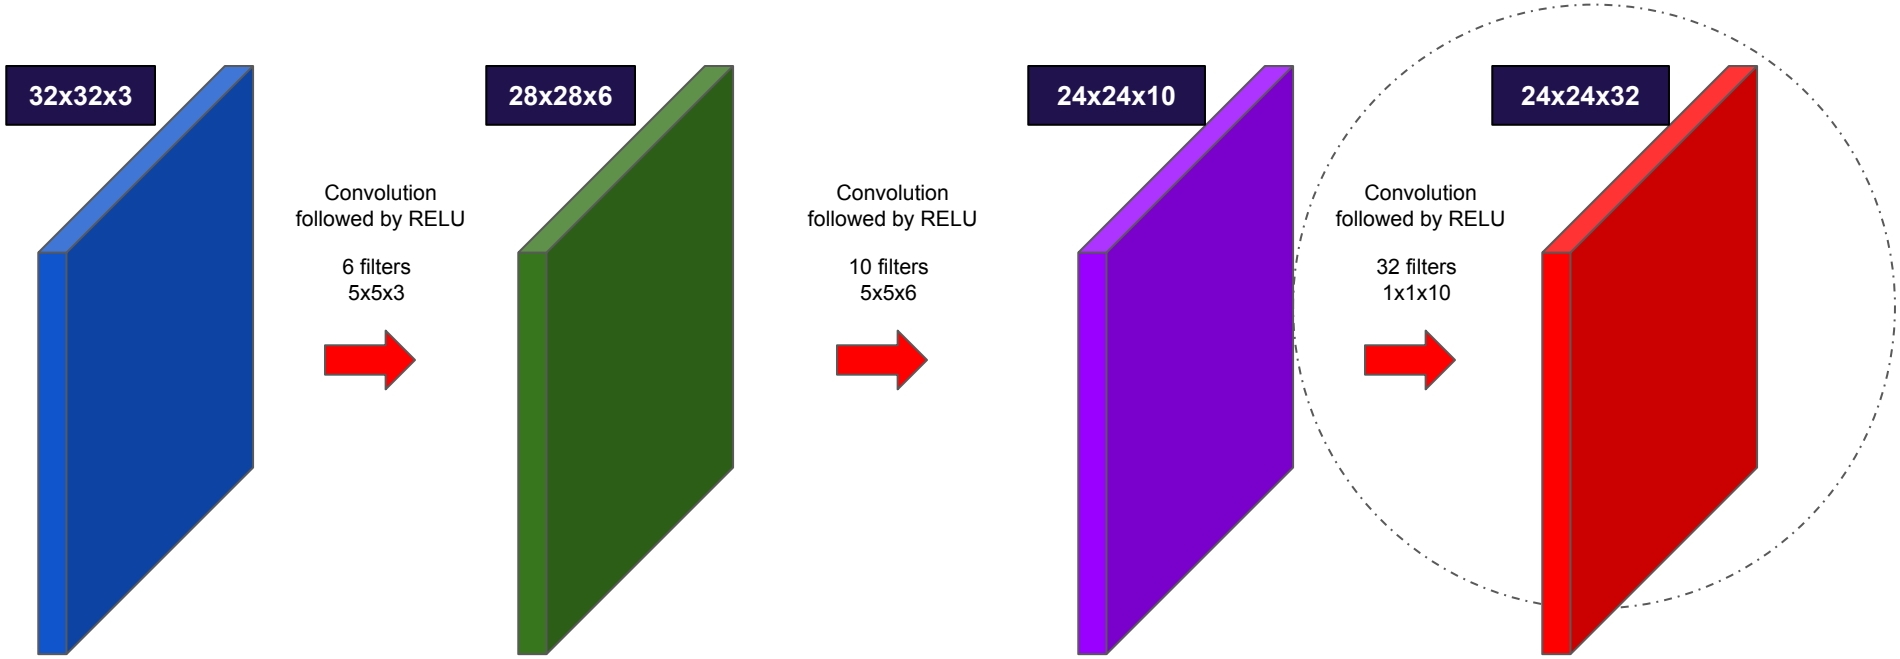

Shrinking too fast is not good. It doesn't work well.

How to calculate the output size

Calculating the output size

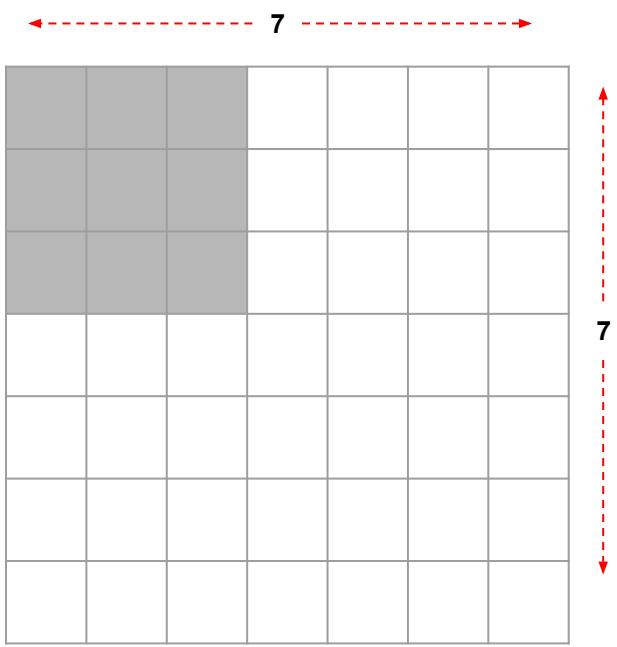

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Stride = the step of the sliding window

Example: Stride = 1

It means we will move the sliding window one position towards the right.

Movements: 1

How to calculate the output size

Calculating the output size

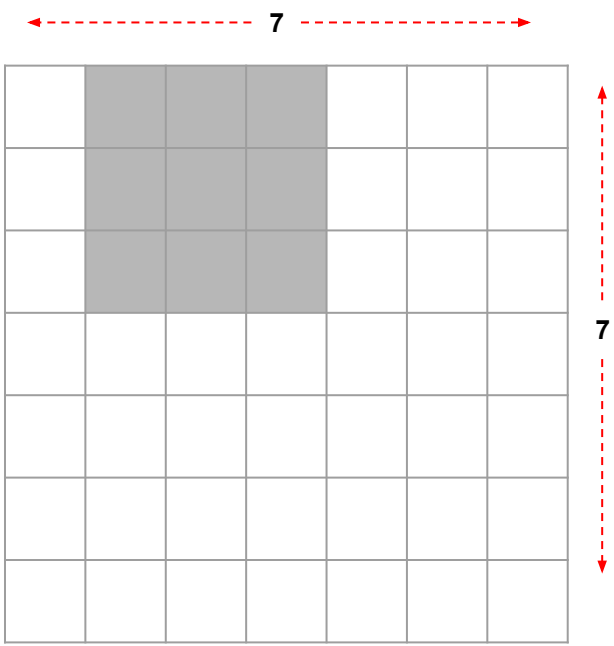

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Stride = the step of the sliding window

Example: Stride = 1

It means we will move the sliding window one position towards the right.

Movements: 2

How to calculate the output size

Calculating the output size

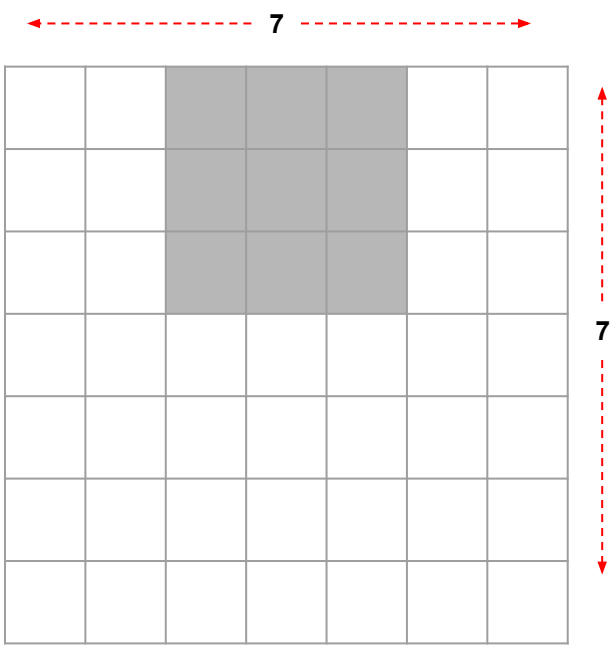

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Stride = the step of the sliding window

Example: Stride = 1

It means we will move the sliding window one position towards the right.

Movements: 3

How to calculate the output size

Calculating the output size

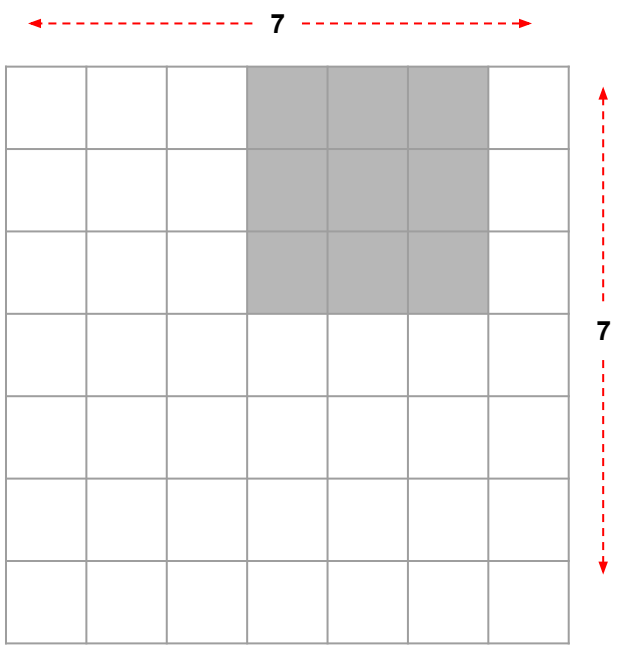

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Stride = the step of the sliding window

Example: Stride = 1

It means we will move the sliding window one position towards the right.

Movements: 4

How to calculate the output size

Calculating the output size

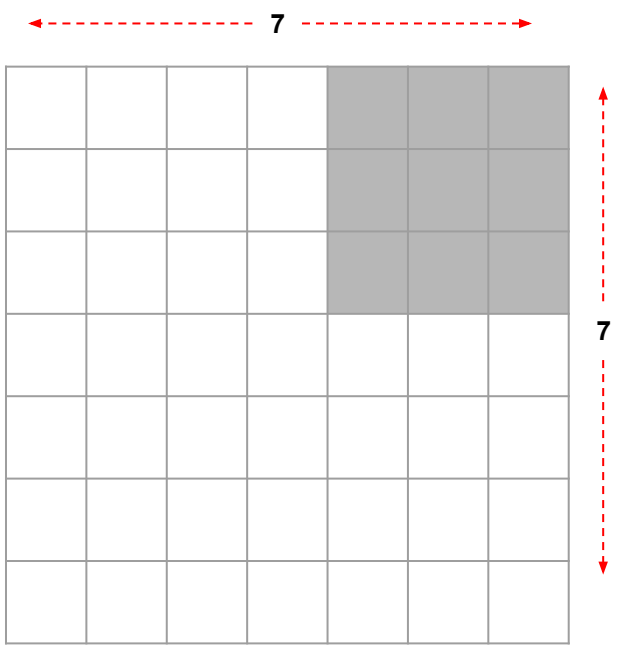

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Stride = the step of the sliding window

Example: Stride = 1

It means we will move the sliding window one position towards the right.

Output Size: 5x5

How to calculate the output size

Calculating the output size

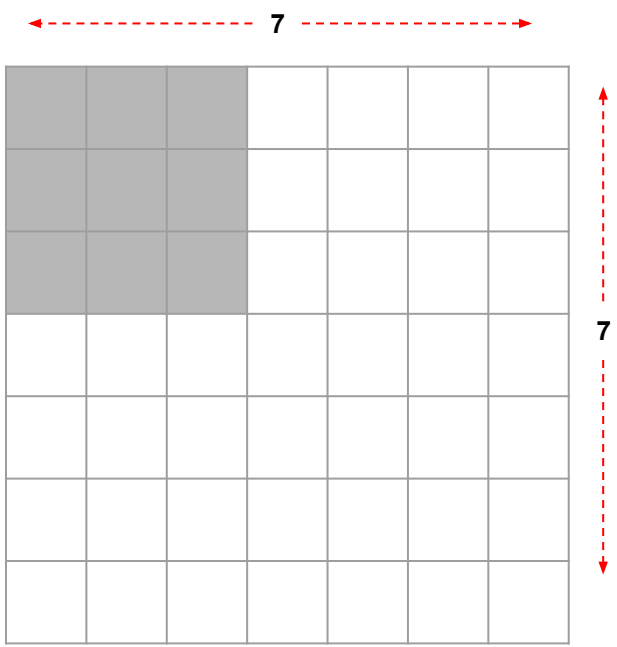

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 2

It means we will move the sliding window two positions towards the right.

Movements: 1

How to calculate the output size

Calculating the output size

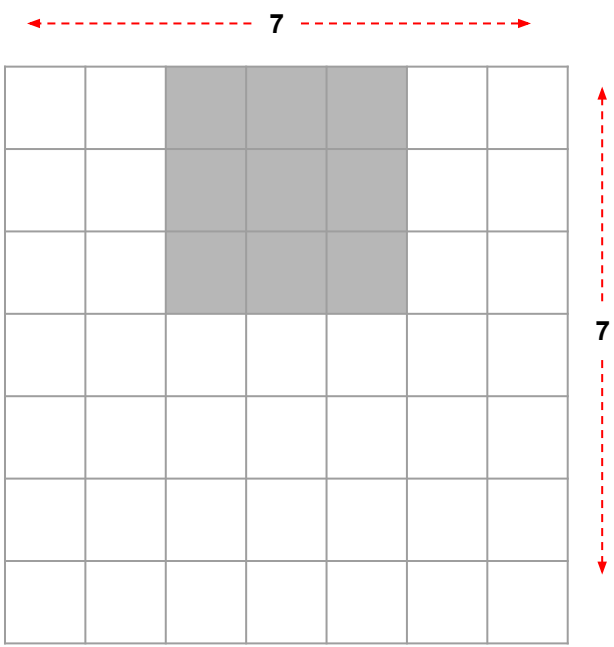

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 2

It means we will move the sliding window two positions towards the right.

Movements: 2

How to calculate the output size

Calculating the output size

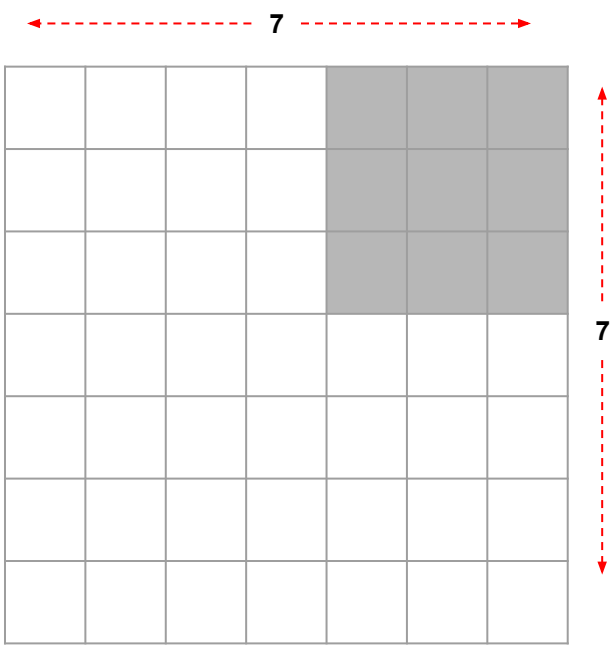

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 2

It means we will move the sliding window two positions towards the right.

Output Size: 3x3

How to calculate the output size

Calculating the output size

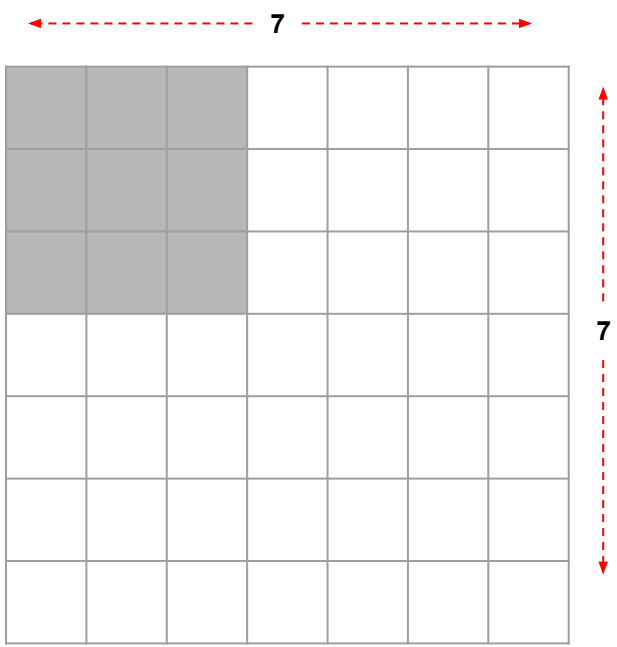

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 3

It means we will move the sliding window three positions towards the right.

Output Size: ?

How to calculate the output size

Calculating the output size

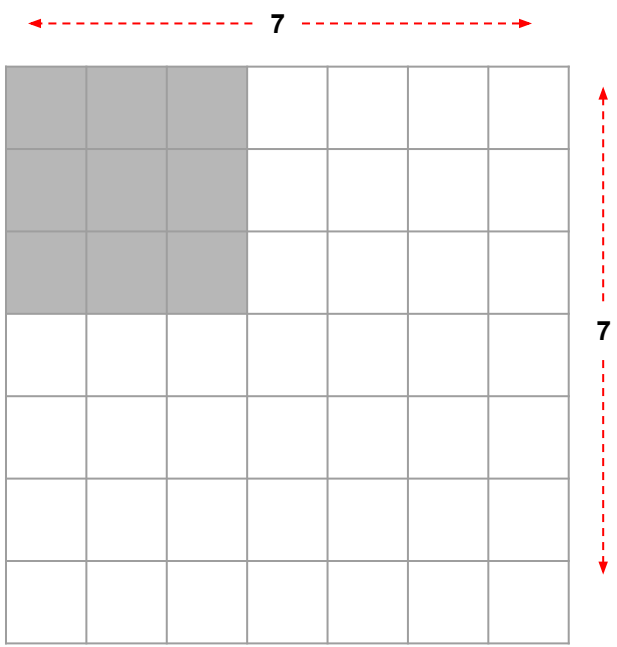

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 3

It means we will move the sliding window three positions towards the right.

Movements: 1

How to calculate the output size

Calculating the output size

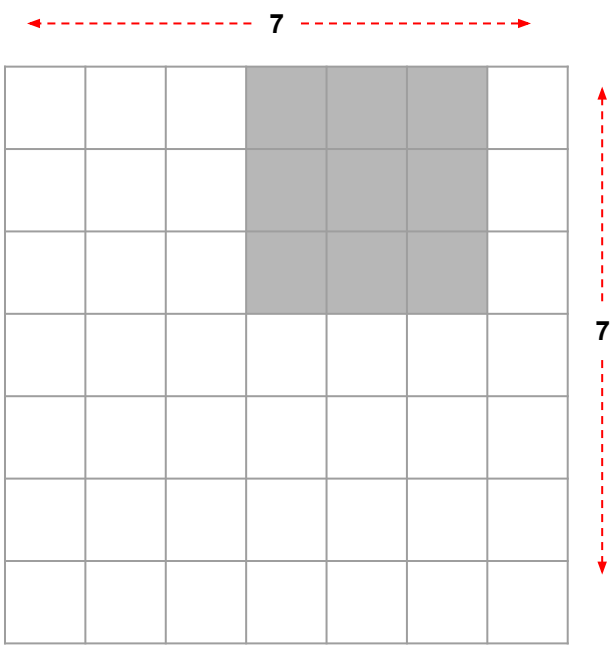

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 3

It means we will move the sliding window three positions towards the right.

Movements: 2

How to calculate the output size

Calculating the output size

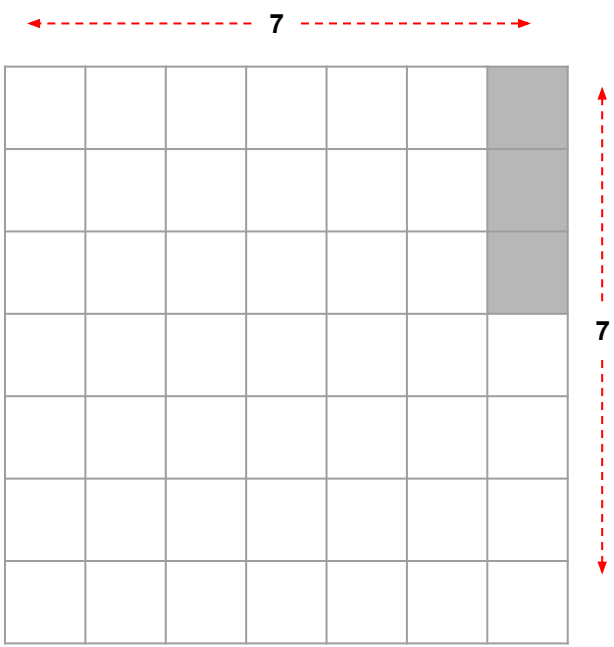

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 3

It means we will move the sliding window three positions towards the right.

Movements: 0.3333

How to calculate the output size

Calculating the output size

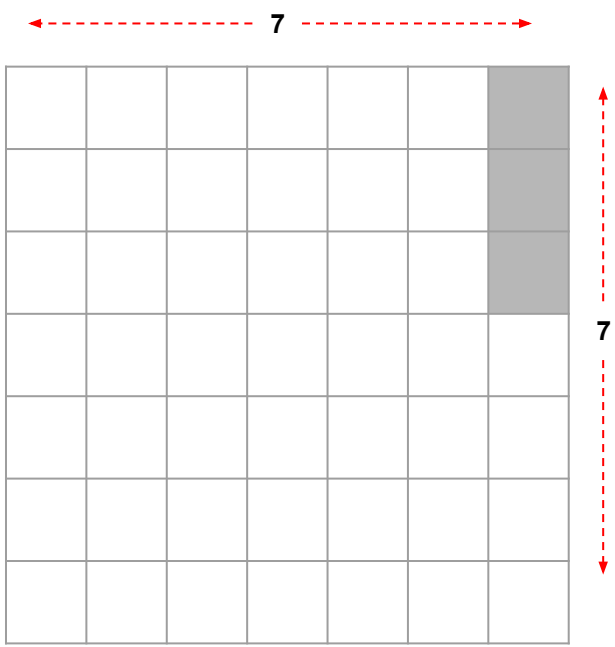

- Suppose our input patch is: 7x7
- Suppose the filter size is: 3x3

Example: Stride = 3

It means we will move the sliding window three positions towards the right.

It doesn't fit !!!

I cannot apply a 3x3 filter on a 7x7 input with stride 3

How to calculate the output size

Calculating the output size

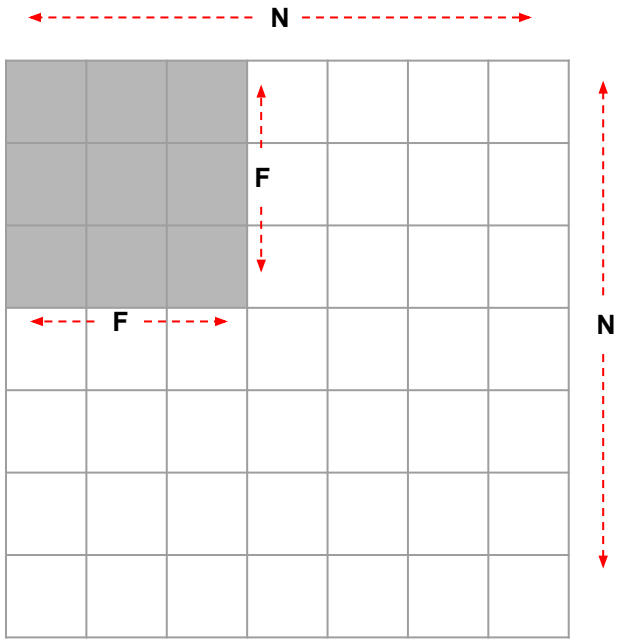

Output size = (N - F) / stride + 1

Example: N = 7 F = 3

- *Stride 1 → (7 - 3) / 1 + 1 = 5 → Output: 5x5*
- *Stride 2 → (7 - 3) / 2 + 1 = 3 → Output: 3x3*
- *Stride 3 → (7 - 3) / 3 + 1 = 2.33 → Doesn't fit !*

Padding

Zero Padding

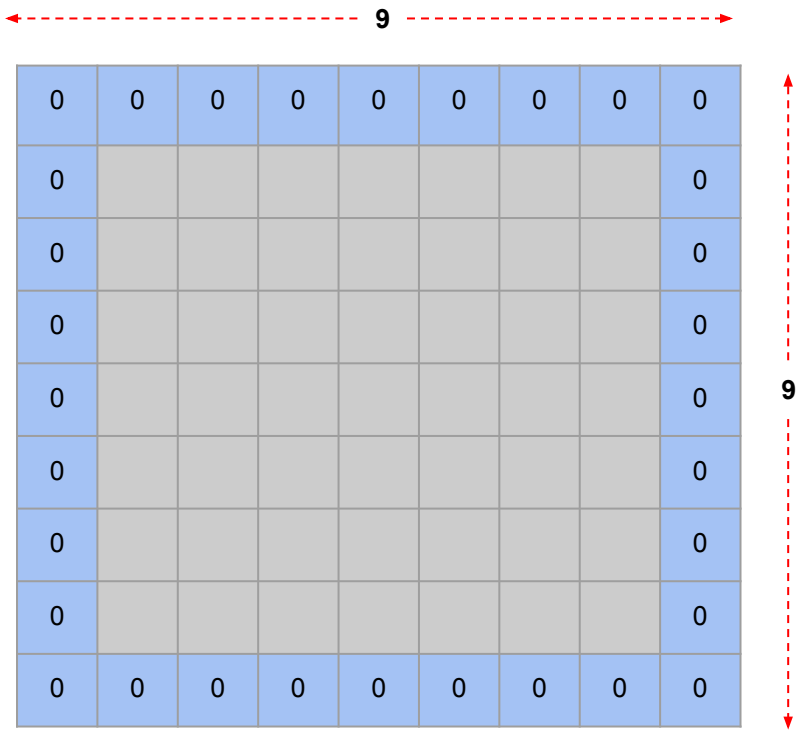

Example

- ***Input: 7x7***
- ***Filter: 3x3***
- ***Stride: 1***
- ***Padding: 1***

Output size = (N - F) / stride + 1

⇩

Output size = (7 + 2 - 3) / 1 + 1 *(padding applied)*

⇩

Output Size: 7x7

Padding

Zero Padding ***(imagine the matrix below is 36x36)***

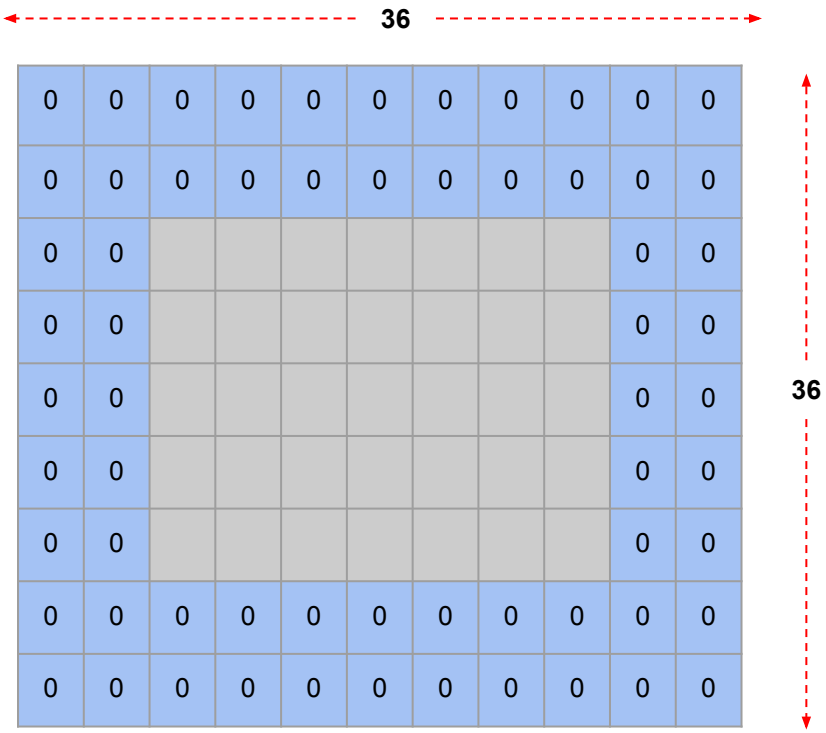

Example

- ***Input: 32x32***
- ***Filter: 5x5***
- ***Stride: 1***
- ***Padding: 2***

Output size = (N - F) / stride + 1

⇩

Output size = (32 + 4 - 5) / 1 + 1 *(padding applied)*

⇩

Output Size: 32x32

How to calculate the *Number of Parameters*

Calculating the number of parameters

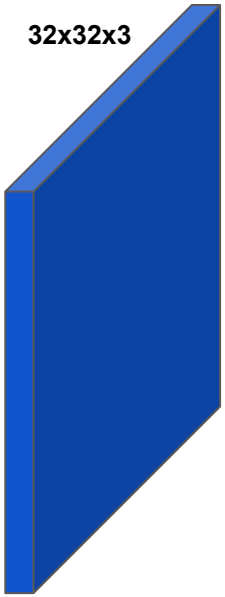

Filters bank:
10 filters of size 5x5x3

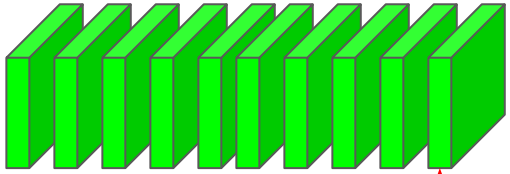

- Input Volume: **32x32x3**
- **10** filters of size **5x5x3**
- Stride 1, padding 2

Number of parameters per filter

Each filter has **5x5x3** parameters = **75 + 1** *(plus one for the bias)* = **76** parameters

Number of parameters in the Layer

76 parameters x **10** filters = **760** parameters

Pooling Layers

Makes the representations smaller and more manageable without losing too much information

Max Pooling *(most common)*

Get the max value in the neighborhood, and pass it forward.

4x4

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

The maximum value

Max Pooling
Filter size = **2x2**
Stride = **2**

2x2

6	8
3	4

Types

- **Max Pooling**
- **Average Pooling**
- *etc*

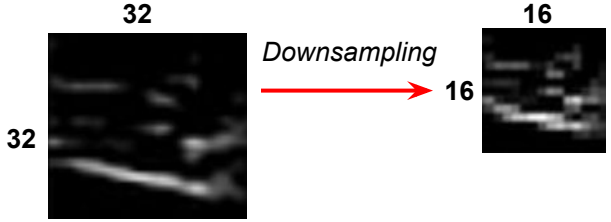

ReLU Layers

Think like a Proton → Stay positive → If negative, then set to zero.

1	0	2	4
5	-1	7	8
3	2	-22	0
1	-4	3	4

Negative values are set to zero

1	0	2	4
5	0	7	8
3	2	0	0
1	0	3	4

Fully Connected Layers

Contain neurons that connect to the entire input volume, as in ordinary *Multi-Layered Perceptrons*

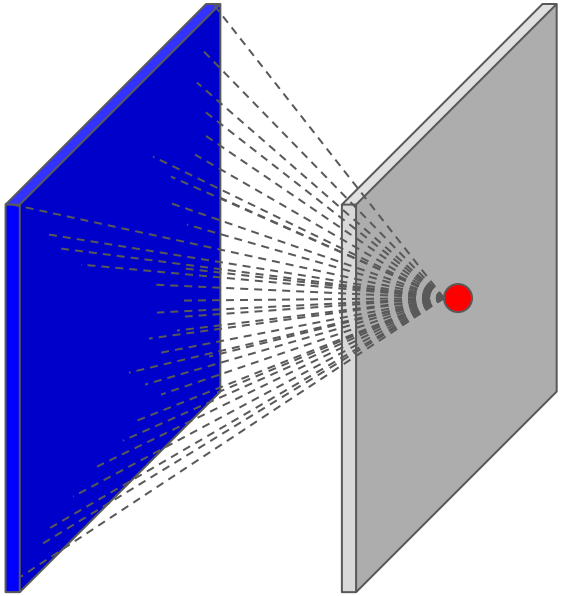

This neuron is connected to all of the neurons available in the previous layer

Architectures

LeNet-5

[LeCun et al., 1998]

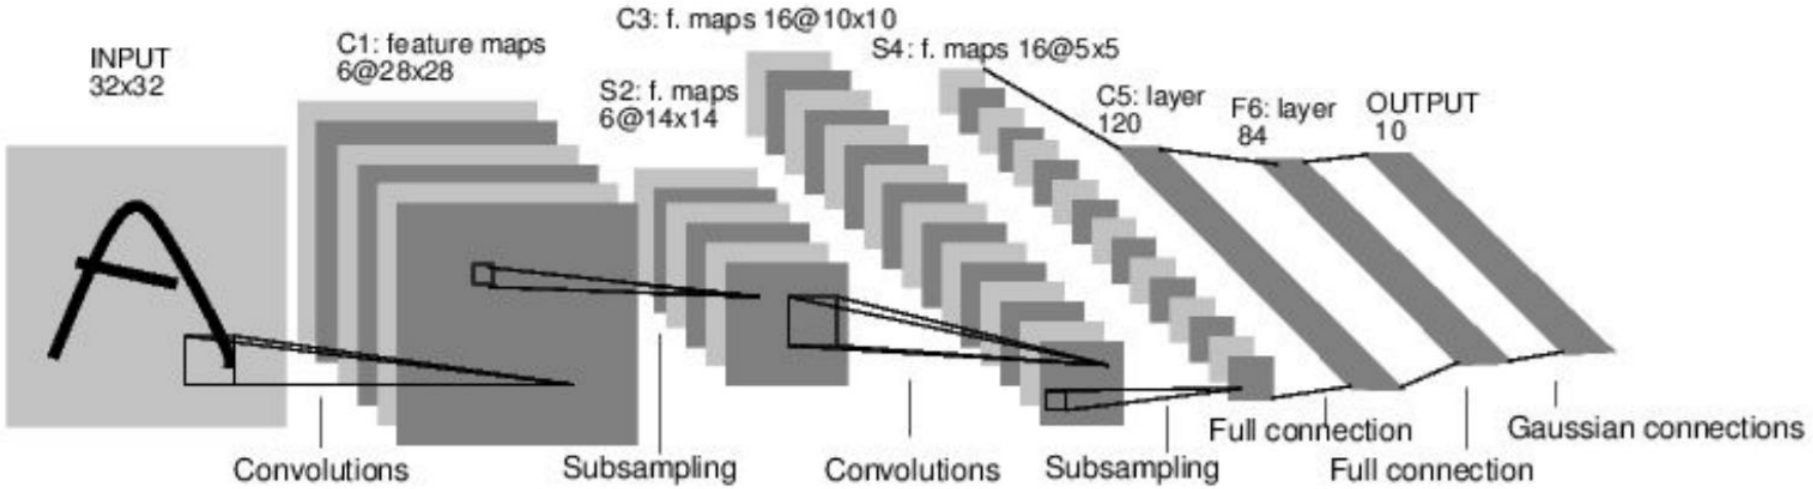

Conv filters were 5x5, applied at stride 1
Subsampling (Pooling) layers were 2x2 applied at stride 2

Case Study: GoogLeNet

[Szegedy et al., 2014]

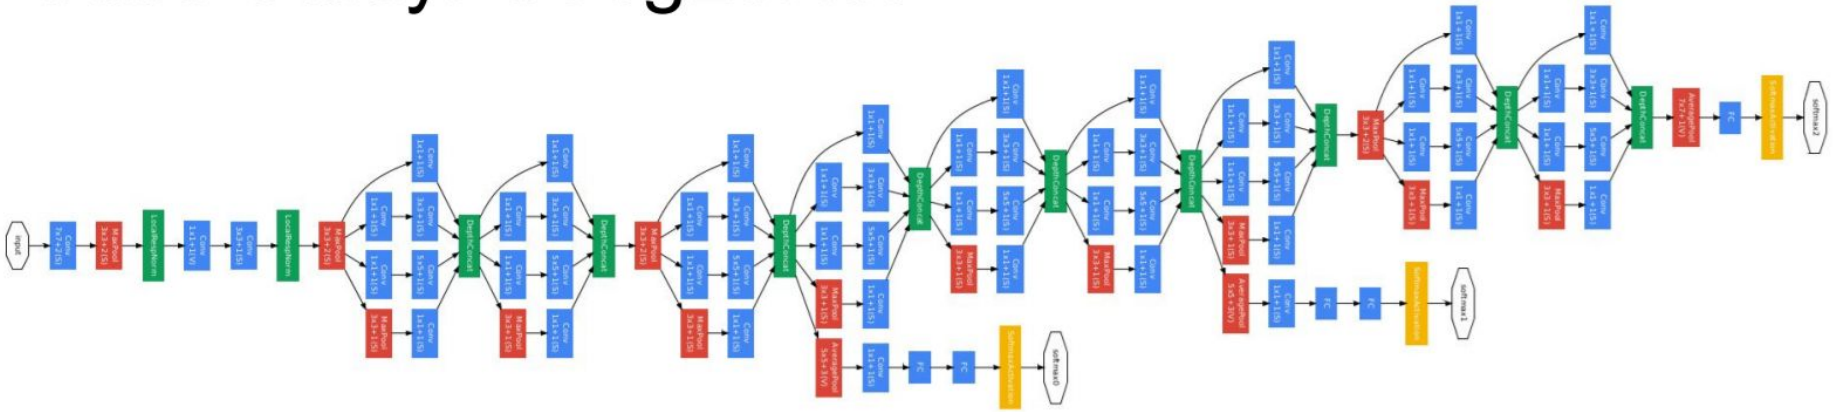

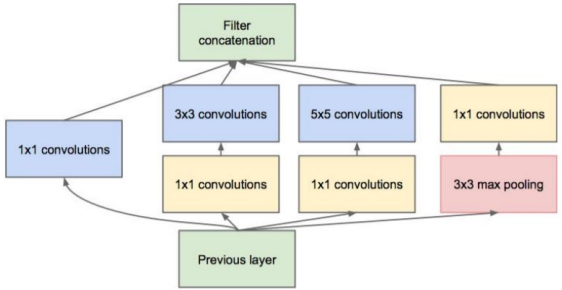

Inception module

ILSVRC 2014 winner (6.7% top 5 error)

Can networks become creative?

As seen on the Press / TV

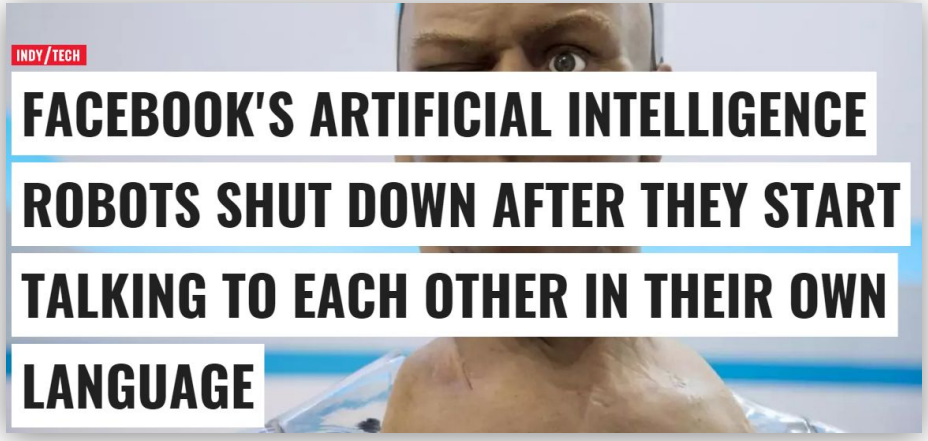

Can networks become creative?

Yes.

More science than fiction.

Generating images

Our network learned to identify people, cars, cats and other classes.

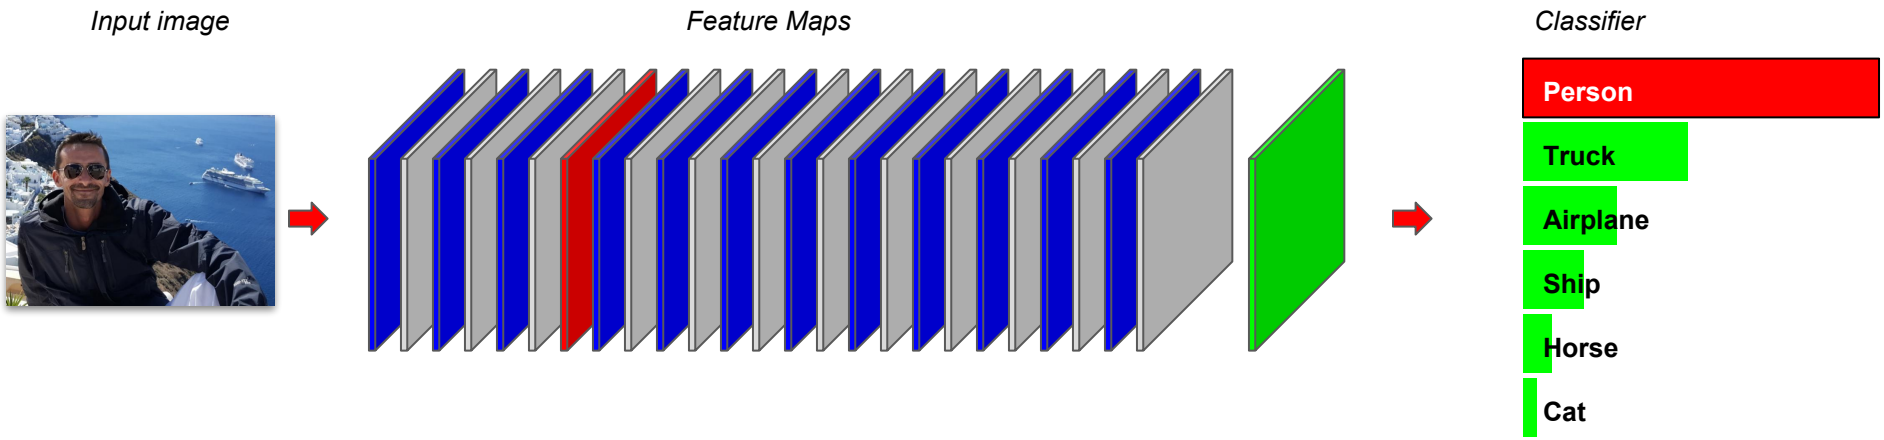

Can we generate a new image from scratch?

Generating images

The goal is to maximize the score of a given class (e.g. “Person”)

Input image

(This is a matrix. All values were set to zero)

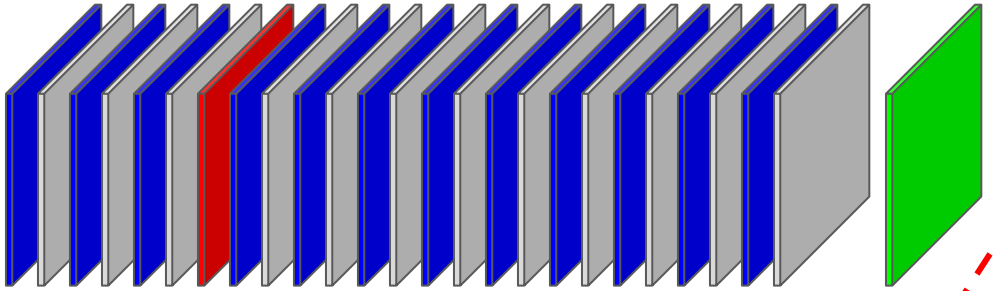

Gradients of the Score vector

1	Person
0	Truck
0	Airplane
0	Ship
0	Horse
0	Cat

Step #1:
Set the gradient of the scores vector to be [0, 0,....1,....,0]

Generating images

The goal is to maximize the score of a given class (e.g. “Person”)

Input image

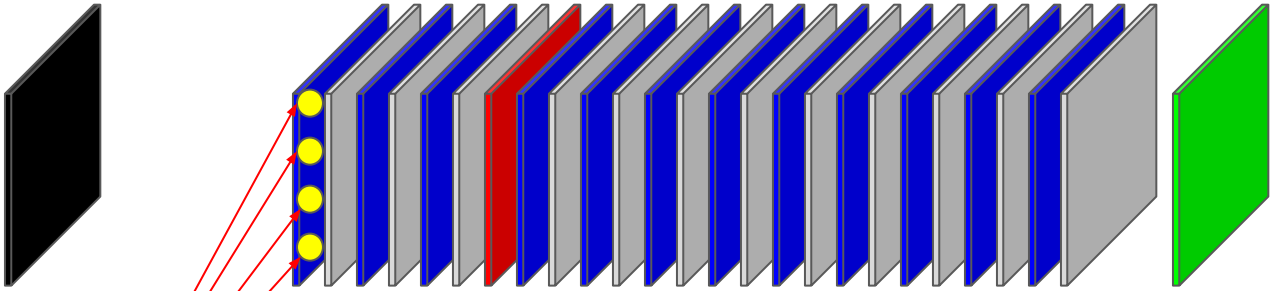

(This is a matrix. All values were set to zero)

Backpropagation

Gradients of the Score vector

1	Person
0	Truck
0	Airplane
0	Ship
0	Horse
0	Cat

Hey! I’ve got some gradients here!

Step #2:
Propagate the error backwards (compute the gradients)

Generating images

The goal is to maximize the score of a given class (e.g. “Person”)

Input image

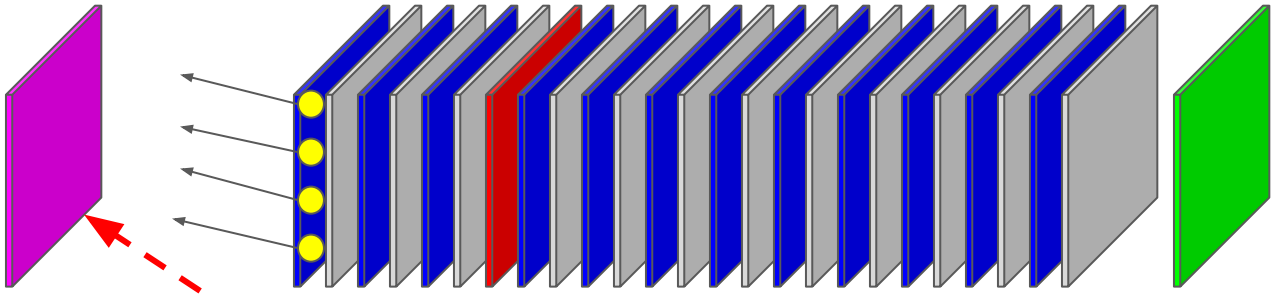

The image is no longer zero

Gradients of the Score vector

1	Person
0	Truck
0	Airplane
0	Ship
0	Horse
0	Cat

Step #3:
Update the input image

Generating images

The goal is to maximize the score of a given class (e.g. “Person”)

Input image

Forward pass

Gradients of the Score vector

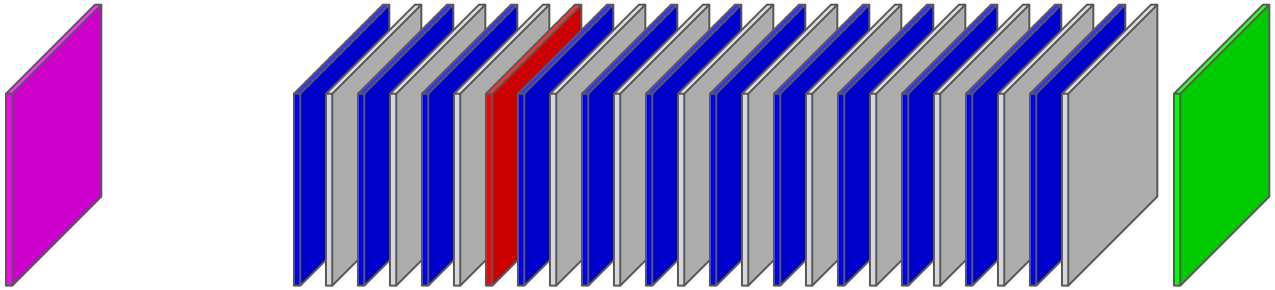

The image is no longer zero

1	Person
0	Truck
0	Airplane
0	Ship
0	Horse
0	Cat

Step #4:
Forward the image through the network.

Generating images

The goal is to maximize the score of a given class (e.g. “Person”)

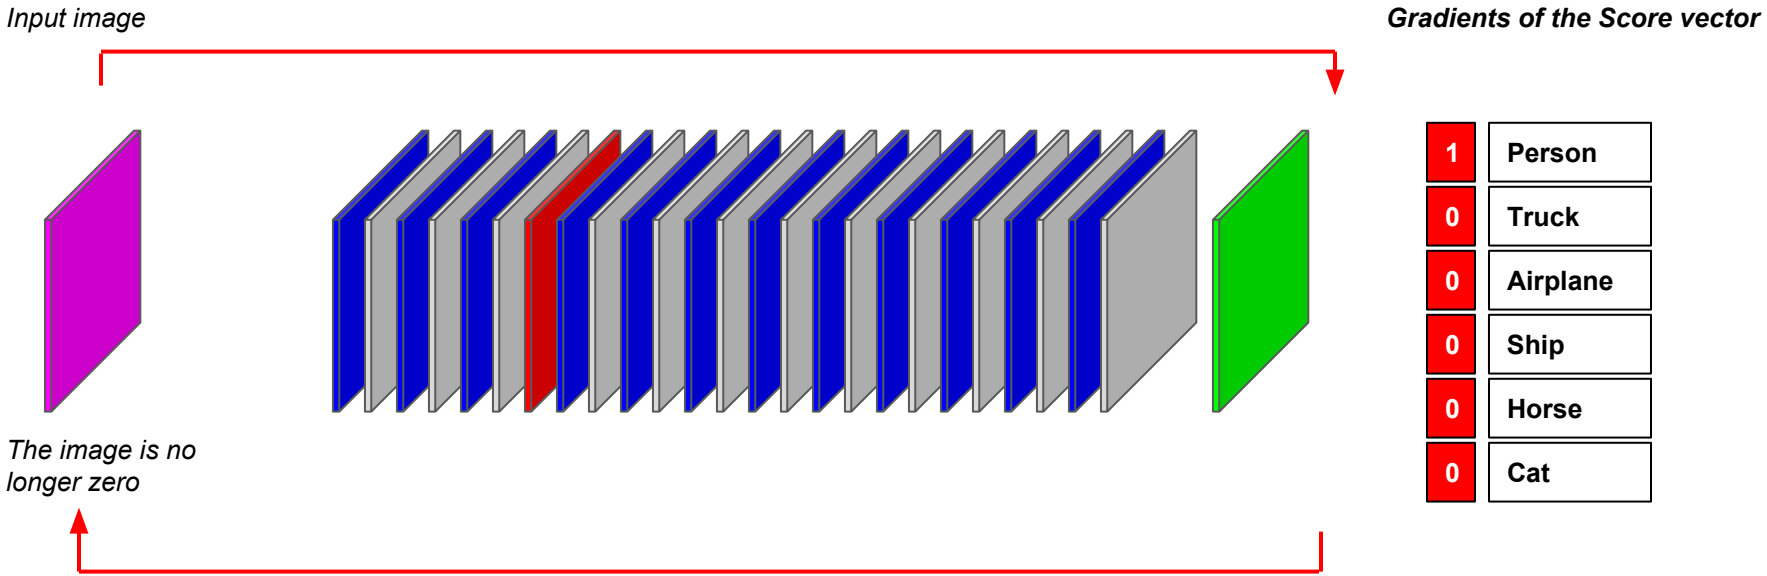

Go back to step #1 (repeat the process again)
Set the gradient of the scores vector to be [0, 0,....1,....,0]

Generating images *(summary)*

The goal is to maximize the score of a given class (e.g. “Person”)

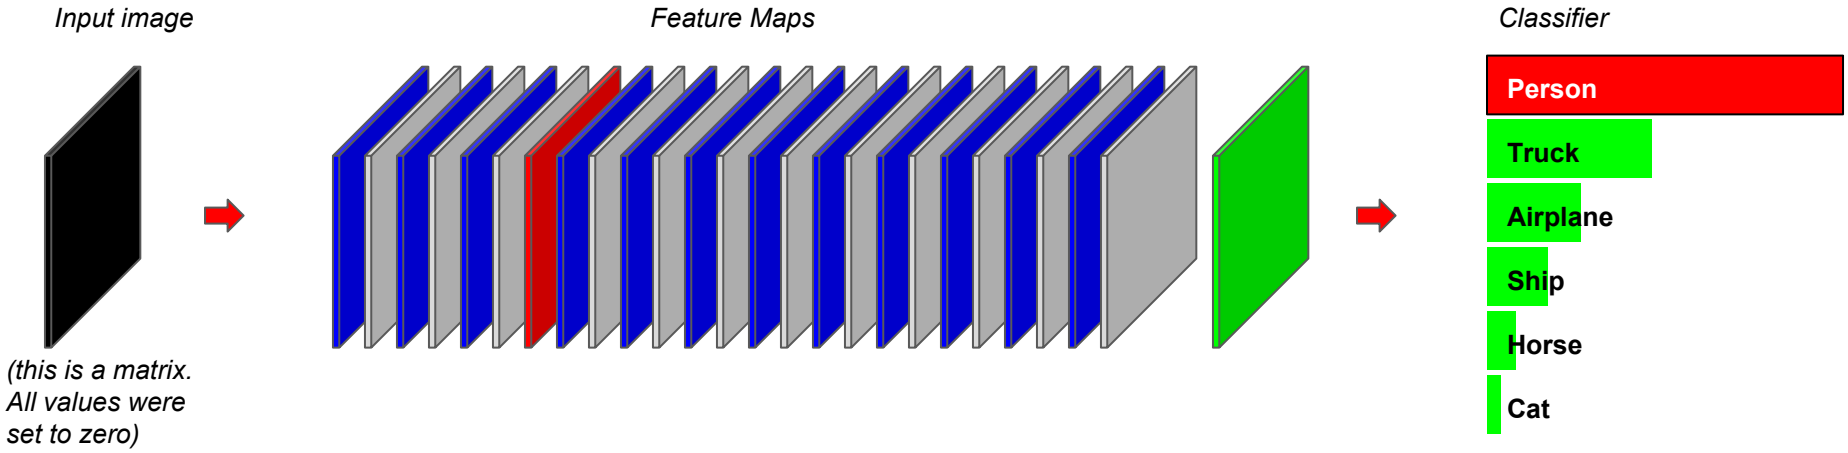

Step #1: Set the gradient of the scores vector to be [0, 0,....1,....,0]
Step #2: Backpropagate the error
Step #3: do a small “image update”
Step #4: forward the image through the network.
Step #5: go back to step #1

Generating images *(examples)*

dumbbell

cup

dalmatian

bell pepper

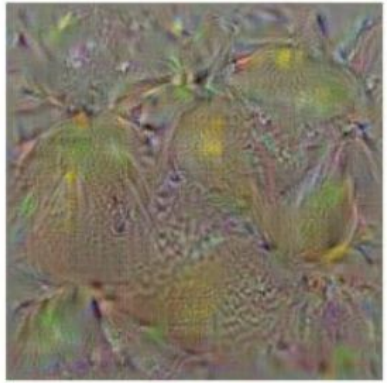

lemon

husky

Generating images

Hey! We could also boost specific activations and neurons

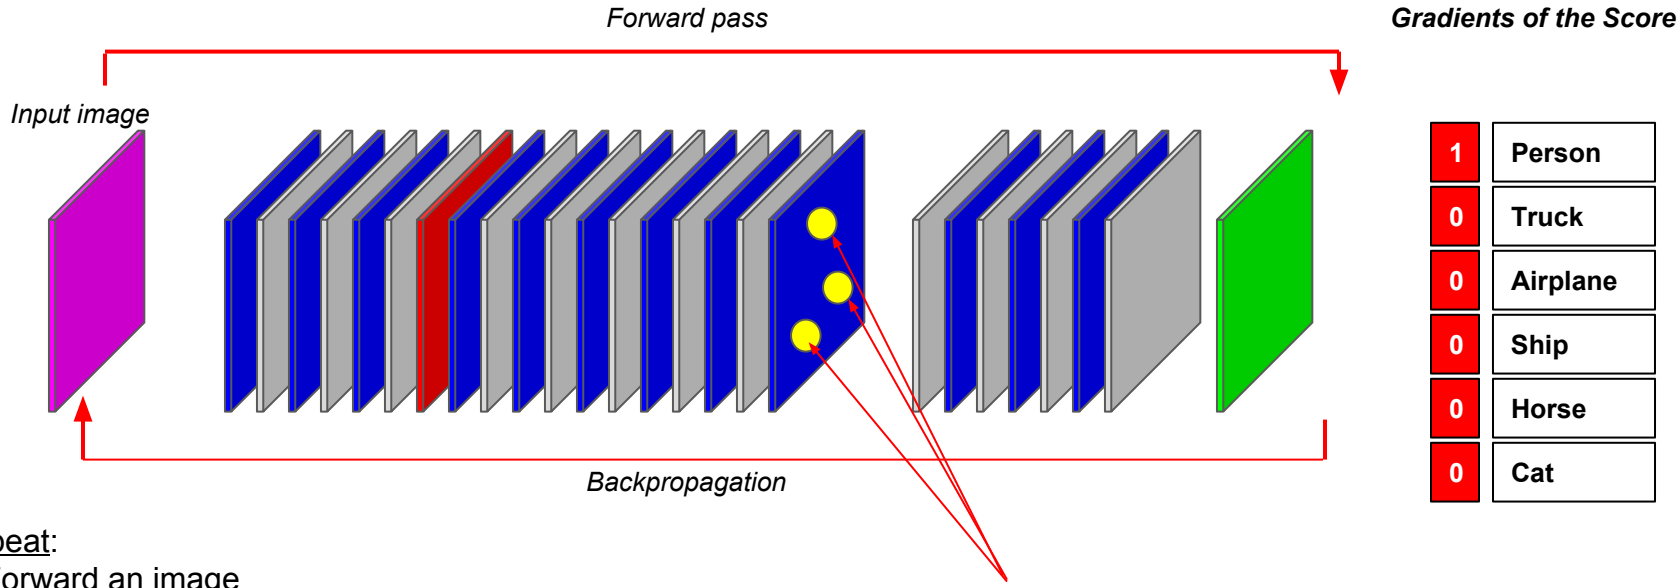

Repeat:

1. Forward an image
2. Set activations in layer of interest to all zero, except for a 1.0 for a neuron of interest
3. Backprop to image
4. Do an “image update”

Generating images

Hey! We could also boost specific activations and neurons

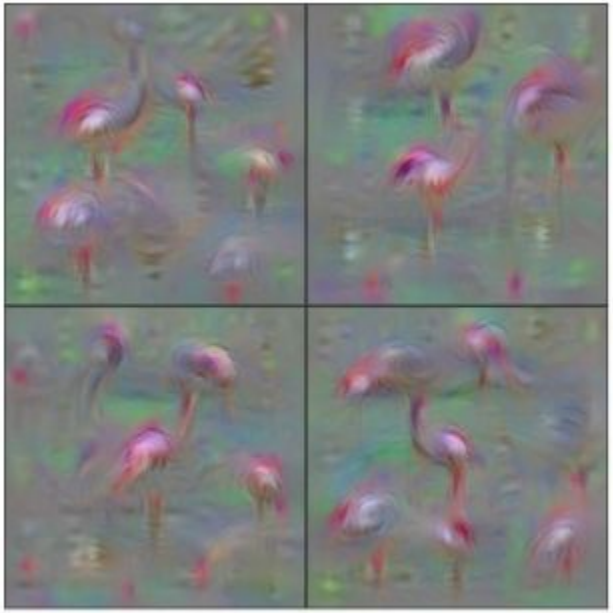

Flamingo

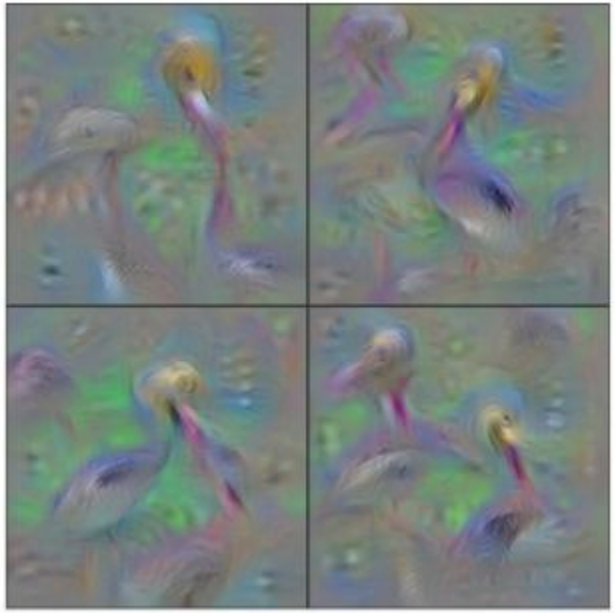

Pelican

Dreams *(and nightmares)*

It happens when you boost *(maximize/minimize)* all activations, at any layer

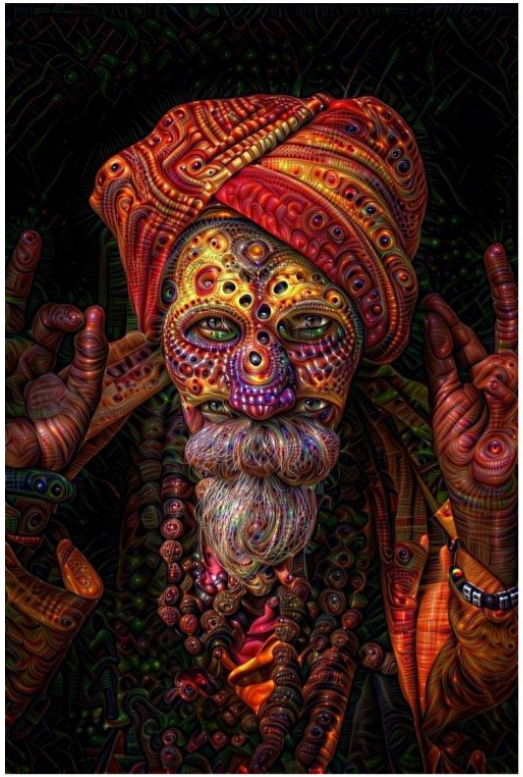

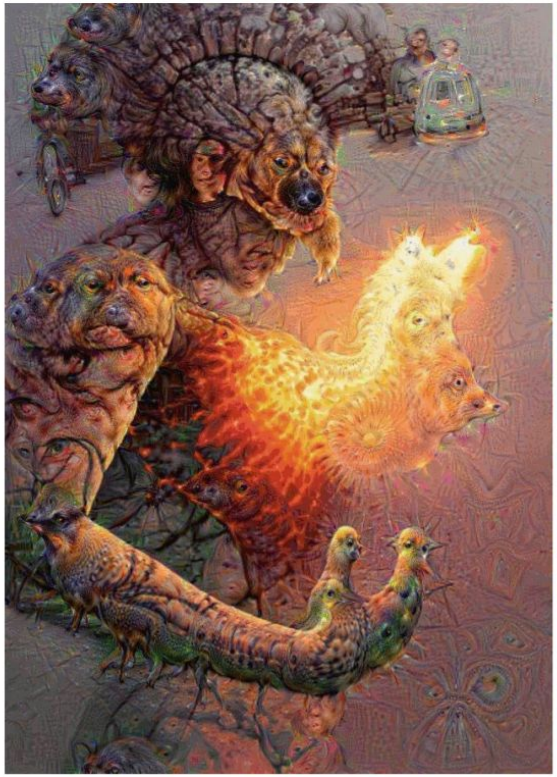

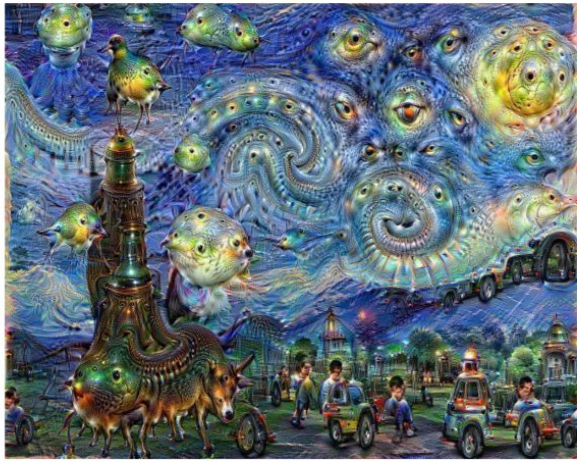

Dreams *(and nightmares)*

It happens when you boost *(maximize/minimize)* all activations, at any layer

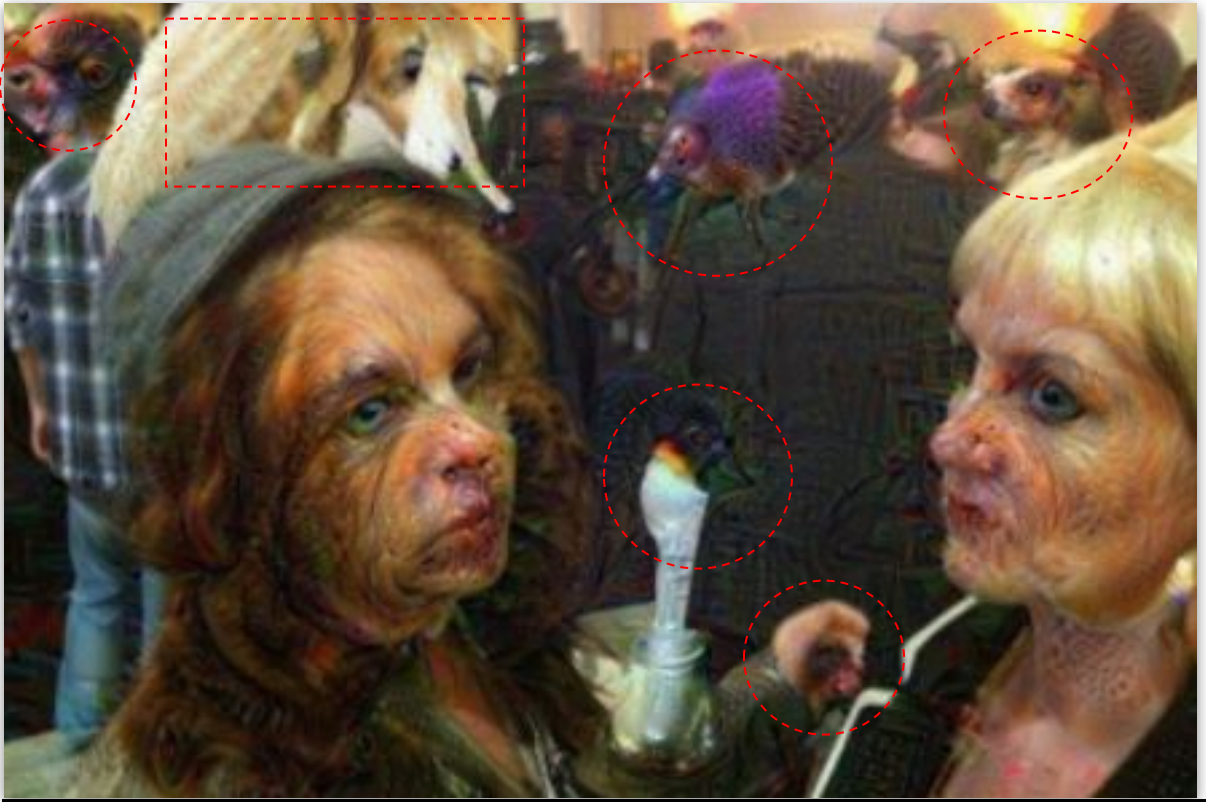

This creates a feedback loop.

For example: any slightly detected dog face will be made more and more dog like over time.

Dreams *(and nightmares)*

It happens when you boost *(maximize/minimize)* all activations, at any layer

Image learned by the network

The dream produced by the Neural Network

Dreams *(and nightmares)*

It happens when you boost *(maximize/minimize)* all activations, at any layer

Image learned by the network

The dream produced by the Neural Network

Dreams *(and nightmares)*

It happens when you boost *(maximize/minimize)* all activations, at any layer

Image learned by the network

The dream produced by the Neural Network

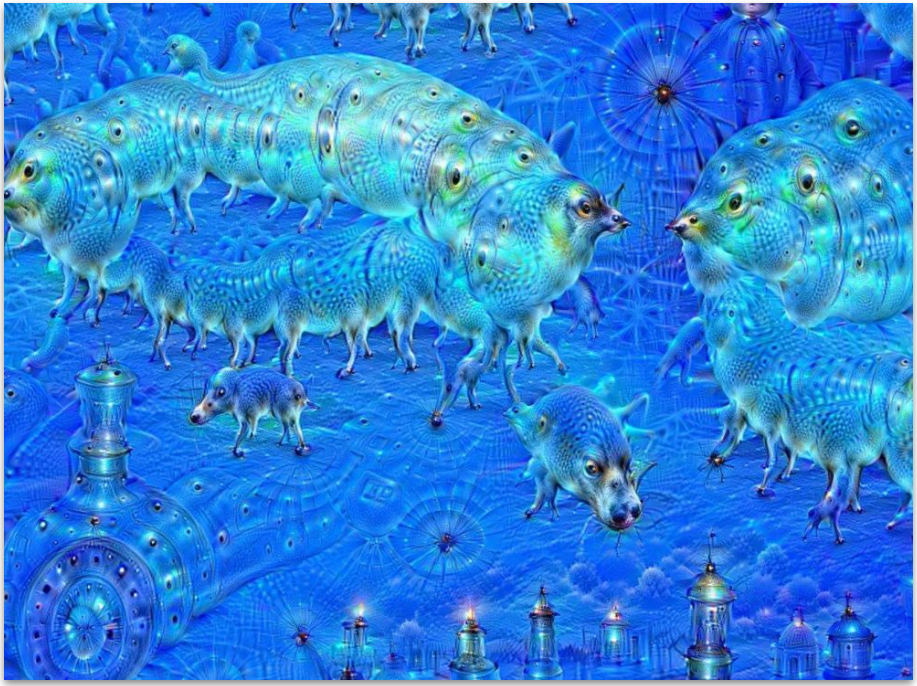

Dreams *(and nightmares)*

How it works

The network has been trained to identify things like *cars, cats, people, airplanes and other objects.*

Input image

Feature Maps

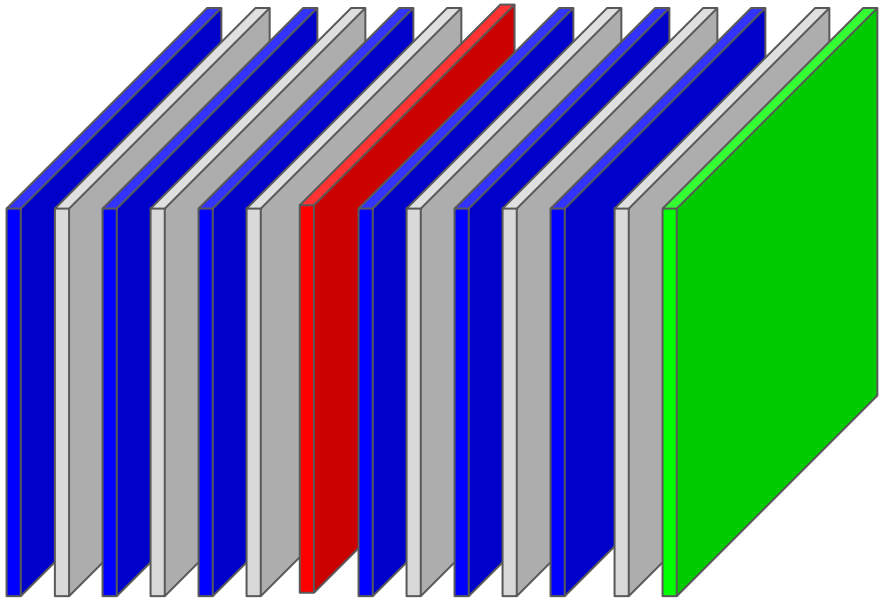

Classifier

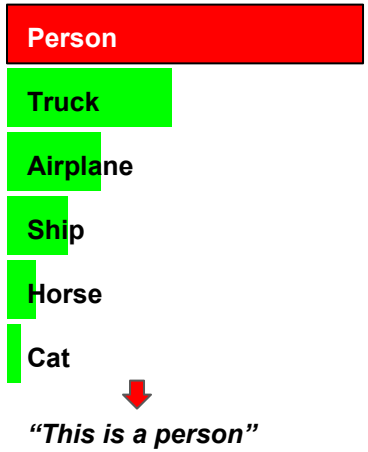

Conv
ReLU
Pooling
Fully Connected

Dreams *(and nightmares)*

How it works

It knows *what filters need to be applied* in order to produce each of the feature maps below.
It learned that through backpropagation.

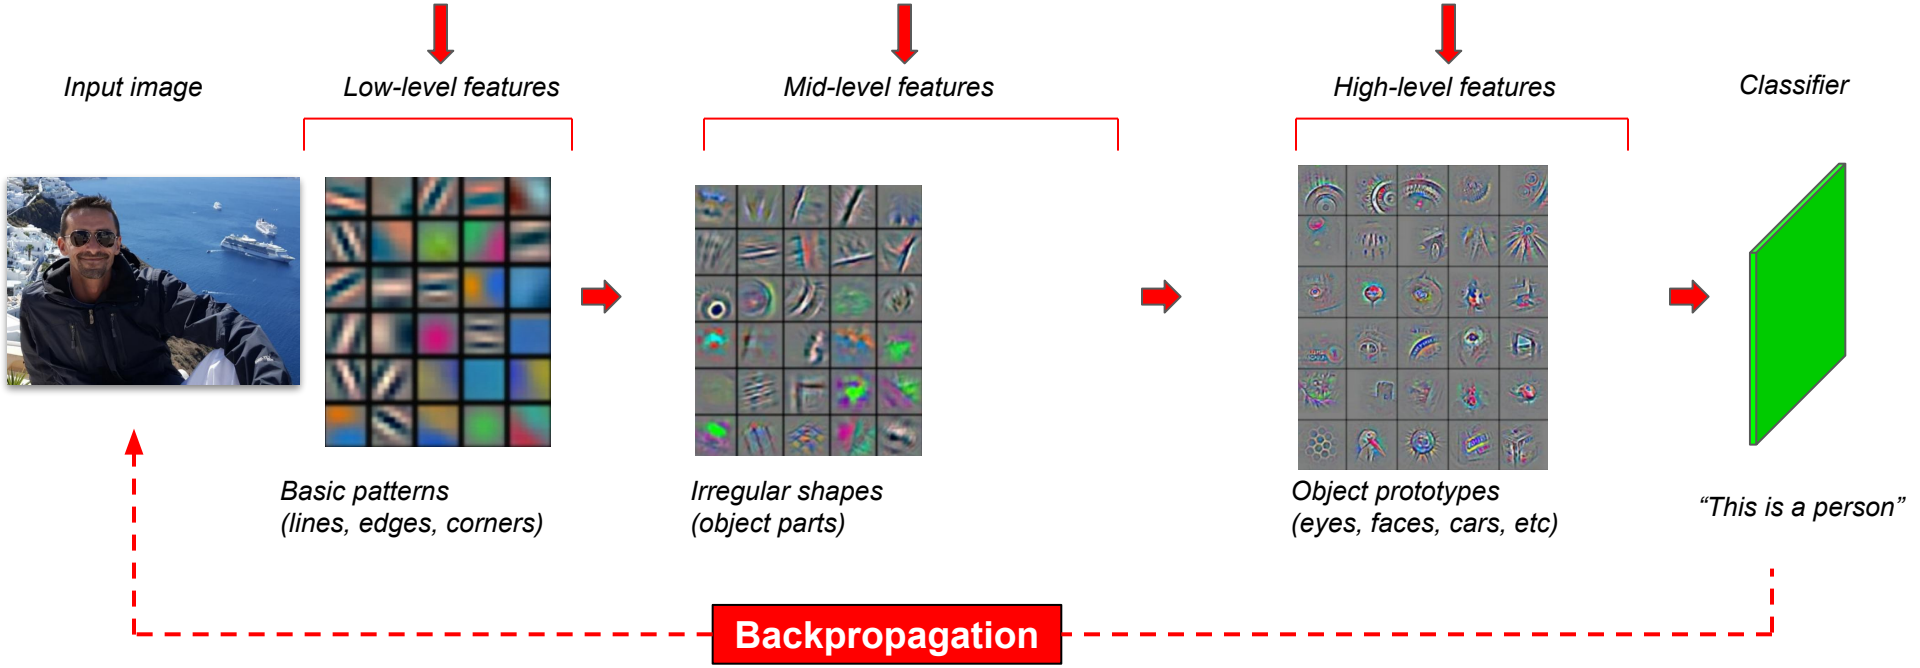

Dreams *(and nightmares)*

How it works

It knows *what filters need to be applied* in order to produce each of the feature maps below.
It learned that through backpropagation.

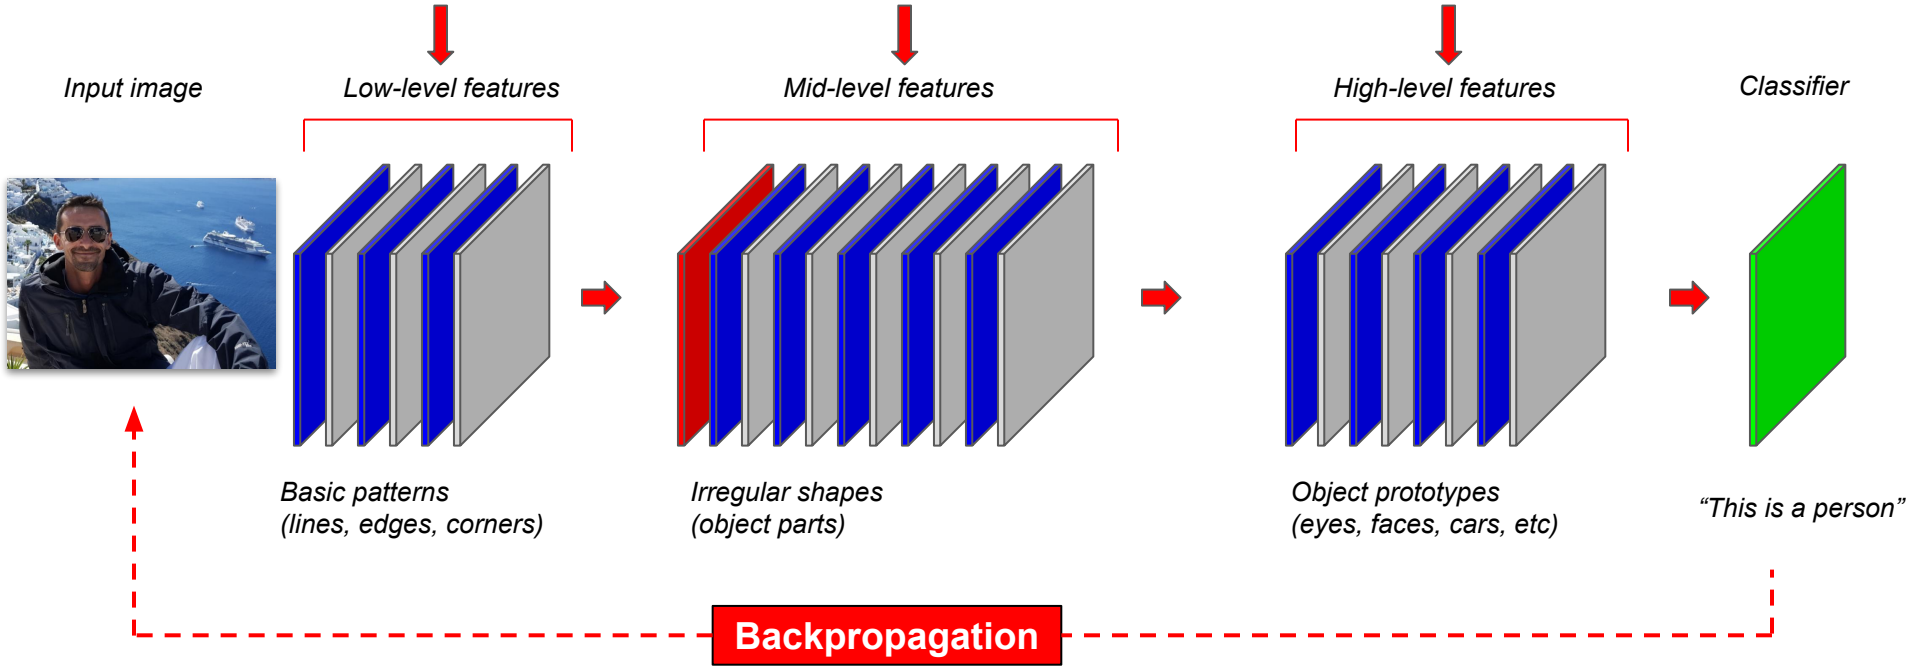

Dreams *(and nightmares)*

So what happens if I boost the values of the high-level features?

For example → "Do it more cat-like", "smaller eyes", "bigger ears"

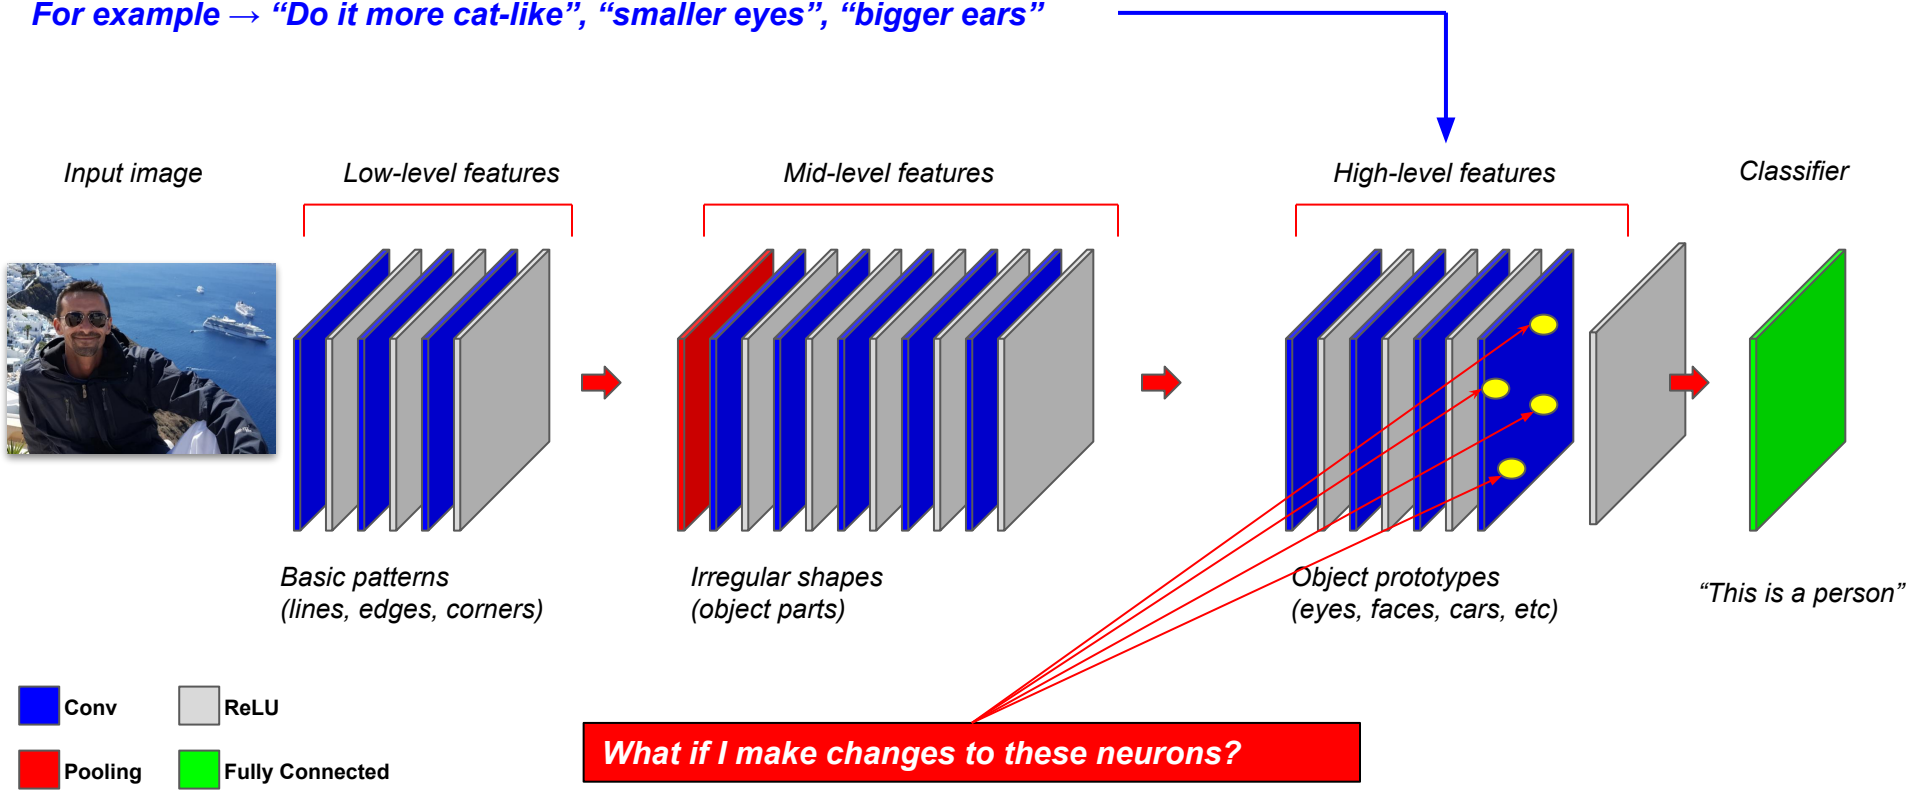

Dreams *(and nightmares)*

So what happens if I boost the values of the high-level features?

For example → "Do it more cat-like", "smaller eyes", "bigger ears"

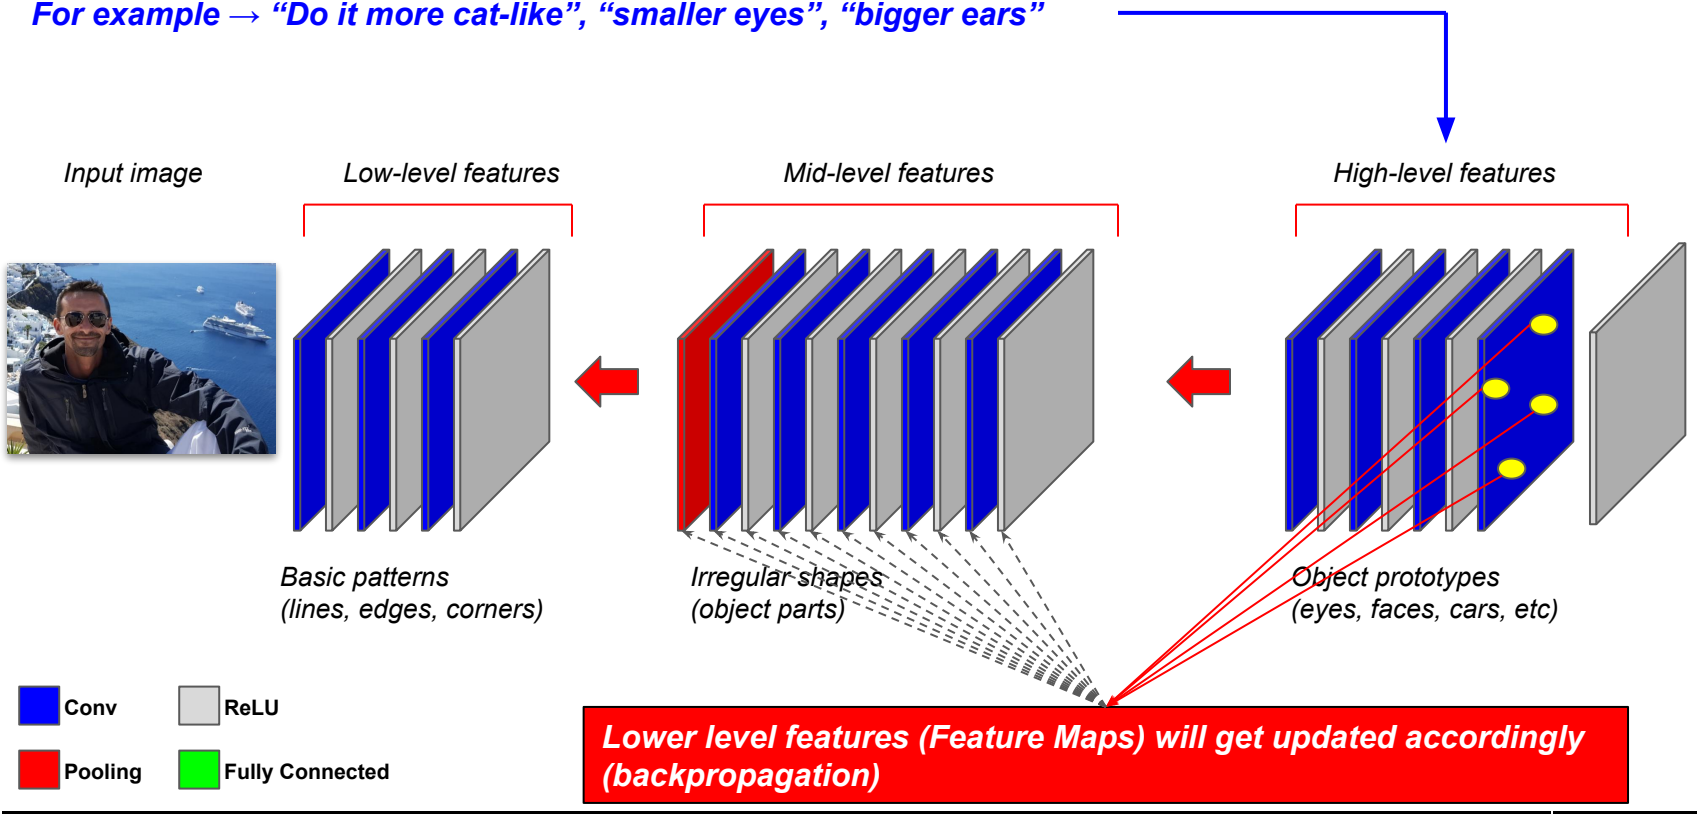

Dreams *(and nightmares)*

How about this? What if I boost mid-level features?

"inhibit rectangles", "maximize circular shapes"

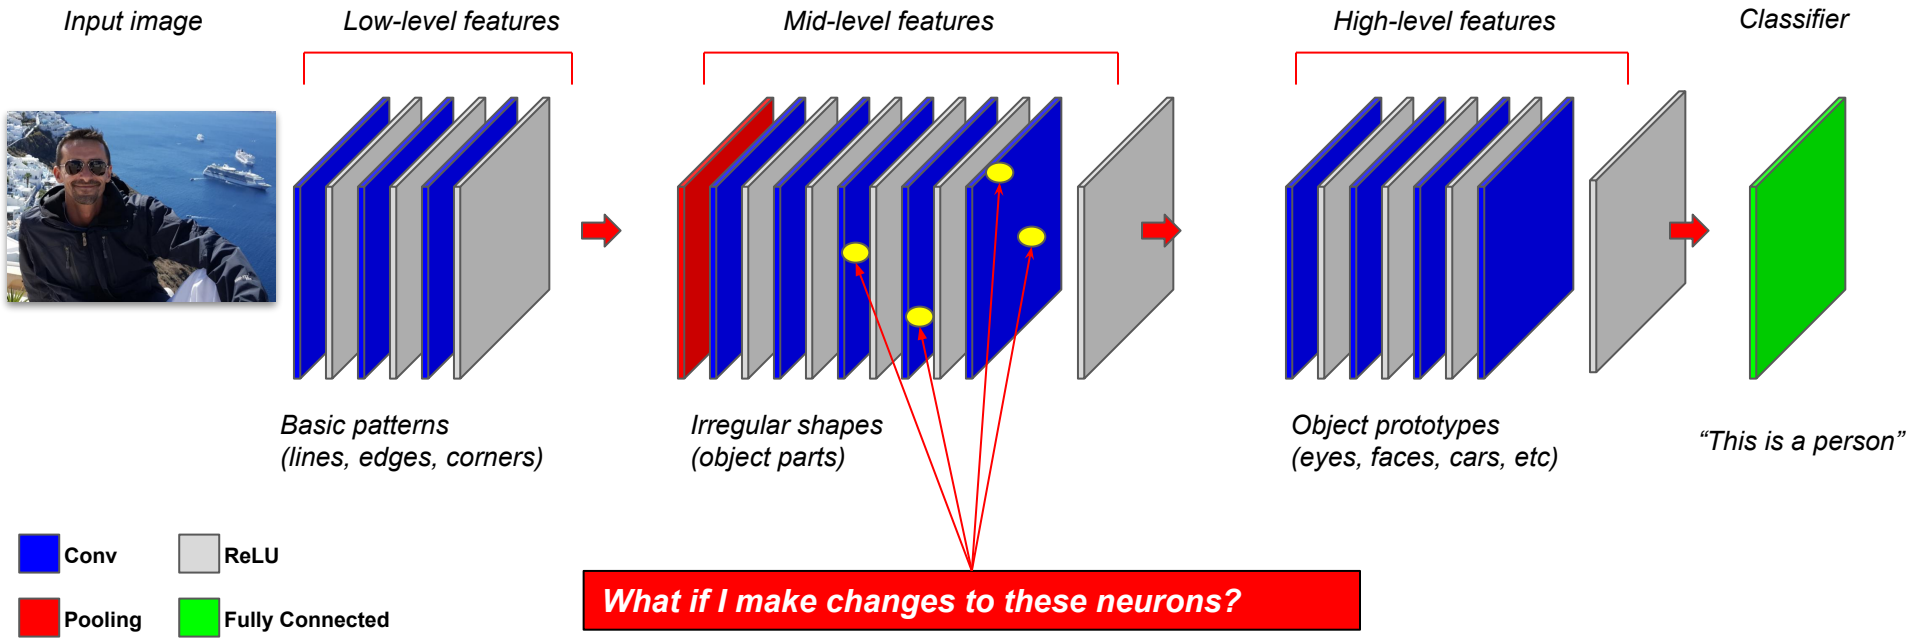

Dreams *(and nightmares)*

How about this? What if I boost mid-level features?

"inhibit rectangles", "maximize circular shapes"

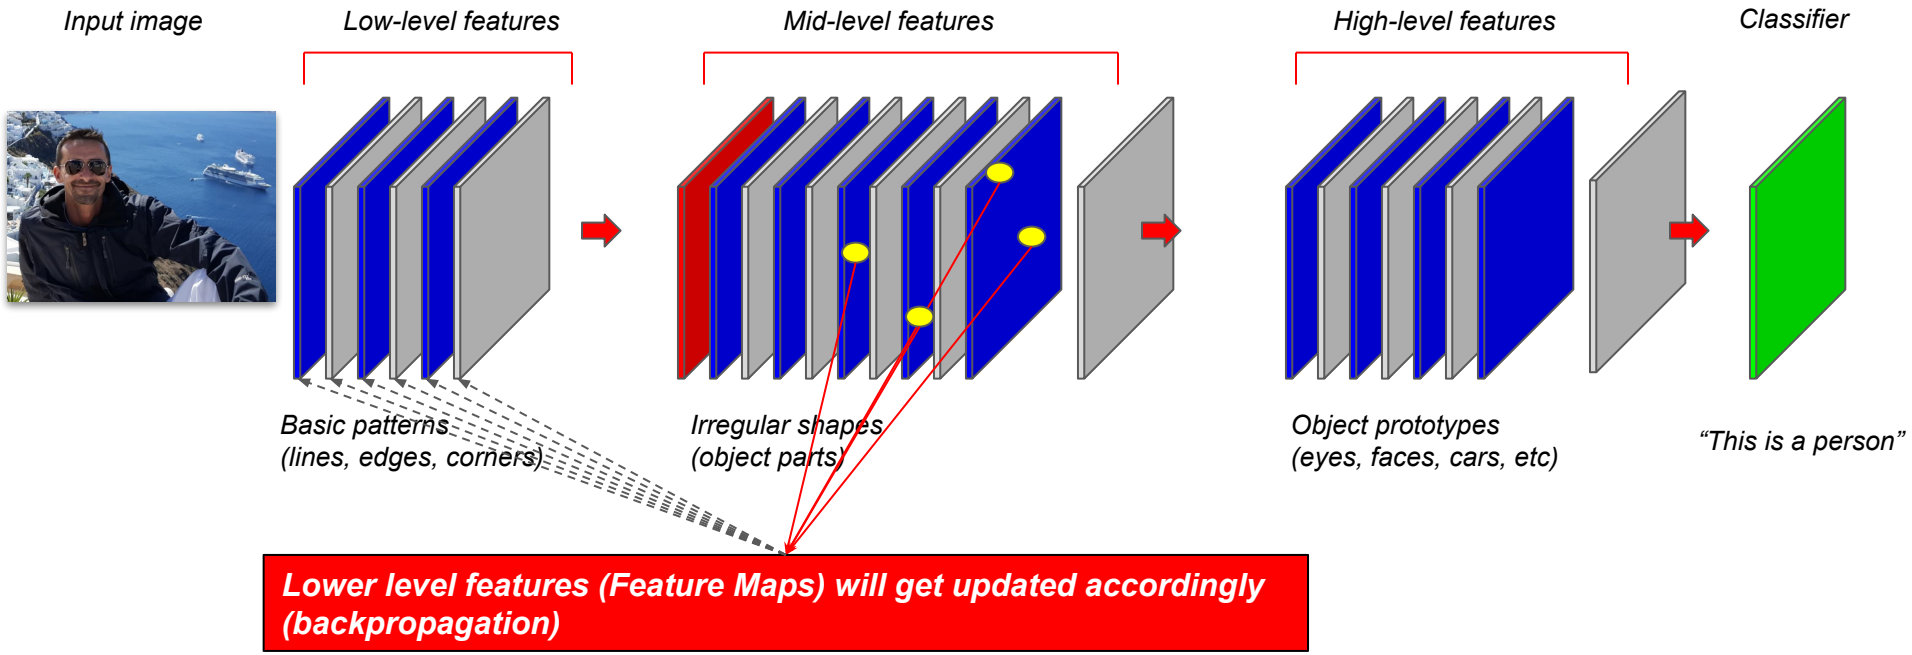

Lower level features (Feature Maps) will get updated accordingly (backpropagation)

Dreams - *Summary*

Dreaming consists in a boost of the values produced by inner feature maps.

Example → *“Any slightly detected bird face will be made more and more bird-like over time”*

Next Steps

What you get is the one you see	How to do it	What you get
Lazy mode	Download someone else's code. It takes 10 seconds. Try to learn about it.	**Just click on the *"Download"* button.** Quick and easy. Enjoy your copy/paste. *Don't forget to brag about it !!!*
A better way	Use a mature technology or framework: • *TensorFlow (my suggestion)* • *Caffe, Lasagne, Keras, ConvNetJs, etc.* *Google them!*	**Enjoy what others have created for you** • Click on the *"Download"* button. • Learn how to use the tool at hand. • Create your Neural Networks. **It's a great start. You should take the offer!**
The hard way, ***(for crazy people like me)*** *"Eat well, Spartan. Because this night we're gonna dine in **hell**"*	Create your own CNN from scratch, in any language of your preference. • C++, Java, C#, etc… **Earn your wings, Spartan!**	**This is Sparta!** Learning and fun is 110% guaranteed. You can trust me on this. ***For road warriors ONLY!***

Thank you !